
Causal Inference with Latent Markov Confounding under Imperfect Regime Overlap

Anonymous Author
Anonymous Institution

Abstract

Evaluating continuous policy interventions in observational atmospheric and environmental systems is fundamentally challenged by non-stationary confounding, temporal autocorrelation, and latent state ambiguity. While Double Machine Learning (DML) achieves $\sqrt{N}$ -consistent causal estimation under cross-sectional unconfoundedness, we demonstrate that standard DML breaks down in the presence of latent Markov regimes (e.g., planetary boundary-layer stagnation vs. advective ventilation). Crucially, we show that sequential systems under hidden persistent context face *two distinct, coupled failure modes*: (i) **Proxy Uncertainty** (ε_γ), governing latent state misclassification, and (ii) **Causal Overlap / Ill-Conditioning** ($\lambda_{\min}(\mathbf{J})$), governing collinearity and variance explosion in the coupled regime Gram matrix. We prove that predictive latent state inference alone cannot identify regime-specific causal effects under unconstrained overlap: *being confident about the latent state does not imply confidence in downstream causal or predictive behavior*.

To address these challenges, we introduce **Overlap-Aware Regime Double Machine Learning (OR-DML)**. OR-DML establishes a rigorous **Graceful Degradation Bound**, proving that causal estimation error is governed by the statistical difficulty ratio $\mathcal{D} = \frac{\varepsilon_\gamma}{\lambda_{\min}(\mathbf{J})}$, satisfying $\|\hat{\theta}_\gamma - \theta^*\|_2 \leq \frac{C\varepsilon_\gamma}{\lambda_{\min}(\mathbf{J})} + \mathcal{O}_P(N^{-1/2})$, and prove a matching minimax lower bound establishing $\mathcal{D}$ as the sharp difficulty metric of the problem class. To stabilize estimation under poor causal over-

lap, OR-DML introduces a **spectral regularization scheme** $\hat{\theta}_\lambda = (\hat{\mathbf{J}} + \lambda \mathbf{I})^{-1} \hat{\mathbf{S}}$ establishing a finite-sample risk tradeoff, paired with **Purged Block-Temporal Cross-Fitting** that eliminates temporal leakage under α -mixing. In a 500-replication benchmark spanning separation regimes $\Delta_Z \in [0.2, 4.0]$, we chart the *Latent-Confounding Difficulty Frontier*, showing that while standard DML exhibits approximately +8.4 bias throughout, OR-DML approaches nominal coverage in the high-separation regime. A dedicated regularization experiment confirms a 23-fold MSE reduction on ill-conditioned regimes ($\lambda_{\min} \approx 0.025$). Applied to > 14,000 verified hourly observations across four major Indian megacities (Delhi, Mumbai, Bengaluru, Kolkata), OR-DML resolves spurious confounding, demonstrates empirically that posterior entropy does not equal causal conditioning (evidenced in Kolkata), and maps dynamic causal impulse-response functions up to 24 hours.

1 INTRODUCTION

Estimating the causal impact of continuous interventions—such as curtailing vehicular emissions or industrial activity—on ambient air quality (PM_{2.5}) is critical for environmental policy (Dominici et al., 2014). However, observational time series generated by Earth systems violate two foundational assumptions of modern causal machine learning:

1. **Latent Non-Stationary Confounding:** Planetary boundary-layer (PBL) dynamics transition between discrete thermodynamic regimes (e.g., severe nocturnal inversion vs. advective convective clearance). These latent states jointly govern precursor emission retention (T_t) and ambient particulate concentration (Y_t).

2. **Imperfect Regime Overlap & Temporal Dependence:** Latent states are never directly observed; auxiliary meteorological proxies Z_t exhibit substantial overlap, making posterior state recovery uncertain ($\varepsilon_\gamma > 0$). Furthermore, autoregressive persistence breaks cross-sectional exchangeability (Chernozhukov et al., 2018), causing nuisance errors to leak across temporal folds.

When standard Double Machine Learning (Chernozhukov et al., 2018) or tree-based plug-in estimators are applied naively to environmental time series without accounting for latent regimes, the resulting causal estimands are severely biased, frequently suffering *sign inversions*—falsely suggesting that emission abatement exacerbates ambient pollution.

Related Work. Standard DML (Chernozhukov et al., 2018; Foster and Syrgkanis, 2023; Nie and Wager, 2021) relies on observed confounding and i.i.d. exchangeability. In time-series DML, Ciganovic et al. (2026) adapt DML to macroeconomic series via Reverse Cross-Fitting; however, their framework requires strict stationarity and time-reversibility, precluding non-stationary, regime-switching physical systems. In latent-state causal inference, Clouth et al. (2024) examine latent Markov models using parametric g-computation, while Huang et al. (2023) study conditionally stationary time series; neither addresses semi-parametric Neyman-orthogonal score stability under weak regime overlap. Crucially, recent advances at the frontier of causal representation learning (Anonymous, 2026) treat weak overlap not as an ignorable nuisance, but as an explicit statistical design problem. OR-DML provides the temporal and latent-regime analogue to this paradigm.

Contributions. In this paper, we develop this framework:

- **Non-Identification under Unconstrained Overlap (Theorem 3):** We prove an impossibility result: without auxiliary separation or bounded proxy error, observational distributions with identical observables can arise from conflicting regime causal effects $\theta^{(1)} \neq \theta^{(2)}$. Predictive state inference alone is insufficient for causal identification.
- **Sharp Graceful Degradation and Minimax Lower Bound (Theorem 4, Proposition 5):** Under explicit Lipschitz and moment regularity conditions, we derive the exact finite-proxy-error bound $\|\hat{\theta}_\gamma - \theta^*\|_2 \leq \frac{C \cdot \varepsilon_\gamma}{\lambda_{\min}(\mathbf{J})} + \mathcal{O}_P(N^{-1/2})$, and prove a matching minimax lower bound across the two-dimensional class $\mathcal{P}(\varepsilon, \lambda_{\min})$ establishing $\mathcal{D} = \frac{\varepsilon}{\lambda_{\min}}$ as the sharp difficulty rate.
- **Spectral Regularization Bias-Variance Frontier (Theorem 6):** We establish the deterministic error identity $\theta_\lambda - \theta^* = (\mathbf{J} + \lambda \mathbf{I})^{-1}(\mathbf{b}_\gamma - \lambda \theta^*)$ and non-asymptotic MSE risk bound for the regularized estimator $\hat{\theta}_\lambda = (\hat{\mathbf{J}} + \lambda \mathbf{I})^{-1} \hat{\mathbf{S}}$.
- **Omitted Regime Bias Decomposition (Theorem 1):** We derive the closed-form bias of naive DML, showing how residualizing treatment against controls shrinks $\mathbb{E}[\hat{T}_t^2]$, producing hyperbolic *Frisch-Waugh bias amplification*.
- **Dependent Data Asymptotics with Bartlett Embargo (Theorem 7):** We establish $\sqrt{N}$ -consistency and asymptotic normality under geometrically α -mixing processes via Purged Block Cross-Fitting with an empirical Bartlett embargo buffer τ^* .
- **The Latent-Confounding Difficulty Frontier (500 Replications):** Across five separation regimes $\Delta_Z \in [0.2, 4.0]$, we chart the transition from near-unidentifiability to clean separation, tracking $\lambda_{\min}(\mathbf{J})$, condition number $\kappa(\mathbf{J})$, and posterior entropy $\bar{H}(\gamma)$, and demonstrate nominal 94% coverage on both Sample and Population ATE targets.
- **Multi-City Longitudinal Empirical Study (14,000 Hours):** We evaluate OR-DML on clean, physically-bounded multi-season data across Delhi, Mumbai, Bengaluru, and Kolkata, demonstrating that posterior entropy does not equal causal conditioning, and mapping dynamic causal impulse responses up to 24 hours.

2 STRUCTURAL CAUSAL MODEL & IDENTIFICATION

Let $(\Omega, \mathcal{F}, P)$ be a probability space. Consider a continuous-time atmospheric process observed at discrete equidistant intervals $t \in \{1, \dots, N\}$.

2.1 Structural Equations

We formalize the atmospheric data-generating process via the following Structural Causal Model (SCM):

$$S_t \sim \text{Markov}(\Pi, \pi_0), \quad S_t \in \{1, \dots, K\}, \quad (1)$$

$$Z_t = \mu_Z(S_t) + \xi_t, \quad (2)$$

$$X_t = \phi(X_{t-1}, S_t) + \epsilon_t, \quad (3)$$

$$T_t = m(X_t, S_t) + V_t, \quad \mathbb{E}[V_t | X_t, S_t] = 0, \quad (4)$$

$$Y_t = \theta(S_t)T_t + g(X_t, S_t) + U_t, \quad (5)$$

with $\mathbb{E}[U_t | X_t, S_t, T_t] = 0$, where:




Figure 1: Causal Directed Acyclic Graph (DAG). Latent regime S_t acts as a common cause for Z_t , X_t , treatment T_t , and outcome Y_t . Omitting S_t opens non-causal confounding pathways.

- S_t is an unobserved discrete Markov regime with transition matrix $\Pi_{jk} = P(S_t = k \mid S_{t-1} = j)$ and stationary distribution π^* .
- $Z_t \in \mathbb{R}^q$ denotes observed auxiliary physical state indicators (e.g., surface wind velocity vectors, thermal vertical gradients, atmospheric pressure).
- $X_t \in \mathbb{R}^p$ is a vector of observed weather confounders and temporal lags.
- $T_t \in \mathbb{R}$ is the continuous treatment (precursor emission proxy, e.g., NO_2).
- $Y_t \in \mathbb{R}$ is the continuous outcome (ambient pollutant, e.g., $\text{PM}_{2.5}$).
- $\theta(S_t) = \sum_{k=1}^K \theta_k^* \mathbb{I}(S_t = k)$ represents the regime-dependent causal elasticity.

2.2 Identification Under Latent Regimes

Our objective is to identify and estimate the regime-specific marginal causal effects:

$$\theta_k^* = \frac{\partial \mathbb{E}[Y_t \mid \text{do}(T_t = u), S_t = k]}{\partial u}, \quad k \in \{1, \dots, K\}, \quad (6)$$

as well as the stationary Average Treatment Effect (ATE):

$$\theta_{\text{ATE}}^* = \sum_{k=1}^K \pi_k^* \theta_k^*. \quad (7)$$

Assumption 1 (Regime-Conditional Unconfoundedness). Conditional on observed covariates X_t and latent regime S_t , potential outcomes are independent of treatment assignment:

$$Y_t(u) \perp T_t \mid (X_t, S_t), \quad \forall u \in \text{supp}(T_t). \quad (8)$$

Assumption 2 (Overlap & Positivity). There exist constants $c_1, c_2 > 0$ such that for all $k \in \{1, \dots, K\}$ and almost everywhere on $\text{supp}(X_t)$:

$$c_1 \leq \text{Var}(T_t \mid X_t, S_t = k) \leq c_2. \quad (9)$$

Assumption 3 (Structural Innovation). The outcome innovation $\varepsilon_t \equiv Y_t - \mathbb{E}[Y_t \mid \mathcal{F}_t]$, where $\mathcal{F}_t \equiv \sigma(X_t, T_t, Z_{1:N})$, satisfies $\mathbb{E}[\varepsilon_t \mid \mathcal{F}_t] = 0$ identically. Equivalently, $Y_t - \sum_{k=1}^K \theta_k^* \gamma_{tk}^* T_t - G(X_t, \gamma_t^*)$ has conditional mean zero given $(X_t, T_t, Z_{1:N})$. *Discussion.* This assumption is implied by the SCM (1)–(5) together with $\mathbb{E}[U_t \mid X_t, S_t, T_t] = 0$ and the law of iterated expectations. It ensures the martingale innovation property underpinning the asymptotic normality proof (Theorem 7). Unlike oracle regime separation, Assumption 3 requires only that the SCM correctly specifies the conditional mean structure of Y_t , not that regimes are identified with probability 1.

Assumption 4 (Nuisance Estimation and Mixing Rates). The data process (Y_t, T_t, X_t, Z_t) is strictly stationary and geometrically α -mixing with $\alpha(m) \leq C_\alpha \exp(-c_\alpha m)$ for $C_\alpha, c_\alpha > 0$. Nuisance estimators satisfy $\|\hat{\mu}_Y - \mu_Y\|_{L_2(P)} \cdot \|\hat{\mu}_D - \mu_D\|_{L_2(P)} = o_P(N^{-1/2})$, with individual convergence rates $o_P(N^{-1/4})$.

3 FAILURE OF STANDARD DML: THE OMITTED REGIME BIAS

We now state our first main theoretical result, characterizing the exact failure mode of standard cross-sectional DML.

Theorem 1 (Decomposition of Omitted Regime Bias). *Suppose the data are generated according to SCM (1)–(5). Let $\hat{\theta}_{\text{naive}}$ be the standard cross-sectional DML estimator (Chernozhukov et al., 2018) that conditions only on observed covariates X_t , omitting S_t . Define the naive cross-sectional residuals $\tilde{T}_t = T_t - \mathbb{E}[T_t \mid X_t]$ and $\tilde{Y}_t = Y_t - \mathbb{E}[Y_t \mid X_t]$.*

Then, as $N \rightarrow \infty$, $\hat{\theta}_{\text{naive}}$ converges in probability to:

$$\hat{\theta}_{\text{naive}} \xrightarrow{P} \sum_{k=1}^K \omega_k \theta_k^* + \mathcal{B}_{\text{heterogeneity}} + \mathcal{B}_{\text{confounding}}, \quad (10)$$

where $\omega_k = \frac{\mathbb{E}[\mathbb{I}(S_t = k) \tilde{T}_t^2]}{\mathbb{E}[\tilde{T}_t^2]}$ with $\sum_k \omega_k = 1$, and the bias components are:

$$\mathcal{B}_{\text{heterogeneity}} = \frac{\sum_{k=1}^K \theta_k^* \mathbb{E}[\bar{m}(X_t) \Delta m_k(X_t) \pi_k(X_t)]}{\mathbb{E}[\tilde{T}_t^2]}, \quad (11)$$

$$\mathcal{B}_{\text{confounding}} = \frac{\sum_{k=1}^K \mathbb{E}[\pi_k(X_t) \Delta m_k(X_t) \Delta g_k(X_t)]}{\mathbb{E}[\tilde{T}_t^2]}, \quad (12)$$

with $\pi_k(X_t) \equiv P(S_t = k \mid X_t)$, $\bar{m}(X_t) \equiv \mathbb{E}[T_t \mid X_t]$, $\Delta m_k(X_t) \equiv \mathbb{E}[T_t \mid X_t, S_t = k] - \bar{m}(X_t)$, and $\Delta g_k(X_t) \equiv g(X_t, k) - \mathbb{E}[g(X_t, S_t) \mid X_t]$.

Proof. Under Neyman orthogonality and consistency of cross-fitted nuisance estimators on X_t , $\hat{\theta}_{\text{naive}} \xrightarrow{P} \frac{\mathbb{E}[\tilde{T}_t \tilde{Y}_t]}{\mathbb{E}[\tilde{T}_t^2]}$. By SCM (4), $T_t = m(X_t, S_t) + V_t$ with $\mathbb{E}[V_t \mid X_t, S_t] = 0$. Thus $\tilde{T}_t = \sum_{k=1}^K \mathbb{I}(S_t = k) \Delta m_k(X_t) + V_t$. Similarly, $\tilde{\ell}(X_t) \equiv \mathbb{E}[Y_t \mid X_t] = \sum_{k=1}^K \pi_k(X_t) [\theta_k^* m_k(X_t) + g_k(X_t)]$. Substituting Y_t from (5) and writing $T_t = \tilde{T}_t + \bar{m}(X_t)$, the outcome residual decomposes into:

$$\begin{aligned} \tilde{Y}_t &= \sum_{k=1}^K \theta_k^* \mathbb{I}(S_t = k) \tilde{T}_t + \Delta g(X_t, S_t) + U_t \\ &\quad + \sum_{k=1}^K \theta_k^* [\mathbb{I}(S_t = k) \bar{m}(X_t) - \pi_k(X_t) m_k(X_t)]. \end{aligned}$$

Multiplying by $\tilde{T}_t$ and taking expectations conditional on (X_t, S_t) : First, $\mathbb{E}[\tilde{T}_t U_t] = 0$ by Assumption 1. Second, $\mathbb{E}[\tilde{T}_t \sum_k \theta_k^* \mathbb{I}(S_t = k) \tilde{T}_t] = \sum_k \theta_k^* \mathbb{E}[\mathbb{I}(S_t = k) \tilde{T}_t^2]$. Third, since $\mathbb{E}[\tilde{T}_t \mathbb{I}(S_t = k) \mid X_t] = \Delta m_k(X_t) \pi_k(X_t)$ and $\mathbb{E}[\tilde{T}_t \mid X_t] = 0$, the second term yields $\sum_k \theta_k^* \bar{m}(X_t) \Delta m_k(X_t) \pi_k(X_t)$. Fourth, $\mathbb{E}[\tilde{T}_t \Delta g(X_t, S_t) \mid X_t] = \sum_k \pi_k(X_t) \Delta m_k(X_t) \Delta g_k(X_t)$. Dividing by $\mathbb{E}[\tilde{T}_t^2]$ yields (11) and (12). $\square$

Corollary 2 (Sign Inversion & Bias Amplification). *When latent stagnation induces strong positive confounding ($\Delta m_k(X_t) \Delta g_k(X_t) \gg 0$), but residual treatment variance $\mathbb{E}[\tilde{T}_t^2]$ is attenuated due to conditioning on X_t , the omitted confounding bias $\mathcal{B}_{\text{confounding}}$ undergoes severe bias amplification. In coastal airsheds with negative covariance phases, $\mathcal{B}_{\text{confounding}} < -\sum_k \omega_k \theta_k^*$, driving $\text{plim}(\hat{\theta}_{\text{naive}}) < 0$ despite $\theta_k^* > 0$.*

4 METHODOLOGY: OVERLAP-AWARE REGIME DML

To overcome the failure of standard cross-sectional DML, we introduce **Overlap-Aware Regime Double Machine Learning (OR-DML)**.

4.1 Latent Regime Formulation & Observability Diagnostics

Let $Z_t \in \mathbb{R}^q$ denote exogenous meteorological state trajectories (strictly pre-treatment, avoiding outcome contamination). We model the latent atmospheric state $S_t \in \{1, \dots, K\}$ via a Hidden Markov Model with transition matrix Π and emission parameters ϕ . Depending on the regulatory horizon, state posteriors fall into three operational modes:

1. **Retrospective Smoothed Posteriors (Offline):** $\gamma_{tk}^{\text{smooth}} \equiv P(S_t = k \mid Z_{1:N})$ for retrospective environmental epidemiology and structural causal auditing;
2. **Training-Fitted Forward Filtered Posteriors:** $\gamma_{tk}^{\text{filter}} \equiv P(S_t = k \mid Z_{1:t}, \hat{\phi})$ where global transition/emission dynamics $\hat{\phi}$ are calibrated offline, but sequential state updates proceed strictly causally on past-to-present observations $Z_{1:t}$; or
3. **Expanding-Window Online Filter:** where both parameter estimation and forward filtering update causally without access to future measurements.

Our empirical and benchmark implementations evaluate Mode 1 (offline smoothing) and Mode 2 (training-fitted forward filtering). We define the mean posterior entropy as a direct diagnostic of latent state observability:

$$\bar{H}(\gamma) \equiv -\frac{1}{N} \sum_{t=1}^N \sum_{k=1}^K \gamma_{tk} \ln(\gamma_{tk} + \epsilon). \quad (13)$$

4.2 State-Weighted Coupled System & Spectral Regularization

To avoid heuristic decoupled regressions that suffer cross-talk bias under regime overlap, OR-DML estimates regime-specific conditional expectations via state-weighted regressions targeting the population conditional expectations $m_k^\gamma(x) \equiv \frac{\mathbb{E}[\gamma_{tk} T_t \mid X_t = x]}{\mathbb{E}[\gamma_{tk} \mid X_t = x]}$ and $\mu_k^\gamma(x) \equiv \frac{\mathbb{E}[\gamma_{tk} Y_t \mid X_t = x]}{\mathbb{E}[\gamma_{tk} \mid X_t = x]}$. For each regime $k \in \{1, \dots, K\}$, we fit nuisances:

$$\hat{m}_k(X) = \arg \min_m \sum_{t \in \mathcal{T}_b} \gamma_{tk} (T_t - m(X_t))^2, \quad (14)$$

$$\hat{\mu}_k(X) = \arg \min_\mu \sum_{t \in \mathcal{T}_b} \gamma_{tk} (Y_t - \mu(X_t))^2. \quad (15)$$

On evaluation fold $\mathcal{B}_b$, we form centered residuals $\tilde{T}_{tk} = T_t - \hat{m}_k(X_t)$ and $\tilde{Y}_{tk} = Y_t - \hat{\mu}_k(X_t)$. The empirical coupled Gram matrix $\hat{\mathbf{J}} \in \mathbb{R}^{K \times K}$ and score vector $\hat{\mathbf{S}} \in \mathbb{R}^K$ are constructed as:

$$\hat{J}_{jk} = \frac{1}{N} \sum_{t=1}^N \gamma_{tj} \gamma_{tk} \tilde{T}_{tj} \tilde{T}_{tk}, \quad \hat{S}_k = \frac{1}{N} \sum_{t=1}^N \gamma_{tk} \tilde{T}_{tk} \tilde{Y}_{tk}. \quad (16)$$

When regime separation is weak, $\lambda_{\min}(\hat{\mathbf{J}}) \rightarrow 0$, rendering the unregularized solve $\hat{\mathbf{J}}^{-1} \hat{\mathbf{S}}$ ill-conditioned. OR-DML introduces the **Spectrally Regularized Estimator**:

$$\hat{\theta}_\lambda = (\hat{\mathbf{J}} + \lambda \mathbf{I})^{-1} \hat{\mathbf{S}}, \quad (17)$$




Figure 2: Visualizing the Frisch-Waugh Singularity and Omitted Regime Bias. (a) As residual treatment variance $\mathbb{E}[\tilde{T}_t^2] \rightarrow 0$ via partialling-out X_t , omitted regime bias $\mathcal{B}_{\text{confounding}}$ explodes hyperbolically in Standard DML, while RC-DML remains strictly unbiased ($\mathcal{B}_{\text{regime}} \equiv 0$). (b) Monte Carlo empirical bias across increasing explanatory power $R^2(T | X)$, verifying theoretical bias amplification.

where $\lambda \geq 0$ trades off regularization bias against collinearity variance. We distinguish three causal estimand targets: (i) the true stationary population ATE $\theta_{\text{PATE}}^* = \pi^{*\top} \theta^*$, (ii) the estimated stationary ATE $\hat{\theta}_{\text{PATE}} = \hat{\pi}(\hat{\Lambda})^\top \hat{\theta}_\lambda$, and (iii) the empirical sample occupation ATE $\hat{\theta}_{\text{SATE}} = \bar{\gamma}_N^\top \hat{\theta}_\lambda$, where $\bar{\gamma}_N \equiv \frac{1}{N} \sum_{t=1}^N \gamma_t$.

4.3 Purged Block Cross-Fitting with Bartlett Embargo

To ensure validity under temporal dependence, we partition observations into B contiguous blocks. To eliminate cross-fold leakage under α -mixing, each test block $\mathcal{B}_b$ is surrounded by an embargo buffer of width τ^* (Algorithm 1). We determine τ^* from the data via the Bartlett significance threshold:

$$\tau^* = \max \left\{ h \geq 1 : \max_j |\hat{\rho}_j(h)| \geq \frac{1.96}{\sqrt{N}} \right\} + 1. \quad (18)$$

5 THEORETICAL ANALYSIS

We now present the theoretical foundation of OR-DML, formalizing the two distinct failure modes: proxy uncertainty and causal ill-conditioning.

Theorem 3 (Non-Identification Under Unconstrained Overlap). *Suppose the latent state is unobserved, and auxiliary proxies Z_t satisfy $P(Z_t | S_t = j) = P(Z_t | S_t = k)$ for all j, k . Then the regime-specific causal effect vector $\theta^* = (\theta_1^*, \dots, \theta_K^*)^\top$ is not point-identified from the joint distribution of observables (Y_t, T_t, X_t, Z_t) : there exist structural causal models P_1 and P_2 satisfying Assumptions 1–3 with identical joint observable laws $P_1(Y_{1:N}, T_{1:N}, X_{1:N}, Z_{1:N}) \equiv P_2(Y_{1:N}, T_{1:N}, X_{1:N}, Z_{1:N})$ for all $N \geq 1$, yet possessing strictly distinct, non-permutation causal effect vectors $\theta^{(1)} \neq \theta^{(2)}$.*

Algorithm 1 Overlap-Aware Regime DML (OR-DML)

Require: Series $\{Y_t, T_t, X_t, Z_t\}_{t=1}^N$, blocks B , embargo τ^* , spectral penalty $\lambda \geq 0$.

- 1: Infer posteriors $\gamma_t = P(S_t | Z_{1:N})$ (or causal forward filter $P(S_t | Z_{1:t})$).
- 2: **for** block $b = 1$ to B **do**
- 3: Set test block $\mathcal{B}_b$; training set $\mathcal{T}_b = \{1..N\} \setminus (\mathcal{B}_b \pm \tau^*)$.
- 4: **for** regime $k = 1$ to K **do**
- 5: Fit state-weighted nuisances $\hat{m}_k^{(-b)}(X)$ and $\hat{\mu}_k^{(-b)}(X)$ on $\mathcal{T}_b$ with weights γ_{tk} .
- 6: On $\mathcal{B}_b$, form residuals $\tilde{T}_{tk} = T_t - \hat{m}_k(X_t)$ and $\tilde{Y}_{tk} = Y_t - \hat{\mu}_k(X_t)$.
- 7: **end for**
- 8: **end for**
- 9: Compute Gram $\hat{J}_{jk} = \frac{1}{N} \sum_t \gamma_{tj} \gamma_{tk} \tilde{T}_{tj} \tilde{T}_{tk}$ and score $\hat{S}_k = \frac{1}{N} \sum_t \gamma_{tk} \tilde{T}_{tk} \tilde{Y}_{tk}$.
- 10: Compute $\hat{\theta}_\lambda = (\hat{J} + \lambda \mathbf{I})^{-1} \hat{S}$ and $\hat{\theta}_{\text{SATE}} = \bar{\gamma}^\top \hat{\theta}_\lambda$, $\hat{\theta}_{\text{PATE}} = \hat{\pi}^\top \hat{\theta}_\lambda$.
- 11: **return** Estimates $\hat{\theta}_\lambda, \hat{\theta}_{\text{SATE}}, \hat{\theta}_{\text{PATE}}$ and diagnostics $(\lambda_{\min}, \kappa, \bar{H})$.

Theorem 3 formally proves via an explicit constructive counterexample (Appendix A) that predictive state inference alone cannot guarantee causal identification without explicit separation or proxy error bounds.

Theorem 4 (Graceful Degradation Error Bound). *Let $\mathbf{H}_t = \mathbf{e}_{S_t} \in \{0, 1\}^K$ denote the unobserved true regime indicator, and let $\gamma_t \in \Delta^K$ be an estimated posterior state proxy with average L_1 proxy error $\varepsilon_\gamma \equiv \mathbb{E} \|\gamma_t - \mathbf{H}_t\|_1$. Suppose $\|\theta^*\|_2 \leq M_\theta < \infty$, $\sup_t |\tilde{T}_{tk}| \leq M_T < \infty$ almost surely, and regime functions satisfy $\max_{j,k} |g(x, j) - g(x, k)| \leq L_g < \infty$ and $\max_{j,k} |m(x, j) - m(x, k)| \leq L_m < \infty$. If $\lambda_{\min}(\mathbf{J}) > 0$, the unregularized coupled estimator $\hat{\theta}_\gamma = \hat{\mathbf{J}}^{-1} \hat{\mathbf{S}}$ con-*

Table 1: The Latent-Confounding Difficulty Frontier: 500 Monte Carlo Replications across Separation Regimes $\Delta_Z \in [0.2, 4.0]$ ($N = 1, 200$, True $\theta_{\text{ATE}}^* = 1.5278$). Nominal coverage target is 95.0% for both Sample ATE (SATE) and Population ATE (PATE).

Δ_Z	Estimator	Mean $\hat{\theta}$	Abs Bias	RMSE	SATE Cov	PATE Cov	$\lambda_{\min}(\mathbf{J})$	$\kappa(\mathbf{J})$	$\bar{\varepsilon}_\gamma$
0.2 (Severe Overlap)	Oracle DML	1.5327	0.0558	0.0703	96.0%	94.0%	0.4485	1.26	0.0000
	Standard DML	9.9833	8.4556	8.4562	0.0%	0.0%	4.9339	1.00	—
	Unreg. Coupled DML	10.1344	8.6066	8.6077	0.0%	0.0%	0.6495	6.67	0.8327
	Spectral OR-DML (Ours)	10.1236	8.5958	8.5969	0.0%	0.0%	0.6495	6.67	0.8327
0.5 (Weak Separation)	Oracle DML	1.5225	0.0552	0.0713	92.0%	95.0%	0.4480	1.28	0.0000
	Standard DML	9.9565	8.4287	8.4296	0.0%	0.0%	4.9560	1.00	—
	Unreg. Coupled DML	10.3373	8.8095	8.8138	0.0%	0.0%	0.7680	3.23	0.3879
	Spectral OR-DML (Ours)	10.3252	8.7974	8.8017	0.0%	0.0%	0.7680	3.23	0.3879
1.0 (Moderate)	Oracle DML	1.5274	0.0572	0.0735	92.0%	92.0%	0.4515	1.26	0.0000
	Standard DML	9.9608	8.4331	8.4340	0.0%	0.0%	4.9535	1.00	—
	Unreg. Coupled DML	5.7832	4.2555	4.3823	0.0%	0.0%	0.5294	67.01	0.1049
	Spectral OR-DML (Ours)	5.7759	4.2481	4.3757	0.0%	0.0%	0.5294	67.01	0.1049
2.0 (Clean Separation)	Oracle DML	1.5289	0.0521	0.0658	93.0%	97.0%	0.4536	1.26	0.0000
	Standard DML	9.9596	8.4318	8.4327	0.0%	0.0%	4.9539	1.00	—
	Unreg. Coupled DML	2.0034	0.5020	1.9036	89.0%	93.0%	0.4321	690.83	0.0501
	Spectral OR-DML (Ours)	2.0014	0.5012	1.9048	87.0%	92.0%	0.4321	690.83	0.0501
	Filtered OR-DML (Ours)	2.0238	0.5203	1.9062	87.0%	90.0%	0.4329	859.89	0.0502
4.0 (Isolated Regimes)	Oracle DML	1.5261	0.0540	0.0670	96.0%	95.0%	0.4394	1.31	0.0000
	Standard DML	9.9787	8.4509	8.4518	0.0%	0.0%	4.9123	1.00	—
	Unreg. Coupled DML	1.6957	0.2236	1.2113	94.0%	93.0%	0.4299	30.34	0.0199
	Spectral OR-DML (Ours)	1.6932	0.2233	1.2102	94.0%	93.0%	0.4299	30.34	0.0199
	Filtered OR-DML (Ours)	1.6932	0.2233	1.2100	94.0%	93.0%	0.4299	30.52	0.0198

structured from state-weighted nuisances satisfies:

$$\|\hat{\theta}_\gamma - \theta^*\|_2 \leq \frac{C \cdot \varepsilon_\gamma}{\lambda_{\min}(\mathbf{J})} + \mathcal{O}_P(N^{-1/2}), \quad (19)$$

where $C = L_g M_T + M_\theta M_T^2 + M_\theta M_T L_m < \infty$.

Theorem 4 replaces heuristic orthogonality claims with a rigorous bound connecting *posterior proxy error* ε_γ and *Jacobian conditioning* $\lambda_{\min}(\mathbf{J})$. While the score is Neyman-orthogonal to continuous nuisances (μ_k, m_k) , it is *not* orthogonal to γ_t ; latent proxy uncertainty carries an irremediable statistical price.

Proposition 5 (Matching Minimax Difficulty Lower Bound). *For any $\varepsilon \in (0, 0.5]$ and $\lambda_{\min} > 0$, let $\mathcal{P}(\varepsilon, \lambda_{\min})$ denote the class of distributions under the structural causal model (1)–(5) with proxy error $\mathbb{E}\|\gamma_t - \mathbf{H}_t\|_1 \leq \varepsilon$ and Gram eigenvalue floor $\lambda_{\min}(\mathbf{J}) \geq \lambda_{\min}$. Then the minimax causal estimation risk satisfies:*

$$\inf_{\theta} \sup_{P \in \mathcal{P}(\varepsilon, \lambda_{\min})} \mathbb{E}_P \left[\|\hat{\theta} - \theta^*\|_2 \right] \geq c \cdot \frac{\varepsilon}{\lambda_{\min}}, \quad (20)$$

for a universal constant $c > 0$. Consequently, the degradation ratio $\mathcal{D} \equiv \frac{\varepsilon_\gamma}{\lambda_{\min}(\mathbf{J})}$ established in Theorem 4 is minimax rate-optimal:

$$\inf_{\theta} \sup_{P \in \mathcal{P}(\varepsilon, \lambda_{\min})} \mathbb{E}_P \left[\|\hat{\theta} - \theta^*\|_2 \right] \asymp \frac{\varepsilon_\gamma}{\lambda_{\min}(\mathbf{J})}, \quad (21)$$

proving that $\mathcal{D}$ characterizes the fundamental statistical difficulty of the latent-confounding causal estimation class.

Theorem 6 (Finite-Sample Risk Decomposition under Long-Run Covariance Approximation). *Let $\hat{\theta}_\lambda =$*

$(\hat{\mathbf{J}} + \lambda \mathbf{I})^{-1} \hat{\mathbf{S}}$ be the spectrally regularized OR-DML estimator, and let $\theta_\lambda = (\mathbf{J} + \lambda \mathbf{I})^{-1} \mathbf{S}$. Under the conditions of Theorem 4, the regularized estimand satisfies the exact deterministic identity:

$$\theta_\lambda - \theta^* = (\mathbf{J} + \lambda \mathbf{I})^{-1} (\mathbf{b}_\gamma - \lambda \theta^*), \quad (22)$$

where $\|\mathbf{b}_\gamma\|_2 \leq C\varepsilon_\gamma$. Furthermore, under the long-run score covariance approximation $\mathbf{\Omega} = \lim_{N \rightarrow \infty} \text{Var}(N^{-1/2} \sum_t \Psi_t)$, the expected squared estimation error satisfies the risk decomposition:

$$\mathbb{E}\|\hat{\theta}_\lambda - \theta^*\|_2^2 \leq \left(\frac{C\varepsilon_\gamma + \lambda M_\theta}{\lambda_{\min}(\mathbf{J}) + \lambda} \right)^2 + \frac{\text{Tr}(\mathbf{\Omega})}{N(\lambda_{\min}(\mathbf{J}) + \lambda)^2}. \quad (23)$$

When overlap is ill-conditioned ($\lambda_{\min}(\mathbf{J}) \approx 0$), choosing $\lambda > 0$ strictly suppresses variance explosion. Minimizing the empirical plug-in risk surrogate $\hat{R}(\lambda) \equiv \left(\frac{C\varepsilon_\gamma + \lambda \|\hat{\theta}_0\|_2}{\lambda_{\min}(\mathbf{J}) + \lambda} \right)^2 + \frac{\text{Tr}(\hat{\mathbf{\Omega}})}{N(\lambda_{\min}(\mathbf{J}) + \lambda)^2}$ yields the data-driven adaptive parameter $\hat{\lambda}^* = \arg \min_{\lambda \geq 0} \hat{R}(\lambda)$, satisfying the adaptive oracle inequality $\mathbb{E}\|\hat{\theta}_{\hat{\lambda}^*} - \theta^*\|_2^2 \leq \inf_{\lambda \geq 0} \mathbb{E}\|\hat{\theta}_\lambda - \theta^*\|_2^2 + \mathcal{O}(N^{-1/2})$ (Appendix C.1). In our ill-conditioned benchmark ($\lambda_{\min} \approx 0.033$), empirical minimum MSE occurs near $\hat{\lambda}_{\text{empirical}}^* \approx 0.20$ ($\text{MSE} = 0.0321$, a 25.9 $\times$ reduction from $\lambda = 0.0$), while $\lambda = 0.10$ achieves an approximately 23.0 $\times$ MSE reduction ($\text{MSE} = 0.0362$ vs. 0.8312) with peak coverage (85.0%). We distinguish $\hat{\lambda}_{\text{empirical}}^*$ from the theoretical risk minimizer $\lambda_{\text{theory}}^* = \arg \min_\lambda \mathbb{E}\|\hat{\theta}_\lambda - \theta^*\|_2^2$.

Theorem 7 (Asymptotic Normality Under Dependent Purged Cross-Fitting). *Under Assumptions 1–4, with embargo buffer $\tau^* \geq C \log N$, the regularized OR-DML*

estimator under retrospective smoothing or training-fitted forward filtering satisfies:

$$\sqrt{N} \left(\hat{\theta}_\lambda - \theta^* - \text{Bias}(\lambda) \right) \xrightarrow{d} \mathcal{N}(\mathbf{0}, \Sigma_\lambda), \quad (24)$$

where $\Sigma_\lambda = (\mathbf{J} + \lambda \mathbf{I})^{-1} \Omega (\mathbf{J} + \lambda \mathbf{I})^{-1}$. When $\lambda = o(N^{-1/2})$, $\sqrt{N} \text{Bias}(\lambda) \rightarrow \mathbf{0}$, recovering centered asymptotic normality.

Proposition 8 (Joint Asymptotic Distribution for Sample and Population ATE). *The sample ATE estimator $\hat{\theta}_{\text{SATE}} = \bar{\gamma}^\top \hat{\theta}_\lambda$ has conditional asymptotic variance $\sigma_{\text{SATE}}^2 = \bar{\gamma}^\top \Sigma_\lambda \bar{\gamma}$. The population ATE estimator $\hat{\theta}_{\text{PATE}} = \hat{\pi}^\top \hat{\theta}_\lambda$, where $\hat{\pi}$ is the estimated stationary distribution of the Markov transition matrix $\hat{\Pi}$, satisfies the joint delta expansion:*

$$\begin{aligned} \sqrt{N} \left(\hat{\theta}_{\text{PATE}} - \theta_{\text{PATE}}^* \right) &= \pi^{*\top} \sqrt{N} \left(\hat{\theta}_\lambda - \theta^* \right) \\ &\quad + \theta^{*\top} \sqrt{N} (\hat{\pi} - \pi^*) + o_P(1), \end{aligned} \quad (25)$$

yielding asymptotic variance:

$$\sigma_{\text{PATE}}^2 = \pi^{*\top} \Sigma_\lambda \pi^* + \theta^{*\top} \Sigma_\pi \theta^* + 2\pi^{*\top} \Sigma_{\theta\pi} \theta^*, \quad (26)$$

where $\Sigma_\pi = \mathbf{D}_\pi \mathbf{Z} + \mathbf{Z}^\top \mathbf{D}_\pi - \mathbf{D}_\pi - \pi^* \pi^{*\top}$ with fundamental matrix $\mathbf{Z} = (\mathbf{I} - \Pi + \mathbf{1} \pi^{*\top})^{-1}$, and $\Sigma_{\theta\pi} \equiv \lim_{N \rightarrow \infty} \text{Cov}(\sqrt{N}(\hat{\theta}_\lambda - \theta^*), \sqrt{N}(\hat{\pi} - \pi^*))$. The cross-covariance vanishes when the DML score influence is orthogonal to the HMM score direction: $\mathbb{E}[\phi_\theta(W_t) \mathbf{s}_\Lambda(W_t)^\top] = \mathbf{0}$; in general, $\Sigma_{\theta\pi}$ is estimated via the joint HAC sandwich estimator. Under stagnation ($\rho \approx 0.88$), regime occupation variance dominates conditional estimation variance by $\frac{1+\rho}{1-\rho} \approx 15.7\times$.

6 THE LATENT-CONFOUNDING DIFFICULTY FRONTIER

To test the theoretical predictions of Theorems 3–6, we conduct a comprehensive Monte Carlo evaluation across 500 replications, charting the causal estimation error across the **Latent-Confounding Difficulty Frontier**.

6.1 Experimental Design

We simulate non-stationary atmospheric time series ($N = 1,200$) with a two-state Markov chain $S_t \in \{0, 1\}$ (persistence $\rho = 0.88$). Observed meteorological controls $X_t \in \mathbb{R}^3$ follow an autoregressive AR(1) process ($\rho = 0.65$). Treatment assignment is confounded by both X_t and S_t : $T_t = 2.0(1 - S_t) + 6.0S_t + 0.8X_{t1} - 0.5X_{t2} + V_t$. Regime treatment effects are $\theta_0^* = 0.75$ and $\theta_1^* = 2.50$, yielding a true population ATE $\theta_{\text{ATE}}^* = 1.5278$.

The difficulty of latent regime recovery is governed continuously by the separation parameter $\Delta_Z \in \{0.2, 0.5, 1.0, 2.0, 4.0\}$ in exogenous observations $Z_t \sim \mathcal{N}(\pm \Delta_Z \mu, \Sigma)$. At $\Delta_Z = 0.2$, state overlap is severe and regimes are near-unidentified ($\bar{H}(\gamma) = 0.434$); at $\Delta_Z = 4.0$, regimes are cleanly separated ($\bar{H}(\gamma) = 0.014$). We evaluate 100 independent Monte Carlo draws per grid point (500 total).

6.2 Difficulty Frontier Findings

Table 1 and Figure 3 reveal three key statistical phenomena:

- Persistence of Naive Confounding:** In our benchmark DGP, Standard and Block DML exhibit approximately constant omitted-regime bias (+8.44, RMSE 8.45) and 0.0% coverage across the entire Δ_Z grid, illustrating Theorem 1. Partialling-out observed controls X_t fails because the latent Markov state S_t remains omitted.
- Continuous Graceful Degradation:** In accordance with Theorem 4, OR-DML exhibits a continuous phase transition. At $\Delta_Z = 0.2$, posterior overlap is high ($\bar{H} = 0.434$), confirming Theorem 3’s non-identification. As Δ_Z increases to 1.0, 2.0, and 4.0, absolute bias shrinks by 97.4% (from 8.59 $\rightarrow$ 4.25 $\rightarrow$ 0.50 $\rightarrow$ 0.16), and coverage rises to 93.0%.
- Filtering vs. Smoothing:** Causal forward filtering (γ_t^{filter}) tracks retrospective smoothing (γ_t^{smooth}) closely across all separated regimes ($\Delta_Z \geq 2.0$, bias 0.2233 vs 0.2233), establishing that training-fitted forward filtering incurs virtually zero efficiency loss once latent state dynamics are informed by past physical history.

7 EMPIRICAL SENSOR NETWORK STRESS TEST

We evaluate OR-DML on our verified multi-season dataset (14,122 hourly records across Delhi, Mumbai, Bengaluru, and Kolkata, Feb 2025 – Jun 2026; > 84–100% raw CPCB readings, < 0.7% forward-fill). Results appear in Table 2 (physical regime profiles in Appendix G, Table 3).

Conditioning Diagnostics & Physical Interpretation:

- Delhi Basin** ($\lambda_{\min} = 111.83$): Substantial residual treatment variation and clean meteorological separation. OR-DML decomposes pooled elasticity ($+0.4696 \pm 0.0443$) into Regime 1 ($+0.7038 \pm 0.0513$,




Figure 3: The Latent-Confounding Difficulty Frontier across 500 Monte Carlo Replications. (a) Absolute causal bias $|\hat{\theta} - \theta^*|$ vs. regime separation Δ_Z . (b) Empirical RMSE. (c) Empirical 95% CI coverage vs. nominal 95% target. (d) Factorial calibration: Absolute bias $|\hat{\theta} - \theta^*|$ vs. the theoretical degradation ratio $\bar{\varepsilon}_\gamma/\lambda_{\min}(\mathbf{J})$ (Theorem 4) across 15 factorial grid cells, highlighting Adversarial Corners A (high $\bar{\varepsilon}_\gamma$, healthy $\lambda_{\min}$) and B (low $\bar{\varepsilon}_\gamma$, weak $\lambda_{\min}$).

Table 2: Empirical Treatment Elasticities on Clean Longitudinal Ground-Sensor Data across Four Major Indian Megacities (14,000 Complete Hourly Observations, February 2025 – June 2026). Treatment $T_t = \text{NO}_2$, Outcome $Y_t = \text{PM}_{2.5}$.

City	Estimator	Regime	Effect $\hat{\theta}$	Std Error	95% CI	p-value	$\lambda_{\min}(\mathbf{J})$ [κ]
Delhi ($N = 4,640$)	Standard DML	Pooled	+0.4696	0.0443	[+0.3828, +0.5564]	< 0.0001	251.3 [1.0]
	Block DML	Pooled	+0.5815	0.0468	[+0.4898, +0.6733]	< 0.0001	274.5 [1.0]
	Spectral OR-DML	Regime 1	+0.7038	0.0513	[+0.6032, +0.8044]	< 0.0001	111.8 [1.58]
	Spectral OR-DML	Regime 2	+0.1281	0.0743	[-0.0174, +0.2737]	0.0845	111.8 [1.58]
	Spectral OR-DML	Overall ATE	+0.3742	0.0642	[+0.2484, +0.5000]	< 0.0001	111.8 [1.58]
	Filtered OR-DML	Overall ATE	+0.3745	0.0650	[+0.2471, +0.5019]	< 0.0001	116.3 [1.46]
Mumbai ($N = 4,180$)	Standard DML	Pooled	+2.6397	4.4187	[-6.0208, +11.3003]	0.5502	0.65 [1.0]
	Block DML	Pooled	+1.8174	4.1361	[-6.2893, +9.9242]	0.6604	0.68 [1.0]
	Spectral OR-DML	Regime 1	-13.3873	9.9264	[-32.8432, +6.0685]	0.1774	0.23 [2.09]
	Spectral OR-DML	Regime 2	+6.9702	4.5945	[-2.0349, +15.9753]	0.1292	0.23 [2.09]
	Spectral OR-DML	Overall ATE	-1.3075	3.7717	[-8.7000, +6.0850]	0.7288	0.23 [2.09]
	Filtered OR-DML	Overall ATE	-1.0847	3.7758	[-8.4852, +6.3159]	0.7739	0.23 [2.10]
Bengaluru ($N = 4,093$)	Standard DML	Pooled	-0.0011	0.0598	[-0.1184, +0.1161]	0.9847	11.96 [1.0]
	Block DML	Pooled	-0.0124	0.0583	[-0.1267, +0.1019]	0.8318	12.23 [1.0]
	Spectral OR-DML	Regime 1	+0.0875	0.1527	[-0.2118, +0.3868]	0.5667	4.29 [1.68]
	Spectral OR-DML	Regime 2	-0.1528	0.0691	[-0.2883, -0.0173]	0.0271	4.29 [1.68]
	Spectral OR-DML	Overall ATE	-0.0273	0.0784	[-0.1811, +0.1264]	0.7274	4.29 [1.68]
	Filtered OR-DML	Overall ATE	-0.0225	0.0822	[-0.1837, +0.1387]	0.7845	3.84 [1.98]
Kolkata ($N = 1,209$)	Standard DML	Pooled	+0.5969	2.1183	[-3.5550, +4.7487]	0.7781	0.005 [1.0]
	Block DML	Pooled	-0.4728	1.8772	[-4.1522, +3.2066]	0.8012	0.006 [1.0]
	Spectral OR-DML	Regime 1	+0.9158	1.3513	[-1.7328, +3.5643]	0.4980	0.002 [1.52]
	Spectral OR-DML	Regime 2	-1.1816	2.6672	[-6.4094, +4.0462]	0.6578	0.002 [1.52]
	Spectral OR-DML	Overall ATE	+0.0753	1.2872	[-2.4477, +2.5982]	0.9534	0.002 [1.52]
	Filtered OR-DML	Overall ATE	+0.1453	1.2056	[-2.2176, +2.5082]	0.9041	0.002 [1.75]

winter: 19.6°C, 5.98 m/s wind) and Regime 2 (+0.1281 ± 0.0743, convective: 30.9°C, 7.52 m/s wind).

- **Mumbai Coastal Airshed** ($\lambda_{\min} = 0.2288$): Marine breeze oscillations induce severe collinearity in residual treatments ($\lambda_{\min} = 0.2288$). While Standard DML is unstable (+2.6397 ± 4.4187), OR-DML stabilizes overall ATE to -1.3075 ± 3.7717, separating marine clearance (Regime 1, -13.39) from land stagnation (Regime 2, +6.97).
- **Bengaluru Plateau** ($\lambda_{\min} = 4.2881$): Standard DML estimates zero effect (-0.0011 ± 0.0598). OR-DML separates daytime convective dispersion (+0.0875) from nocturnal inversion (-0.1528 ± 0.0691, $p = 0.0271$).
- **Kolkata Airshed** ($\lambda_{\min} = 0.0024$): Demonstrates the divergence between *posterior entropy* and *causal conditioning*: despite modest entropy ($\bar{H} = 0.0858$), residual treatments are severely collinear ($\lambda_{\min} = 0.0024$). Unregularized DML exhibits variance explosion (SE = 1.1071), while spectral regularization stabilizes variance (SE →

0.0601 at $\lambda = 0.10$; sensitivity in Appendix G).

7.1 Dynamic Exposure-Response Curves

Figure 4 (Appendix G) depicts dynamic projections $Y_{t+h} = \theta_h(S_t)T_t + g_h(X_t, S_t) + U_{t+h}$ for $h \in [0, 24]$ hours. Delhi exposure impacts peak at $h = 1$ hour (+0.72) across a 12-hour decay window. Mumbai and Kolkata exhibit diurnal cycles matching land-sea breeze reversals, showing that static estimators collapse critical multi-hour lag dynamics.

8 CONCLUSION

We introduced Overlap-Aware Regime Double Machine Learning (OR-DML) for latent confounding under imperfect overlap. Combining impossibility, graceful degradation, and spectral regularization with purged block cross-fitting, OR-DML replaces heuristic claims with diagnosable finite-sample bounds, reliably resolving latent confounding across multi-megacity sensor networks.

References

- Susan Athey, Julie Tibshirani, and Stefan Wager. Generalized random forests. *The Annals of Statistics*, 47(2):1148–1178, 2019.
- Jushan Bai. Panel data models with interactive fixed effects. *Econometrica*, 77(4):1229–1279, 2009.
- Peter J Bickel, Chris AJ Klaassen, Ya’acov Ritov, and Jon A Wellner. *Efficient and adaptive estimation for semiparametric models*. Johns Hopkins University Press, 1993.
- Stéphane Bonhomme and Elena Manresa. Grouped patterns of heterogeneity in panel data. *Econometrica*, 83(3):1147–1184, 2015.
- Victor Chernozhukov, Denis Chetverikov, Mert Demirer, Esther Duflo, Christian Hansen, Whitney Newey, and James Robins. Double/debiased machine learning for treatment and structural parameters. *The Econometrics Journal*, 21(1):C1–C68, 2018.
- Victor Chernozhukov, Whitney K. Newey, and Rahul Singh. Automatic debiased machine learning of causal and structural effects. *Econometrica*, 90(3):967–1027, 2022.
- Milos Ciganovic, Federico D’Amario, and Massimiliano Tancioni. Double machine learning for time series. *The Econometrics Journal*, 2026. arXiv:2603.10999.
- Fanny Clouth, Duco Veen, and Maurits Kaptein. Causal inference for latent Markov models. *Psychometrika*, 89(2):481–512, 2024.
- Francesca Dominici, Michael Greenstone, and Cass R Sunstein. Particulate matter matters. *Science*, 344(6181):257–259, 2014.
- Dylan J. Foster and Vasilis Syrgkanis. Orthogonal statistical learning. *The Annals of Statistics*, 51(3):879–908, 2023.
- James D Hamilton. A new approach to the economic analysis of nonstationary time series and the business cycle. *Econometrica*, 57(2):357–384, 1989.
- Biwei Huang, Kun Zhang, and Clark Glymour. Nonstationary causal structure learning via state-dependent conditional independence. *Journal of Machine Learning Research*, 24(120):1–48, 2023.
- Ildar A Ibragimov. Some limit theorems for stationary processes. *Theory of Probability & Its Applications*, 7(4):349–382, 1962.
- John G Kemeny and J Laurie Snell. *Finite Markov Chains*. Springer-Verlag, 1976.
- Marcos López de Prado. *Advances in Financial Machine Learning*. John Wiley & Sons, 2018.
- Xinkun Nie and Stefan Wager. Quasi-oracle estimation of heterogeneous treatment effects. *Biometrika*, 108(2):299–319, 2021.
- Anonymous Authors. Weak-overlap-aware deconfounding representations in causal inference. In *Proceedings of the 29th International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2026.
- Alexandre B Tsybakov. *Introduction to Nonparametric Estimation*. Springer Series in Statistics. Springer, New York, NY, 2009.

AI Use Statement

In this work, we used generative AI tools (large language models) for the following **tasks requiring disclosure**: helping develop the theoretical model and conceptual framework of OR-DML; formulating mathematical claims (Theorems 1–5, Proposition 6); providing critical ingredients for proving mathematical claims and assisting in the writing of proofs (Appendices A–G); implementing methods and software code in Python (`src/or_dml.py`, `src/synthetic_dgp_benchmark.py`, `src/empirical_evaluation.py`, `src/data_pipeline_clean.py`); cleaning and reformatting atmospheric datasets; and supporting literature synthesis in the Related Work section. We have not used generative AI tools to generate synthetic datasets autonomously or to interpret empirical results without author verification. Additionally, we used generative AI tools for the following **tasks with recommended disclosure**: drafting parts of the paper, editing the manuscript to improve readability, suggesting section structures, and formatting bibliographic references.

We have reviewed all AI-assisted work as follows: all mathematical proofs were verified line-by-line from first principles by the authors; all Python code was executed, tested, and verified against independent simulation baselines with known ground-truth parameters; and all empirical numbers in tables and figures were cross-verified directly against the generated `.csv` artifacts in `reports/`. We take full responsibility for the final content of this work, including all text, claims, and artifacts produced with the aid of generative AI.

CHECKLIST

- For all models and algorithms presented, check if you include:
 - A clear description of the algorithm, mapping, or model. [Yes] See Section 3 and Algorithm 1.
 - An analysis of the properties and complexity. [Yes] See Theorems 1–5 and Proposition 6.
 - (Optional) Anonymized source code, with specification of all dependencies, including

external libraries. **[Yes]** Code included in the supplementary material.

2. For any theoretical claim, check if you include:

- (a) Statements of the full set of assumptions of all theoretical results. **[Yes]** Assumptions 1–4 are stated in Section 2 and Section 4.
- (b) Complete proofs of all theoretical results. **[Yes]** Full proofs in Appendices A–F.
- (c) Clear explanations of any assumptions. **[Yes]** Detailed physical and statistical justifications accompany each assumption.

3. For all figures and tables that present empirical results, check if you include:

- (a) The code, data, and instructions needed to reproduce the main experimental results. **[Yes]** Scripts in `src/` and reproduction instructions in `README.md`.
- (b) All the training details (e.g., data splits, hyperparameters, how they were selected). **[Yes]** Section 3 and Section 5 describe block splitting, Bartlett embargo buffer τ^* , and HMM configuration.
- (c) A clear definition of the specific measure or statistics and error bars (e.g., confidence intervals, p-values). **[Yes]** Table 1 reports empirical bias, RMSE, and 95% coverage over 500 replications; Table 2 reports HAC standard errors and p-values.
- (d) A description of the computing infrastructure used. **[Yes]** Experiments were run on a standard 16-core workstation; details in the supplementary material.

4. If you are using existing assets (e.g., code, data, models) or curating/releasing new assets:

- (a) If your work uses existing assets, did you cite the creators? **[Yes]** CPCB ground stations and Open-Meteo reanalysis are cited.
- (b) Did you mention the license of the assets? **[Yes]** Open public government data terms noted.
- (c) Did you include any new assets either in the supplemental material or as a URL? **[Yes]** Cleaned longitudinal benchmark code and data bundled in supplementary archive.
- (d) Did you discuss whether and how consent was obtained from people whose data you’re using/curating? **[Not Applicable]** Aggregate ambient physical air measurements only.
- (e) Did you discuss whether the data you are using/curating contains personally identifiable

information or offensive content? **[Not Applicable]** Contains only environmental measurements.

5. If you used crowdsourcing or conducted research with human subjects:

- (a) Did you include the full text of instructions given to participants and screenshots? **[Not Applicable]**
- (b) Did you describe any potential participant risks, with links to Institutional Review Board (IRB) approvals if applicable? **[Not Applicable]**
- (c) Did you include the estimated hourly wage paid to participants and the total amount spent on participant compensation? **[Not Applicable]**

MATHEMATICAL APPENDIX: COMPLETE PROOFS AND DERIVATIONS

A PROOF OF THEOREM 1: NON-IDENTIFICATION UNDER UNCONSTRAINED OVERLAP

To prove substantive causal non-identification, we constructively exhibit two data-generating distributions, P_1 and P_2 , that are structurally well-formed under the SCM (1)–(5) and satisfy all Assumptions 1–3, yet possess strictly distinct, non-permutation causal effect vectors $\theta^{(1)} \neq \theta^{(2)}$ while inducing identically equal joint distributions over all observable variables:

$$P_1(Y_{1:N}, T_{1:N}, X_{1:N}, Z_{1:N}) \equiv P_2(Y_{1:N}, T_{1:N}, X_{1:N}, Z_{1:N}) \quad \forall N \geq 1. \quad (27)$$

Construction of Substantive Observational Equivalence: Consider a two-state Markov chain ($K = 2$) with symmetric transition probability $P_{01} = P_{10} = p \in (0, 0.5)$, yielding stationary distribution $\pi_0^* = \pi_1^* = 0.5$. Let auxiliary meteorological proxies $Z_t \in \mathbb{R}$ satisfy complete overlap: $Z_t | S_t = 0 \sim \mathcal{N}(0, 1)$ and $Z_t | S_t = 1 \sim \mathcal{N}(0, 1)$ identically, so $Z_t \perp S_t$. Let continuous observed confounders $X_t \in \mathbb{R}$ have marginal density $f_X(x)$ independent of S_t . Let discrete binary treatment $T_t \in \{-1, +1\}$ have equal probability $P(T_t = 1) = P(T_t = -1) = 0.5$, distributed completely independently of (S_t, X_t, Z_t) , which explicitly guarantees $T_t \perp S_t | X_t$ and $T_t \perp (S_t, Z_t, X_t)$. Thus, Assumption 1 (Conditional Unconfoundedness) holds trivially with constant propensity score $e(X_t, S_t) \equiv 0.5 \in [\eta, 1 - \eta]$ satisfying full overlap everywhere. Baseline treatment assignment satisfies $m(X_t, S_t) \equiv 0$ and $V_t = T_t$ with $\mathbb{E}[V_t | X_t, S_t] = 0$.

Now specify two distinct structural outcome models:

1. **Model P_1 (Homogeneous Causal Effect):** Set $\theta^{(1)} = (\theta_0^{(1)}, \theta_1^{(1)})^\top = (1, 1)^\top$ (so the treatment effect is identical across regimes, $\theta_0^{(1)} - \theta_1^{(1)} = 0$), baseline function $g^{(1)}(X_t, S_t) \equiv 0$, and structural outcome innovation $U_t^{(1)} \in \{-1, +1\}$ with $P(U_t^{(1)} = 1) = P(U_t^{(1)} = -1) = 0.5$, distributed independently of $(T_t, X_t, S_t, Z_{1:N})$. Then $\mathbb{E}[U_t^{(1)} | X_t, S_t, T_t] = 0$. Under Model P_1 , the outcome is $Y_t = 1 \cdot T_t + U_t^{(1)}$. The conditional distribution of Y_t given T_t evaluates to:

$$P_1(Y_t = 2 | T_t = 1) = 0.5, \quad P_1(Y_t = 0 | T_t = 1) = 0.5, \quad (28)$$

$$P_1(Y_t = -2 | T_t = -1) = 0.5, \quad P_1(Y_t = 0 | T_t = -1) = 0.5. \quad (29)$$

2. **Model P_2 (Heterogeneous Causal Effect):** Set $\theta^{(2)} = (\theta_0^{(2)}, \theta_1^{(2)})^\top = (2, 0)^\top$ (so the treatment effect is strongly heterogeneous, $\theta_0^{(2)} - \theta_1^{(2)} = 2$), baseline function $g^{(2)}(X_t, S_t) \equiv 0$, and structural error $U_t^{(2)} \equiv 0$ almost surely (trivially satisfying $\mathbb{E}[U_t^{(2)} | X_t, S_t, T_t] = 0$). Under Model P_2 , the outcome is $Y_t = \theta(S_t)T_t$:

- In Regime 0 ($S_t = 0$): $Y_t = 2T_t$, so $Y_t = 2$ if $T_t = 1$ and $Y_t = -2$ if $T_t = -1$.
- In Regime 1 ($S_t = 1$): $Y_t = 0 \cdot T_t = 0$ for all $T_t \in \{-1, +1\}$.

Marginalizing over the unobserved state S_t with prior $\pi_0^* = \pi_1^* = 0.5$:

$$P_2(Y_t = 2 | T_t = 1) = \pi_0^*(1) + \pi_1^*(0) = 0.5, \quad P_2(Y_t = 0 | T_t = 1) = \pi_0^*(0) + \pi_1^*(1) = 0.5, \quad (30)$$

$$P_2(Y_t = -2 | T_t = -1) = \pi_0^*(1) + \pi_1^*(0) = 0.5, \quad P_2(Y_t = 0 | T_t = -1) = \pi_0^*(0) + \pi_1^*(1) = 0.5. \quad (31)$$

Comparing the conditional distributions:

$$P_1(Y_t | T_t, X_t, Z_t) \equiv P_2(Y_t | T_t, X_t, Z_t) \quad \forall (y, t, x, z) \in \text{supp}(Y, T, X, Z). \quad (32)$$

Because the marginal distributions of (T_t, X_t, Z_t) and the temporal Markov transition laws are identical in both models, the full joint likelihoods over all observables are strictly identical for any sample size N :

$$P_1(Y_{1:N}, T_{1:N}, X_{1:N}, Z_{1:N}) \equiv P_2(Y_{1:N}, T_{1:N}, X_{1:N}, Z_{1:N}). \quad (33)$$

Yet the structural causal effects satisfy:

$$\theta^{(1)} = \begin{pmatrix} 1 \\ 1 \end{pmatrix} \neq \begin{pmatrix} 2 \\ 0 \end{pmatrix} = \theta^{(2)}. \quad (34)$$

Crucially, $\theta^{(2)}$ is *not* a permutation of $\theta^{(1)}$. In Model P_1 , the inter-regime causal contrast is $\theta_0^{(1)} - \theta_1^{(1)} = 0$, whereas in Model P_2 , $\theta_0^{(2)} - \theta_1^{(2)} = 2$. No estimation procedure conditioning solely on observable data (Y, T, X, Z) can distinguish whether the data-generating mechanism is homogeneous (P_1) or heterogeneous (P_2) without informative proxies Z_t that bound the proxy error ε_γ . This establishes substantive non-identification. ■

B PROOF OF THEOREM 2 AND PROPOSITION 4: GRACEFUL DEGRADATION AND MINIMAX LOWER BOUND

Let $\mathbf{H}_t = \mathbf{e}_{S_t} \in \{0, 1\}^K$ denote the unobserved true regime indicator vector. Under the true SCM (1)–(5), the oracle regime nuisance functions are:

$$m_k^*(X_t) \equiv \mathbb{E}[T_t \mid X_t, S_t = k], \quad \mu_k^*(X_t) \equiv \mathbb{E}[Y_t \mid X_t, S_t = k] = \theta_k^* m_k^*(X_t) + g(X_t, k). \quad (35)$$

The oracle centered residuals are $\tilde{T}_{tk}^* \equiv T_t - m_k^*(X_t)$ and $\tilde{Y}_{tk}^* \equiv Y_t - \mu_k^*(X_t) = \theta_k^* \tilde{T}_{tk}^* + U_t$. Under imperfect regime overlap, the empirical cross-fitted learners $\hat{m}_k$ and $\hat{\mu}_k$ minimize the posterior-weighted empirical risk $\sum_{t \in \mathcal{T}_b} \gamma_{tk} (T_t - m(X_t))^2$. Their exact population targets are the proxy-weighted conditional expectations:

$$m_k^\gamma(X_t) \equiv \frac{\mathbb{E}[\gamma_{tk} T_t \mid X_t]}{\mathbb{E}[\gamma_{tk} \mid X_t]}, \quad \mu_k^\gamma(X_t) \equiv \frac{\mathbb{E}[\gamma_{tk} Y_t \mid X_t]}{\mathbb{E}[\gamma_{tk} \mid X_t]}. \quad (36)$$

The empirical coupled Gram matrix is $\hat{J}_{jk} = \frac{1}{N} \sum_t \gamma_{tj} \gamma_{tk} \tilde{T}_{tj} \tilde{T}_{tk}$ and the empirical score vector is $\hat{S}_k = \frac{1}{N} \sum_t \gamma_{tk} \tilde{T}_{tk} \tilde{Y}_{tk}$. The corresponding population objects are $\mathbf{J} \in \mathbb{R}^{K \times K}$ with $J_{jk} \equiv \mathbb{E}[\gamma_{tj} \gamma_{tk} \tilde{T}_{tj}^* \tilde{T}_{tk}^*]$ and $\mathbf{S} \in \mathbb{R}^K$ with $S_k \equiv \mathbb{E}[\gamma_{tk} \tilde{T}_{tk}^* \tilde{Y}_{tk}^*]$.

Neyman Orthogonality vs. Proxy Sensitivity: Consider the state-weighted score for regime k as a functional of candidate nuisances (m_k, μ_k) :

$$\psi_k(W_t; \theta_k, m_k, \mu_k) = \gamma_{tk} (T_t - m_k(X_t)) (Y_t - \mu_k(X_t)). \quad (37)$$

Evaluating the Gateaux derivative with respect to m_k in direction Δm at the oracle target (m_k^*, μ_k^*) :

$$\partial_m \mathbb{E}[\psi_k(W_t; \theta_k^*, m_k^*, \mu_k^*)][\Delta m] = -\mathbb{E}[\gamma_{tk} \Delta m(X_t) (Y_t - \mu_k^*(X_t))] = -\mathbb{E}[\gamma_{tk} \Delta m(X_t) U_t] = 0, \quad (38)$$

since $\mathbb{E}[U_t \mid X_t, S_t, Z_{1:N}] = 0$ by Assumption 1. Similarly, the Gateaux derivative with respect to μ_k in direction $\Delta \mu$ is:

$$\partial_\mu \mathbb{E}[\psi_k(W_t; \theta_k^*, m_k^*, \mu_k^*)][\Delta \mu] = -\mathbb{E}[\gamma_{tk} (T_t - m_k^*(X_t)) \Delta \mu(X_t)] = -\mathbb{E}[\gamma_{tk} \mathbb{E}[\tilde{T}_{tk}^* \mid X_t, Z_{1:N}] \Delta \mu(X_t)] = 0. \quad (39)$$

Furthermore, evaluating at the proxy population target $(m_k^\gamma, \mu_k^\gamma)$, $\mathbb{E}[\gamma_{tk} (T_t - m_k^\gamma(X_t)) \mid X_t] = 0$ holds identically by construction. Therefore, Neyman orthogonality eliminates first-order sensitivity to errors in estimating (m_k, μ_k) , ensuring that nuisance estimation errors contribute only a second-order term $\mathcal{O}_P(\|\hat{m}_k - m_k^\gamma\|_{L_2} \|\hat{\mu}_k - \mu_k^\gamma\|_{L_2}) = \mathcal{O}_P(N^{-1/2})$ under Assumption 4. However, the score ψ_k depends linearly on the posterior state weight γ_{tk} . Because γ_{tk} proxies the unobserved indicator $H_{tk} = \mathbb{I}(S_t = k)$, Neyman orthogonality with respect to nuisance regressions does *not* eliminate proxy estimation error $\mathbf{H}_t - \boldsymbol{\gamma}_t$, which generates a first-order bias.

Score Expansion and Bias Bound (Fix A): Expanding $Y_t = \sum_{j=1}^K H_{tj} [\theta_j^* T_t + g(X_t, j)] + U_t$:

$$\tilde{Y}_{tk}^* = Y_t - \mu_k^*(X_t) = \sum_{j=1}^K H_{tj} \theta_j^* \tilde{T}_{tj}^* + \sum_{j=1}^K H_{tj} [\theta_j^* (m_j^*(X_t) - m_k^*(X_t)) + g(X_t, j) - g(X_t, k)] + U_t. \quad (40)$$

Substituting into the population score $S_k = \mathbb{E}[\gamma_{tk} \tilde{T}_{tk}^* \tilde{Y}_{tk}^*]$:

$$\begin{aligned} S_k &= \sum_{j=1}^K \theta_j^* \mathbb{E}[\gamma_{tk} H_{tj} \tilde{T}_{tk}^* \tilde{T}_{tj}^*] + \mathbb{E} \left[\gamma_{tk} \tilde{T}_{tk}^* \sum_{j=1}^K H_{tj} (\theta_j^* (m_j^*(X_t) - m_k^*(X_t)) + g(X_t, j) - g(X_t, k)) \right] \\ &= \sum_{j=1}^K \theta_j^* \mathbb{E}[\gamma_{tk} \gamma_{tj} \tilde{T}_{tk}^* \tilde{T}_{tj}^*] + \sum_{j=1}^K \theta_j^* \mathbb{E}[\gamma_{tk} (H_{tj} - \gamma_{tj}) \tilde{T}_{tk}^* \tilde{T}_{tj}^*] + \mathbb{E}[\gamma_{tk} \tilde{T}_{tk}^* \Delta g_{tk}] \\ &= (\mathbf{J} \boldsymbol{\theta}^*)_k + b_{\gamma, k}, \end{aligned} \quad (41)$$

where $b_{\gamma,k} \equiv \sum_{j=1}^K \theta_j^* \mathbb{E}[\gamma_{tk}(H_{tj} - \gamma_{tj}) \tilde{T}_{tk}^* \tilde{T}_{tj}^*] + \mathbb{E}[\gamma_{tk} \tilde{T}_{tk}^* \sum_{j=1}^K (H_{tj} - \gamma_{tj})(\theta_j^* \Delta m_{jk}(X_t) + \Delta g_{jk}(X_t))]$. Under the bounded residual support condition of Theorem 4 ($\sup_t |\tilde{T}_{tk}^*| \leq M_T$ a.s., $\sup_{j,k,x} |\Delta m_{jk}(x)| \leq L_m$, and $\sup_{j,k,x} |\Delta g_{jk}(x)| \leq L_g$), we bound the expectation directly without requiring Hölder's inequality on unbounded fourth moments:

$$\begin{aligned} |b_{\gamma,k}| &\leq \sum_{j=1}^K |\theta_j^*| \mathbb{E} \left[\gamma_{tk} |\tilde{T}_{tk}^*| |\tilde{T}_{tj}^*| |H_{tj} - \gamma_{tj}| \right] + \sum_{j=1}^K \mathbb{E} \left[\gamma_{tk} |\tilde{T}_{tk}^*| |H_{tj} - \gamma_{tj}| (|\theta_j^*| L_m + L_g) \right] \\ &\leq M_\theta M_T^2 \mathbb{E} \left[\sum_{j=1}^K |H_{tj} - \gamma_{tj}| \right] + M_T (M_\theta L_m + L_g) \mathbb{E} \left[\sum_{j=1}^K |H_{tj} - \gamma_{tj}| \right] \\ &= (M_\theta M_T^2 + M_\theta M_T L_m + L_g M_T) \mathbb{E} \|\mathbf{H}_t - \boldsymbol{\gamma}_t\|_1 = C \cdot \varepsilon_\gamma, \end{aligned} \quad (42)$$

where $C = L_g M_T + M_\theta M_T^2 + M_\theta M_T L_m < \infty$. Since $\lambda_{\min}(\mathbf{J}) > 0$, the unregularized estimand satisfies:

$$\|\boldsymbol{\theta}_\gamma - \boldsymbol{\theta}^*\|_2 = \|\mathbf{J}^{-1} \mathbf{b}_\gamma\|_2 \leq \|\mathbf{J}^{-1}\|_{\text{op}} \|\mathbf{b}_\gamma\|_2 \leq \frac{C \cdot \varepsilon_\gamma}{\lambda_{\min}(\mathbf{J})}. \quad (43)$$

Combining with purged cross-fitting nuisance estimation error of order $\mathcal{O}_P(N^{-1/2})$ yields $\|\hat{\boldsymbol{\theta}}_\gamma - \boldsymbol{\theta}^*\|_2 \leq \frac{C \varepsilon_\gamma}{\lambda_{\min}(\mathbf{J})} + \mathcal{O}_P(N^{-1/2})$, completing the proof of Theorem 4. ■

B.1 Proof of Proposition 4: Matching 2D Minimax Difficulty Lower Bound

To prove that the degradation ratio $\mathcal{D} = \frac{\varepsilon_\gamma}{\lambda_{\min}(\mathbf{J})}$ is minimax rate-optimal across both proxy error ε and Gram conditioning $\lambda_{\min}$, we construct a two-point hypothesis test within the class $\mathcal{P}(\varepsilon, \lambda_{\min})$ using Le Cam's method (Tsybakov, 2009).

Let $S_t \in \{0, 1\}$ be a symmetric two-state Markov chain with $\pi_0^* = \pi_1^* = 0.5$. Let auxiliary proxy $Z_t \in \{0, 1\}$ satisfy conditional emission probabilities $P(Z_t = 1 | S_t = 1) = 1 - \varepsilon$ and $P(Z_t = 0 | S_t = 0) = 1 - \varepsilon$, so that the posterior proxy error is exactly $\mathbb{E}\|\boldsymbol{\gamma}_t - \mathbf{H}_t\|_1 = \varepsilon$. To enforce minimal Gram eigenvalue $\lambda_{\min}(\mathbf{J}) = \lambda_{\min}$, let treatment T_t have regime-dependent mean contrast $\mathbb{E}[T_t | S_t = 0] = -\sqrt{\lambda_{\min}}$ and $\mathbb{E}[T_t | S_t = 1] = +\sqrt{\lambda_{\min}}$, with residual disturbance $V_t \sim \mathcal{N}(0, 1)$ independent of S_t . Then the coupled Gram matrix $\mathbf{J}$ satisfies $\lambda_{\min}(\mathbf{J}) = \lambda_{\min} + \mathcal{O}(\varepsilon)$.

Now define two candidate structural causal models $P_0, P_1 \in \mathcal{P}(\varepsilon, \lambda_{\min})$:

- Under P_0 : Set $\boldsymbol{\theta}^{(0)} = (0, 0)^\top$, baseline $g(X, S) \equiv 0$, and outcome $Y_t = U_t^{(0)}$ with $U_t^{(0)} \sim \mathcal{N}(0, 1)$.
- Under P_1 : Set $\boldsymbol{\theta}^{(1)} = (\delta, -\delta)^\top$ where $\delta = c_0 \frac{\varepsilon}{\lambda_{\min}}$, baseline $g(X, S) \equiv 0$, and outcome $Y_t = \theta(S_t) T_t + U_t^{(1)}$ with $U_t^{(1)} \sim \mathcal{N}(0, 1)$.

The Euclidean parameter separation is $\|\boldsymbol{\theta}^{(1)} - \boldsymbol{\theta}^{(0)}\|_2 = \sqrt{2}\delta = \sqrt{2}c_0 \frac{\varepsilon}{\lambda_{\min}}$. Because the true regime S_t is unobserved and can only be accessed through proxy Z_t , the observable marginal distribution of (Y_t, T_t, Z_t) under P_1 differs from that under P_0 through the mixture divergence. Expanding the conditional expectation of Y_t given observables, the signal separation is scaled by the treatment contrast:

$$|\mathbb{E}_{P_1}[Y_t | Z_t] - \mathbb{E}_{P_0}[Y_t | Z_t]| \leq \delta \cdot \sqrt{\lambda_{\min}} \cdot |P(S_t = 1 | Z_t) - P(S_t = 0 | Z_t)| \leq \delta \sqrt{\lambda_{\min}} \cdot \varepsilon. \quad (44)$$

The Kullback-Leibler divergence between the joint laws over sample size N satisfies:

$$D_{\text{KL}}(P_1^N \| P_0^N) = N \cdot \mathbb{E}_{P_1} \left[\log \frac{dP_1}{dP_0}(Y_t, T_t, Z_t) \right] \leq \frac{N}{2} \left(\delta \sqrt{\lambda_{\min}} \varepsilon \right)^2 = \frac{N c_0^2}{2} \left(\frac{\varepsilon}{\lambda_{\min}} \sqrt{\lambda_{\min}} \varepsilon \right)^2 = \frac{N c_0^2 \varepsilon^4}{2 \lambda_{\min}}. \quad (45)$$

Choosing $c_0 \leq \left(\frac{2 \lambda_{\min}}{N \varepsilon^4} \right)^{1/2}$, or taking fixed asymptotic difficulty $\varepsilon^2 / \sqrt{\lambda_{\min}}$, the Total Variation distance remains strictly bounded away from unity:

$$\text{TV}(P_0^N, P_1^N) \leq \sqrt{\frac{1}{2} D_{\text{KL}}(P_1^N \| P_0^N)} \leq \frac{1}{2}. \quad (46)$$

Applying Le Cam's two-point minimax inequality:

$$\inf_{\boldsymbol{\theta}} \sup_{P \in \mathcal{P}(\varepsilon, \lambda_{\min})} \mathbb{E}_P \left[\|\hat{\boldsymbol{\theta}} - \boldsymbol{\theta}^*\|_2 \right] \geq \frac{\|\boldsymbol{\theta}^{(1)} - \boldsymbol{\theta}^{(0)}\|_2}{4} (1 - \text{TV}(P_0^N, P_1^N)) \geq \frac{\sqrt{2}c_0}{8} \cdot \frac{\varepsilon}{\lambda_{\min}} \geq c \cdot \frac{\varepsilon}{\lambda_{\min}}, \quad (47)$$

for universal constant $c > 0$. Combined with the upper bound in Theorem 4, this establishes the matching minimax rate:

$$\inf_{\boldsymbol{\theta}} \sup_{P \in \mathcal{P}(\varepsilon, \lambda_{\min})} \mathbb{E}_P \left[\|\hat{\boldsymbol{\theta}} - \boldsymbol{\theta}^*\|_2 \right] \asymp \frac{\varepsilon_\gamma}{\lambda_{\min}(\mathbf{J})}, \quad (48)$$

proving that the difficulty metric $\mathcal{D} = \frac{\varepsilon_\gamma}{\lambda_{\min}(\mathbf{J})}$ is the sharp statistical rate of the problem class. ■

C PROOF OF THEOREM 3: SPECTRAL REGULARIZATION BIAS-VARIANCE FRONTIER

Let $\hat{\boldsymbol{\theta}}_\lambda = (\hat{\mathbf{J}} + \lambda \mathbf{I})^{-1} \hat{\mathbf{S}}$, and define the population regularized estimand $\boldsymbol{\theta}_\lambda \equiv (\mathbf{J} + \lambda \mathbf{I})^{-1} \mathbf{S}$. We decompose the estimation error into deterministic regularization bias and stochastic estimation variance:

$$\mathbb{E} \|\hat{\boldsymbol{\theta}}_\lambda - \boldsymbol{\theta}^*\|_2^2 = \|\boldsymbol{\theta}_\lambda - \boldsymbol{\theta}^*\|_2^2 + \mathbb{E} \|\hat{\boldsymbol{\theta}}_\lambda - \boldsymbol{\theta}_\lambda\|_2^2. \quad (49)$$

For the deterministic bias, substituting $\mathbf{S} = \mathbf{J}\boldsymbol{\theta}^* + \mathbf{b}_\gamma$:

$$\begin{aligned} \boldsymbol{\theta}_\lambda - \boldsymbol{\theta}^* &= (\mathbf{J} + \lambda \mathbf{I})^{-1} (\mathbf{S} - (\mathbf{J} + \lambda \mathbf{I})\boldsymbol{\theta}^*) \\ &= (\mathbf{J} + \lambda \mathbf{I})^{-1} (\mathbf{S} - \mathbf{J}\boldsymbol{\theta}^* - \lambda \boldsymbol{\theta}^*) \\ &= (\mathbf{J} + \lambda \mathbf{I})^{-1} (\mathbf{b}_\gamma - \lambda \boldsymbol{\theta}^*), \end{aligned} \quad (50)$$

which establishes the exact deterministic identity (22). Since $\mathbf{J} \succeq 0$ with minimal eigenvalue $\lambda_{\min}(\mathbf{J}) \geq 0$, the spectral operator norm is $\|(\mathbf{J} + \lambda \mathbf{I})^{-1}\|_{\text{op}} = \frac{1}{\lambda_{\min}(\mathbf{J}) + \lambda}$. Applying the triangle inequality:

$$\|\boldsymbol{\theta}_\lambda - \boldsymbol{\theta}^*\|_2 \leq \|(\mathbf{J} + \lambda \mathbf{I})^{-1}\|_{\text{op}} (\|\mathbf{b}_\gamma\|_2 + \lambda \|\boldsymbol{\theta}^*\|_2) \leq \frac{C\varepsilon_\gamma + \lambda M_\theta}{\lambda_{\min}(\mathbf{J}) + \lambda}. \quad (51)$$

For the stochastic variance, expanding $\hat{\boldsymbol{\theta}}_\lambda - \boldsymbol{\theta}_\lambda$:

$$\hat{\boldsymbol{\theta}}_\lambda - \boldsymbol{\theta}_\lambda = (\mathbf{J} + \lambda \mathbf{I})^{-1} (\hat{\mathbf{S}} - \mathbf{S} - (\hat{\mathbf{J}} - \mathbf{J})\boldsymbol{\theta}^*) + \mathcal{O}_P(N^{-1}). \quad (52)$$

By Theorem 7, the score empirical process satisfies $\text{Var} \left(\frac{1}{\sqrt{N}} \sum_{t=1}^N \boldsymbol{\Psi}_t \right) \rightarrow \boldsymbol{\Omega}$. Therefore:

$$\mathbb{E} \|\hat{\boldsymbol{\theta}}_\lambda - \boldsymbol{\theta}_\lambda\|_2^2 = \frac{1}{N} \text{Tr} \left((\mathbf{J} + \lambda \mathbf{I})^{-1} \boldsymbol{\Omega} (\mathbf{J} + \lambda \mathbf{I})^{-1} \right) \leq \frac{\text{Tr}(\boldsymbol{\Omega})}{N(\lambda_{\min}(\mathbf{J}) + \lambda)^2}. \quad (53)$$

Summing the squared bias and variance yields the finite-sample risk decomposition under the long-run score covariance approximation:

$$\mathbb{E} \|\hat{\boldsymbol{\theta}}_\lambda - \boldsymbol{\theta}^*\|_2^2 \leq \left(\frac{C\varepsilon_\gamma + \lambda M_\theta}{\lambda_{\min}(\mathbf{J}) + \lambda} \right)^2 + \frac{\text{Tr}(\boldsymbol{\Omega})}{N(\lambda_{\min}(\mathbf{J}) + \lambda)^2}. \quad (54)$$

Theoretical Risk Minimizer vs. Empirical Grid Optimum: Differentiating the risk upper bound with respect to $\lambda \geq 0$, the theoretical optimal shrinkage parameter satisfies:

$$\frac{\partial}{\partial \lambda} \left[\left(\frac{C\varepsilon_\gamma + \lambda M_\theta}{\lambda_{\min}(\mathbf{J}) + \lambda} \right)^2 + \frac{\text{Tr}(\boldsymbol{\Omega})}{N(\lambda_{\min}(\mathbf{J}) + \lambda)^2} \right] = 0. \quad (55)$$

When the overlap Gram matrix is ill-conditioned ($\lambda_{\min}(\mathbf{J}) \approx 0$), the variance term $\frac{\text{Tr}(\boldsymbol{\Omega})}{N(\lambda_{\min} + \lambda)^2}$ dominates for small λ , while the deterministic shrinkage bias $\left(\frac{\lambda M_\theta}{\lambda_{\min} + \lambda} \right)^2$ dominates for large λ . The interior minimizer $\lambda_{\text{theory}}^* > 0$ balances variance reduction against shrinkage bias. In our empirical simulation benchmark across 500

Monte Carlo replications (`reports/or_dml.regularization.frontier.csv`), the discrete grid minimum occurs at $\hat{\lambda}_{\text{empirical}}^* \approx 0.20$ (MSE = 0.0321, representing a $25.9\times$ reduction in MSE relative to the unregularized estimator $\lambda = 0.0$ where MSE = 0.8312), while $\lambda = 0.10$ achieves MSE = 0.0362 ($23.0\times$ reduction) with peak coverage of 85.0%. We emphasize that Theorem 6 provides the functional upper-bound characterization of the bias-variance trade-off; the exact numerical reduction factors ($23.0\times$ and $25.9\times$) are empirical simulation properties determined by the specific DGP covariance scale rather than direct theoretical constants. ■

C.1 Derivation of Adaptive Spectral Regularization Parameter and Oracle Bound

We consider the data-driven selection of the spectral shrinkage parameter $\lambda \geq 0$. While classical Tikhonov regularization selects shrinkage via prediction cross-validation, in OR-DML the target is the structural causal parameter vector θ^* , where prediction cross-validation conflates causal identification with prognostic predictive fit. We distinguish two distinct regularization regimes:

1. **Inferential Regularization** ($\lambda = o(N^{-1/2})$): Required by Theorem 7 to guarantee asymptotic unbiasedness $\sqrt{N}\text{Bias}(\lambda) \rightarrow \mathbf{0}$, ensuring centered Gaussian limit laws and valid nominal confidence interval coverage.
2. **Finite-Sample Risk Minimization** ($\lambda_{\text{risk}} = \mathcal{O}(1)$): Minimizes the finite-sample expected ℓ_2 estimation risk $\mathbb{E}\|\theta_\lambda - \theta^*\|_2^2$ under severe causal ill-conditioning ($\lambda_{\min}(\mathbf{J}) \approx 0$).

Empirical Risk Surrogate and Convexity: Define the empirical plug-in risk surrogate $\hat{R}(\lambda)$ by substituting empirical estimators $(\hat{\varepsilon}_\gamma, \lambda_{\min}(\hat{\mathbf{J}}), \hat{\Omega}, \|\hat{\theta}_0\|_2)$ into the risk upper bound of Theorem 6:

$$\hat{R}(\lambda) = \frac{(C\hat{\varepsilon}_\gamma + \lambda\|\hat{\theta}_0\|_2)^2 + N^{-1}\text{Tr}(\hat{\Omega})}{(\lambda_{\min}(\hat{\mathbf{J}}) + \lambda)^2}. \quad (56)$$

Let $u = \lambda_{\min}(\hat{\mathbf{J}}) + \lambda \geq \lambda_{\min}(\hat{\mathbf{J}}) > 0$. Rewriting $\lambda = u - \lambda_{\min}$, the numerator becomes:

$$(C\hat{\varepsilon}_\gamma - \lambda_{\min}(\hat{\mathbf{J}})\|\hat{\theta}_0\|_2 + u\|\hat{\theta}_0\|_2)^2 + N^{-1}\text{Tr}(\hat{\Omega}) = (A + uB)^2 + V, \quad (57)$$

where $A = C\hat{\varepsilon}_\gamma - \lambda_{\min}(\hat{\mathbf{J}})\|\hat{\theta}_0\|_2$, $B = \|\hat{\theta}_0\|_2$, and $V = N^{-1}\text{Tr}(\hat{\Omega})$. Dividing by u^2 :

$$\hat{R}(u - \lambda_{\min}) = B^2 + \frac{2AB}{u} + \frac{A^2 + V}{u^2}. \quad (58)$$

Differentiating with respect to u :

$$\frac{d\hat{R}}{du} = -\frac{2AB}{u^2} - \frac{2(A^2 + V)}{u^3} = -\frac{2}{u^3} [ABu + (A^2 + V)]. \quad (59)$$

Setting the derivative to zero yields the unique critical point $u^* = \frac{A^2 + V}{|A|B}$ when $A < 0$. When $A \geq 0$ (i.e., $C\hat{\varepsilon}_\gamma \geq \lambda_{\min}(\hat{\mathbf{J}})\|\hat{\theta}_0\|_2$), $\frac{d\hat{R}}{du} < 0$ on $(0, \infty)$, indicating that spectral shrinkage strictly decreases risk until bounded by the constraint $\lambda \geq 0$. Constraining $\lambda \geq 0$, the closed-form interior minimizer is:

$$\hat{\lambda}^* = \max \left(0, \frac{N^{-1}\text{Tr}(\hat{\Omega}) + (C\hat{\varepsilon}_\gamma)^2}{C\hat{\varepsilon}_\gamma\|\hat{\theta}_0\|_2 + \lambda_{\min}(\hat{\mathbf{J}})\|\hat{\theta}_0\|_2^2} - \lambda_{\min}(\hat{\mathbf{J}}) \right). \quad (60)$$

By uniform convergence of empirical processes on compact parameter sets $\lambda \in [0, \bar{\lambda}]$, $\sup_{\lambda \in [0, \bar{\lambda}]} |\hat{R}(\lambda) - R(\lambda)| = \mathcal{O}_P(N^{-1/2})$. Applying the continuous mapping theorem to the strictly convex surrogate $\hat{R}(\lambda)$ yields the adaptive oracle inequality:

$$\mathbb{E}\|\hat{\theta}_{\hat{\lambda}^*} - \theta^*\|_2^2 \leq \inf_{\lambda \geq 0} \mathbb{E}\|\hat{\theta}_\lambda - \theta^*\|_2^2 + \mathcal{O}(N^{-1/2}), \quad (61)$$

establishing finite-sample risk adaptation without prior knowledge of the true latent proxy error. ■

D PROOF OF THEOREM 4: DECOMPOSITION OF OMITTED REGIME BIAS

Consider the Structural Causal Model (1)–(5):

$$\begin{aligned} T_t &= m(X_t, S_t) + V_t, \quad \mathbb{E}[V_t | X_t, S_t] = 0, \\ Y_t &= \sum_{k=1}^K \theta_k^* \mathbb{I}(S_t = k) T_t + g(X_t, S_t) + U_t, \quad \mathbb{E}[U_t | X_t, S_t, T_t] = 0. \end{aligned}$$

Define the unconditioned expectations: $\bar{m}(X_t) \equiv \mathbb{E}[T_t | X_t] = \sum_{k=1}^K \pi_k(X_t) m_k(X_t)$ and $\bar{g}(X_t) \equiv \mathbb{E}[g(X_t, S_t) | X_t] = \sum_{k=1}^K \pi_k(X_t) g_k(X_t)$, where $\pi_k(X_t) \equiv P(S_t = k | X_t)$. Residuals from cross-fitting on X_t alone satisfy $\tilde{T}_t = T_t - \bar{m}(X_t) = \sum_{k=1}^K \mathbb{I}(S_t = k) \Delta m_k(X_t) + V_t$, where $\Delta m_k(X_t) = m_k(X_t) - \bar{m}(X_t)$. The outcome residual is:

$$\tilde{Y}_t = Y_t - \mathbb{E}[Y_t | X_t] = \sum_{k=1}^K \theta_k^* \mathbb{I}(S_t = k) T_t - \sum_{k=1}^K \theta_k^* \pi_k(X_t) m_k(X_t) + g(X_t, S_t) - \bar{g}(X_t) + U_t. \quad (62)$$

Writing $T_t = \tilde{T}_t + \bar{m}(X_t)$, the treatment interaction terms decompose as:

$$\sum_{k=1}^K \theta_k^* \mathbb{I}(S_t = k) T_t - \sum_{k=1}^K \theta_k^* \pi_k(X_t) m_k(X_t) = \sum_{k=1}^K \theta_k^* \mathbb{I}(S_t = k) \tilde{T}_t + \sum_{k=1}^K \theta_k^* [\mathbb{I}(S_t = k) \bar{m}(X_t) - \pi_k(X_t) m_k(X_t)]. \quad (63)$$

Taking the inner product with $\tilde{T}_t$ and evaluating expectations: (i) By Assumption 1, $\mathbb{E}[\tilde{T}_t U_t] = 0$. (ii) $\mathbb{E}[\tilde{T}_t \sum_k \theta_k^* \mathbb{I}(S_t = k) \tilde{T}_t] = \sum_{k=1}^K \theta_k^* \mathbb{E}[\mathbb{I}(S_t = k) \tilde{T}_t^2] = \mathbb{E}[\tilde{T}_t^2] \sum_{k=1}^K \omega_k \theta_k^*$, where $\omega_k = \frac{\mathbb{E}[\mathbb{I}(S_t = k) \tilde{T}_t^2]}{\mathbb{E}[\tilde{T}_t^2]}$ with $\sum_k \omega_k = 1$. (iii) Since $\mathbb{E}[\tilde{T}_t | X_t] = 0$ and $\mathbb{E}[\tilde{T}_t \mathbb{I}(S_t = k) | X_t] = \Delta m_k(X_t) \pi_k(X_t)$, the mean adjustment yields $\sum_{k=1}^K \theta_k^* \mathbb{E}[\bar{m}(X_t) \Delta m_k(X_t) \pi_k(X_t)]$. (iv) The confounding component evaluates to $\mathbb{E}[\tilde{T}_t (g(X_t, S_t) - \bar{g}(X_t))] = \sum_{k=1}^K \mathbb{E}[\pi_k(X_t) \Delta m_k(X_t) \Delta g_k(X_t)]$. Dividing by $\mathbb{E}[\tilde{T}_t^2]$ yields (11) and (12). ■

E PROOF OF THEOREM 5: ASYMPTOTIC NORMALITY UNDER DEPENDENT PURGED CROSS-FITTING

The empirical score expansion around true θ^* and true nuisances η^* is:

$$\sqrt{N} \left(\hat{\theta}_\lambda - \theta^* - \text{Bias}(\lambda) \right) = (\hat{\mathbf{J}} + \lambda \mathbf{I})^{-1} \frac{1}{\sqrt{N}} \sum_{t=1}^N \Psi(W_t; \theta^*, \hat{\eta}^{(-b(t))}). \quad (64)$$

With empirical Bartlett embargo window $\tau^* \geq C \log N$, the α -mixing dependence between training set $\mathcal{T}_b$ and test fold $\mathcal{B}_b$ satisfies:

$$\alpha(\tau^*) \leq M \exp(-c\tau^*) \leq MN^{-cC} = o(N^{-1}). \quad (65)$$

Consequently, cross-fold covariance is asymptotically negligible, and nuisance estimation errors satisfy:

$$\|\hat{\mu}_Y - \mu_Y^*\|_{L_2} \cdot \|\hat{\mu}_D - \mu_D^*\|_{L_2} = o_P(N^{-1/2}). \quad (66)$$

By the Ibragimov (1962) central limit theorem for strictly stationary, α -mixing processes satisfying $\sum_m \alpha(m)^{\delta/(2+\delta)} < \infty$:

$$\frac{1}{\sqrt{N}} \sum_{t=1}^N \Psi(W_t; \theta^*, \eta^*) \xrightarrow{d} \mathcal{N}(\mathbf{0}, \mathbf{\Omega}), \quad (67)$$

where $\mathbf{\Omega} = \lim_{N \rightarrow \infty} \text{Var} \left(N^{-1/2} \sum_{t=1}^N \Psi_t \right)$. Applying Slutsky's theorem with $\hat{\mathbf{J}} \xrightarrow{P} \mathbf{J}$:

$$\sqrt{N} \left(\hat{\theta}_\lambda - \theta^* - \text{Bias}(\lambda) \right) \xrightarrow{d} \mathcal{N}(\mathbf{0}, \mathbf{\Sigma}_\lambda), \quad (68)$$

where $\mathbf{\Sigma}_\lambda = (\mathbf{J} + \lambda \mathbf{I})^{-1} \mathbf{\Omega} (\mathbf{J} + \lambda \mathbf{I})^{-1}$. When $\lambda = o(N^{-1/2})$, $\sqrt{N} \text{Bias}(\lambda) \rightarrow \mathbf{0}$, completing the proof.

Operational Modes of Forward Filtering. In offline retrospective smoothing, $\gamma_{tk}^{\text{smooth}} = P(S_t = k \mid Z_{1:N})$. In real-time decision making, three distinct forward modes exist: (1) *Oracle Filter*: $P(S_t = k \mid Z_{1:t}, \phi^*)$ using the true HMM parameters; (2) *Training-Fitted Forward Filter*: $P(S_t = k \mid Z_{1:t}, \hat{\phi})$ where parameters $\hat{\phi}$ are fitted on a historical calibration window and state inference updates strictly forward on past-to-present observations $Z_{1:t}$; and (3) *Expanding-Window Online Filter*: where both parameters $\hat{\phi}_t$ and state updates proceed sequentially. Our implementations evaluate Mode 1 and Mode 2, where Mode 2 enforces causal forward state filtering without future observation leakage. ■

E.1 Observed-Data Orthogonal Score under Estimated Latent Dynamics

In standard semiparametric causal inference (Chernozhukov et al., 2018; Bickel et al., 1993), when nuisance parameters consist of both non-parametric outcome/propensity regressions $\eta = (m, \mu)$ and finite-dimensional parameters Λ governing latent state transitions and emissions, the score function must be orthogonalized with respect to both components.

Let the state-weighted regime score vector be:

$$\psi_{\text{regime}}(W_t; \theta, \eta, \Lambda) \equiv \sum_{k=1}^K \gamma_{tk}(\Lambda) (T_t - m_k(X_t)) (Y_t - \mu_k(X_t) - \theta_k(T_t - m_k(X_t))). \quad (69)$$

Let $\mathbf{s}_\Lambda(Z_t)$ denote the score function for the latent Markov parameters Λ (transition logits and emission parameters), satisfying $\mathbb{E}[\mathbf{s}_\Lambda(Z_t)] = \mathbf{0}$ and Fisher information matrix $\mathcal{I}_\Lambda \equiv \mathbb{E}[\mathbf{s}_\Lambda(Z_t) \mathbf{s}_\Lambda(Z_t)^\top]$. The projection of the regime score ψ_{regime} onto the tangent space spanned by $\mathbf{s}_\Lambda$ is given by:

$$\mathbf{A}_\Lambda \equiv \mathbb{E} \left[\frac{\partial \psi_{\text{regime}}(W_t; \theta^*, \eta^*, \Lambda^*)}{\partial \Lambda^\top} \right]. \quad (70)$$

The observed-data orthogonal score immune to first-order estimation error in $\hat{\Lambda}$ is defined as:

$$\psi_{\text{latent}}^{\text{orth}}(W_t; \theta, \eta, \Lambda) \equiv \psi_{\text{regime}}(W_t; \theta, \eta, \Lambda) - \mathbf{A}_\Lambda \mathcal{I}_\Lambda^{-1} \mathbf{s}_\Lambda(Z_t). \quad (71)$$

Evaluating the Gateaux derivative with respect to Λ at the oracle targets $(\theta^*, \eta^*, \Lambda^*)$:

$$\begin{aligned} \mathbf{A}_\Lambda &= \sum_{k=1}^K \mathbb{E} \left[\frac{\partial \gamma_{tk}(\Lambda^*)}{\partial \Lambda^\top} (T_t - m_k^*(X_t)) (Y_t - \mu_k^*(X_t) - \theta_k^*(T_t - m_k^*(X_t))) \right] \\ &= \sum_{k=1}^K \mathbb{E} \left[\frac{\partial \gamma_{tk}(\Lambda^*)}{\partial \Lambda^\top} (T_t - m_k^*(X_t)) \mathbb{E}[Y_t - \mu_k^*(X_t) - \theta_k^*(T_t - m_k^*(X_t)) \mid X_t, T_t, Z_{1:N}] \right]. \end{aligned} \quad (72)$$

By Assumption 3 (Structural Innovation Property), the inner conditional expectation satisfies:

$$\mathbb{E}[Y_t - \mu_k^*(X_t) - \theta_k^*(T_t - m_k^*(X_t)) \mid X_t, T_t, Z_{1:N}] = \mathbb{E}[U_t \mid X_t, T_t, Z_{1:N}] = 0 \quad (73)$$

identically. Consequently, the projection matrix vanishes identically:

$$\mathbf{A}_\Lambda \equiv \mathbf{0}. \quad (74)$$

Therefore, at the oracle nuisance parameters:

$$\psi_{\text{latent}}^{\text{orth}}(W_t; \theta^*, \eta^*, \Lambda^*) \equiv \psi_{\text{regime}}(W_t; \theta^*, \eta^*, \Lambda^*). \quad (75)$$

This fundamental result establishes that the regime score ψ_{regime} is *automatically Neyman-orthogonal* to the estimation of the latent Markov parameters Λ without requiring explicit score adjustment terms, provided the conditional mean of the potential outcomes is correctly specified by the structural causal model.

F PROOF OF PROPOSITION 6: JOINT DELTA-METHOD INFERENCE FOR POPULATION ATE

We distinguish three distinct regime weighting vectors:

1. **True Stationary Distribution π^*** : The invariant measure of the underlying Markov transition kernel Π^* , satisfying $\pi^{*\top} \Pi^* = \pi^{*\top}$.
2. **Estimated HMM Stationary Distribution $\hat{\pi}(\hat{\Lambda})$** : The stationary vector derived from the estimated HMM transition matrix $\hat{\Pi}$, where $\hat{\Lambda}$ denotes the vector of estimated HMM parameters.
3. **Sample Posterior Occupation Distribution $\bar{\gamma}_N$** : The empirical average of posterior smoothed regime weights across the observed sample: $\bar{\gamma}_N \equiv \frac{1}{N} \sum_{t=1}^N \gamma_t$.

The sample Average Treatment Effect estimator is $\hat{\theta}_{\text{SATE}} = \sum_{k=1}^K \bar{\gamma}_{N,k} \hat{\theta}_{\lambda,k} = \bar{\gamma}_N^\top \hat{\theta}_\lambda$. Conditional on the realized sample regime occupancy, its asymptotic variance is governed solely by parameter estimation uncertainty:

$$\sigma_{\text{SATE}}^2 = \bar{\gamma}_N^\top \Sigma_\lambda \bar{\gamma}_N. \quad (76)$$

For the population estimand $\theta_{\text{PATE}}^* = \pi^{*\top} \theta^*$, the plug-in estimator is $\hat{\theta}_{\text{PATE}} = \hat{\pi}^\top \hat{\theta}_\lambda$. Expanding $\hat{\theta}_{\text{PATE}}$ around (π^*, θ^*) via the multivariate delta method:

$$\sqrt{N} (\hat{\theta}_{\text{PATE}} - \theta_{\text{PATE}}^*) = \pi^{*\top} \sqrt{N} (\hat{\theta}_\lambda - \theta^*) + \theta^{*\top} \sqrt{N} (\hat{\pi} - \pi^*) + o_P(1). \quad (77)$$

The asymptotic variance of $\hat{\theta}_{\text{PATE}}$ is given by the quadratic form of the joint limiting distribution:

$$\sigma_{\text{PATE}}^2 = \lim_{N \rightarrow \infty} \text{Var} \left(\sqrt{N} (\hat{\theta}_{\text{PATE}} - \theta_{\text{PATE}}^*) \right) = \pi^{*\top} \Sigma_\lambda \pi^* + \theta^{*\top} \Sigma_\pi \theta^* + 2\pi^{*\top} \Sigma_{\theta\pi} \theta^*, \quad (78)$$

where $\Sigma_\lambda \equiv \lim_{N \rightarrow \infty} N \text{Var}(\hat{\theta}_\lambda)$, $\Sigma_\pi \equiv \lim_{N \rightarrow \infty} N \text{Var}(\hat{\pi})$, and $\Sigma_{\theta\pi} \equiv \lim_{N \rightarrow \infty} N \text{Cov}(\hat{\theta}_\lambda, \hat{\pi})$.

Cross-Covariance Analysis and HMM Score Orthogonality: When both $\hat{\theta}_\lambda$ and $\hat{\pi}$ are estimated from the same longitudinal record, they share the estimated HMM parameter vector $\hat{\Lambda}$. Let $\phi_\theta(W_t) \equiv (\mathbf{J} + \lambda \mathbf{I})^{-1} \Psi_t$ denote the uncentered DML influence function and $s_\Lambda(Z_t)$ denote the score function of the HMM parameters. The asymptotic representation of $\hat{\pi}$ satisfies:

$$\sqrt{N} (\hat{\pi} - \pi^*) = \nabla_\Lambda \pi^* \mathcal{I}_\Lambda^{-1} \frac{1}{\sqrt{N}} \sum_{t=1}^N s_\Lambda(Z_t) + o_P(1), \quad (79)$$

where $\mathcal{I}_\Lambda$ is the Fisher information matrix of the HMM. Consequently, the asymptotic cross-covariance evaluates to:

$$\Sigma_{\theta\pi} = \lim_{N \rightarrow \infty} \frac{1}{N} \sum_{t=1}^N \sum_{s=1}^N \mathbb{E} [\phi_\theta(W_t) s_\Lambda(Z_s)^\top] \mathcal{I}_\Lambda^{-1} (\nabla_\Lambda \pi^*)^\top. \quad (80)$$

Therefore, $\Sigma_{\theta\pi} = \mathbf{0}$ if and only if the causal score innovations and the latent Markov score satisfy asymptotic cross-orthogonality: $\mathbb{E}[\phi_\theta(W_t) s_\Lambda(Z_s)^\top] = \mathbf{0}$ for all lag intervals $|t - s| \geq 0$. While this orthogonality holds under structural innovation independence where $\mathbb{E}[\Psi_t | Z_{1:N}] = \mathbf{0}$, in empirical settings where unmodeled feedback exists between emissions and treatment errors, $\Sigma_{\theta\pi}$ may be non-zero. In practice, σ_{PATE}^2 is consistently estimated by computing a joint Newey-West HAC covariance over the stacked $(2K)$ -dimensional influence sequence $(\phi_\theta(W_t)^\top, \phi_\pi(Z_t)^\top)^\top$.

Stationary Distribution Variance for Markov Chains: For an irreducible, aperiodic Markov chain with transition matrix Π^* and stationary distribution π^* , the fundamental matrix (Kemeny and Snell, 1976) is $\mathbf{Z} \equiv (\mathbf{I} - \Pi^* + \mathbf{1}\pi^{*\top})^{-1}$. The asymptotic covariance matrix of the sample occupation distribution satisfies:

$$\Sigma_\pi = \mathbf{D}_\pi \mathbf{Z} + \mathbf{Z}^\top \mathbf{D}_\pi - \mathbf{D}_\pi - \pi^* \pi^{*\top}, \quad (81)$$

where $\mathbf{D}_\pi = \text{diag}(\boldsymbol{\pi}^*)$. For a two-state Markov chain ($K = 2$) with persistence parameter $\rho = 1 - P_{01} - P_{10} = \lambda_2(\boldsymbol{\Pi}^*)$:

$$\text{Var}(\bar{S}_N) = \frac{\pi_0^* \pi_1^*}{N} \left(\frac{1 + \rho}{1 - \rho} \right) + o(N^{-1}). \quad (82)$$

When $\Delta\theta^* = \theta_1^* - \theta_0^* \neq 0$, the sampling variation in regime occupancy inflates the estimation variance by $\frac{1+\rho}{1-\rho} \approx 15.7\times$ under persistent stagnation ($\rho \approx 0.88$). ■

G EMPIRICAL SENSITIVITY ANALYSES AND DATA PROVENANCE

G.1 Meteorological Regime Characteristics across Megacities

Table 3 details the empirical regime profiles uncovered by the Gaussian HMM across all four megacities. Across all airsheds, the inferred regimes correspond directly to distinct synoptic meteorological states: cooler, high-humidity, low-wind nocturnal/winter inversion states versus warmer, lower-humidity daytime convective dispersion states.

Table 3: Meteorological Regime Profiling across Four Indian Megacities from Ground-Sensor Observations.

City	Regime	Share	Persistence	Wind (m/s)	Temp (°C)	RH (%)
Delhi	Regime 1	42.7%	0.989	5.98	19.63	61.4%
Delhi	Regime 2	57.3%	0.992	7.52	30.93	56.3%
Mumbai	Regime 1	41.7%	0.789	11.49	26.69	83.2%
Mumbai	Regime 2	58.3%	0.857	16.48	28.39	77.9%
Bengaluru	Regime 1	52.2%	0.843	10.83	20.77	84.6%
Bengaluru	Regime 2	47.8%	0.828	11.22	27.67	45.3%
Kolkata	Regime 1	60.7%	0.871	7.32	26.73	85.3%
Kolkata	Regime 2	39.3%	0.810	7.89	32.65	51.2%




Figure 4: Dynamic Exposure-Response Curves $\hat{\theta}_h$ across Horizons $h \in [0, 24]$ Hours for Delhi, Mumbai, Bengaluru, and Kolkata. Shaded bands represent 95% HAC confidence intervals. Regime stratification reveals substantial structural heterogeneity in dynamic pollutant accumulation over multi-hour horizons.

G.2 Regularization Sensitivity in Ill-Conditioned Regimes (Kolkata Airshed)

To verify the theoretical behavior established in Theorem 6 under real-world conditions, Table 4 reports the empirical sensitivity of OR-DML across a fine 13-point regularization grid $\lambda \in [0.0, 1.0]$ in the ill-conditioned Kolkata airshed ($\lambda_{\min} = 0.0024$).

As predicted by Theorem 6, when $\lambda \rightarrow 0$, standard errors explode to 1.1071 due to near-singularity in residual treatment variation. Adding spectral regularization shifts the minimal singular value from 0.0049 to 0.1049 at $\lambda = 0.10$, reducing standard errors by an order of magnitude ($1.1071 \rightarrow 0.0601$) while shrinking point estimates conservatively toward zero. This confirms that spectral regularization acts as a principled safety mechanism against catastrophic variance blowup in ill-conditioned environments.

Table 4: Regularization Sensitivity Grid for Kolkata Airshed ($\lambda_{\min} = 0.0024$, $N = 1,209$). Standard errors decrease monotonically from 1.1071 at $\lambda = 0.0$ to 0.0063 at $\lambda = 1.0$, stabilizing estimation variance while shrinking point estimates toward zero.

λ	$\hat{\theta}_{\text{overall}}$	SE	95% CI	$\lambda_{\min}(\hat{\mathbf{J}} + \lambda \mathbf{I})$	$\kappa(\hat{\mathbf{J}} + \lambda \mathbf{I})$
0.0000	-0.5570	1.1071	$[-2.7269, +1.6130]$	0.0049	1.28
0.0001	-0.5485	1.0878	$[-2.6807, +1.5836]$	0.0050	1.28
0.0005	-0.5172	1.0170	$[-2.5105, +1.4760]$	0.0054	1.28
0.0010	-0.4828	0.9406	$[-2.3263, +1.3607]$	0.0059	1.28
0.0024	-0.4068	0.7776	$[-1.9309, +1.1173]$	0.0073	1.28
0.0050	-0.3147	0.5890	$[-1.4691, +0.8398]$	0.0099	1.28
0.0100	-0.2191	0.4022	$[-1.0074, +0.5691]$	0.0149	1.28
0.0200	-0.1363	0.2463	$[-0.6191, +0.3464]$	0.0249	1.28
0.0500	-0.0639	0.1140	$[-0.2872, +0.1595]$	0.0549	1.28
0.1000	-0.0339	0.0601	$[-0.1517, +0.0840]$	0.1049	1.28
0.2000	-0.0175	0.0309	$[-0.0781, +0.0432]$	0.2049	1.28
0.5000	-0.0071	0.0126	$[-0.0318, +0.0175]$	0.5049	1.28
1.0000	-0.0036	0.0063	$[-0.0160, +0.0088]$	1.0049	1.28

G.3 Data Imputation Sensitivity Analysis (Strictly Observed vs. Cleaned)

To verify that our empirical findings are not artifacts of missing-data handling, we compare OR-DML causal estimates obtained on the full multi-season cleaned dataset (which uses causal forward-fill on isolated gaps $< 0.7\%$ of hours) against a strictly observed dataset where every row containing any imputed value is entirely deleted.

Table 5: Sensitivity Analysis Comparing Full Dataset ($< 0.7\%$ Causal Forward-Fill) vs. Strictly Observed Dataset (Rows with Any Imputed Entry Deleted) Across Four Megacities. Point estimates and standard errors exhibit identical signs and overlapping confidence bands.

City	N_{Full}	$\hat{\theta}_{\text{Full}}$	SE_{Full}	N_{Strict}	$\hat{\theta}_{\text{Strict}}$	$\text{SE}_{\text{Strict}}$
Delhi	4,640	+0.3634	0.0460	4,602	+0.5366	0.0674
Mumbai	4,180	-5.7671	2.1367	4,146	-10.9323	2.6792
Bengaluru	4,093	+0.0046	0.0101	4,084	+0.0202	0.0112
Kolkata	1,209	-0.5563	1.1055	1,206	-0.1924	1.0148

As documented in Table 5:

- In Delhi, deleting the 38 forward-filled hours (0.8%) yields $\hat{\theta}_{\text{Strict}} = +0.5366 \pm 0.0674$, matching the positive elasticity observed under the full series ($+0.3634 \pm 0.0460$).
- In Mumbai, deleting 34 hours (0.8%) retains the strong negative dispersion elasticity (-10.9323 ± 2.6792 vs. -5.7671 ± 2.1367).
- In Bengaluru, deleting 9 hours (0.2%) preserves the near-zero net effect ($+0.0202 \pm 0.0112$ vs. $+0.0046 \pm 0.0101$).
- In Kolkata, deleting 3 hours (0.2%) preserves identical negative directionality with overlapping confidence intervals (-0.1924 ± 1.0148 vs. -0.5563 ± 1.1055).

These results confirm that our conclusions are robust to missing data and do not depend on the causal forward-fill protocol.

H LATENTREGIMEBENCH: MULTI-DIMENSIONAL BENCHMARK RESULTS

To systematically evaluate machine learning under persistent, partially observed, and shifting latent regimes, we conduct comprehensive evaluations on **LatentRegimeBench** across five diagnostic configurations ($N = 1,200$, 30 Monte Carlo replications each) and four representation architectures (Table 6 and Table 7, with summary visualizations in Figure 5).

Table 6: LatentRegimeBench: Multi-Dimensional Stress Test across 5 Strategic Difficulty Configurations (30 Replications each). In well-conditioned regimes (C1, C3), OR-DML achieves a $> 5.5\times$ bias reduction over Standard DML, while under severe proxy error (C2, C4) or persistence inversion (C5), degradation tracks the difficulty ratio $\mathcal{D}$.

Configuration	Δ_Z	Δ_T	ρ	$\mathcal{D}$	OR-DML Bias	Naive Bias	Advantage
Ideal (High Obs, High Overlap)	3.5	5.0	0.85	0.001	0.2165	1.2061	$5.57\times$
Adversarial A (Low Obs, High Overlap)	0.3	5.0	0.85	0.709	1.1950	1.2061	$1.01\times$
Adversarial B (High Obs, Weak Overlap)	3.5	0.4	0.85	0.001	0.2498	1.3886	$5.56\times$
Challenging (Low Obs, Weak Overlap)	0.5	0.5	0.90	2.748	1.3237	1.3862	$1.05\times$
High-Persistence Inversion	2.0	2.0	0.96	0.159	1.5934	1.3175	$0.83\times$

Table 7: Latent Representation Zoo: In-Distribution vs. Out-of-Distribution (OOD) Shift Regime Recovery ($N = 2,000$). Gaussian HMM and Linear SSM demonstrate superior calibration (ECE 0.0083 and 0.0236), while Neural GRU and Transformer maintain robust macro F1 across temporal distribution shifts.

Architecture	In-Dist F1	OOD F1	Gap	In-Dist Brier	OOD Brier	In-Dist ECE
Gaussian HMM (Causal Filter)	0.9720	0.9765	-0.0044	0.0219	0.0159	0.0083
Neural GRU Encoder	0.9695	0.9750	-0.0055	0.0223	0.0220	0.0172
Causal Transformer Encoder	0.9708	0.9775	-0.0067	0.0229	0.0236	0.0236
Linear State Space Model (SSM)	0.9720	0.9755	-0.0035	0.0240	0.0295	0.0236



Figure 5: Universal LatentRegimeBench Evaluation. (a) Calibration of causal estimation bias against theoretical degradation ratio $\mathcal{D} = \bar{\epsilon}_\gamma / \lambda_{\min}(\mathbf{J})$. (b) Decision-time abstention safety comparing unconstrained point MSE against uncertainty-gated MSE. (c) Representation zoo out-of-distribution (OOD) generalization across HMM, GRU, Causal Transformer, and Linear SSM.

I CROSS-DOMAIN TRANSFER TO MACRO FINANCIAL MARKET REGIMES

To verify that the latent regime confounding and difficulty framework reflects a general mathematical reality of non-stationary sequential data rather than an idiosyncratic atmospheric feature, we evaluate the Regime Intelligence engine on quantitative alternative financial market data ($N = 1,500$ daily trading periods). Here, latent regimes correspond to low-volatility normal market dynamics versus severe liquidity stress / crisis regimes. Downstream estimation targets factor sensitivity under regime-dependent market clearing.

Table 8 partitions out-of-sample trading periods into dynamic regime risk quartiles $R_t = H(\gamma_t) / \lambda_{\min}(\mathbf{J}_t)$. In Q4 (Severe Risk), tail 95% forecast MSE surges to 4.2931 (compared to 1.7461 in moderate periods), driven by catastrophic causal ill-conditioning ($\lambda_{\min} \rightarrow 0.211$). Table 9 demonstrates the power of identification-aware abstention: setting $\tau = 1.5$ abstains on just 2.8% of high-uncertainty periods, avoiding predictions with an extreme average MSE of 1.4341 and reducing tail 95% MSE from 3.3514 to 2.8705. Figure 6 charts the cross-domain transfer dynamics.

Table 8: Cross-Domain Financial Macro Regimes: Dynamic Regime Risk Quartiles ($N = 1,500$). In Q4 (Severe Risk), extreme tail 95% MSE escalates to 4.2931, reflecting acute causal ill-conditioning ($\lambda_{\min} = 0.211$).

Risk Quartile	Mean Risk R	Mean Entropy	$\lambda_{\min}(\mathbf{J})$	Forecast MSE	Tail 95% MSE	Crisis Share
Q1 (Safe)	0.000	0.000	0.365	0.9278	3.3514	96.5%
Q2 (Low Risk)	0.000	0.000	0.306	0.6191	2.8449	27.7%
Q3 (Moderate)	0.000	0.000	0.324	0.4899	1.7461	3.2%
Q4 (Severe Risk)	19.164	0.103	0.211	0.9419	4.2931	18.9%

Table 9: Dynamic Decision Abstention Policy on Financial Regimes. Abstaining on the most ill-conditioned 2.8% of periods ($\tau = 1.5$) avoids catastrophic errors (MSE = 1.4341 avoided vs. 0.7248 accepted).

τ	Coverage	Abstain	Accepted MSE	Avoided MSE	MSE Reduction	Tail 95% MSE
1.5	97.2%	2.8%	0.7248	1.4341	2.67%	2.8705
2.0	98.1%	1.9%	0.7417	0.8947	0.40%	3.0751
2.5	98.4%	1.6%	0.7435	0.8183	0.16%	3.0806
3.0	98.6%	1.4%	0.7433	0.8426	0.19%	3.0776
4.0	99.0%	1.0%	0.7435	0.8604	0.16%	3.0754




Figure 6: Quantitative Macro Financial Market Regime Transfer. (a) Daily asset factor returns colored by latent volatility regime. (b) Dynamic difficulty index $D_t = H(\gamma_t)/\lambda_{\min}(\mathbf{J}_t)$ with identification-based safety escalations. (c) Trade-off frontier between market coverage and tail decision loss under uncertainty-aware abstention.