
Reliable Causal Estimation under Latent Markov Confounding

Anonymous Author
Anonymous Institution

Abstract

Causal inference from dependent observational time series is complicated when an unobserved, persistent state jointly affects treatment assignment and outcomes. Standard Double Machine Learning (DML) controls for observed covariates through orthogonal residualization, but does not by itself resolve confounding induced by an imperfect proxy for latent state dynamics. We study causal estimation under latent Markov confounding and characterize how latent-state uncertainty and downstream task conditioning propagate into causal error. We develop a regime-aware orthogonal estimator combining probabilistic state inference, temporally purged cross-fitting, and regularized inversion of the cross-regime score Jacobian. Our analysis separates nuisance-estimation error, latent-state recovery error, and variability induced by persistent regime occupation, and gives conditions under which the estimator approaches the corresponding oracle-state estimator as latent-state uncertainty decreases. We evaluate these properties using real-data-calibrated semi-synthetic time series constructed from high-frequency atmospheric observations across heterogeneous Indian cities, preserving persistence, seasonality, missingness, and regime overlap while keeping causal effects known. Across controlled variations in state persistence and observability, the experiments show how latent-state uncertainty predicts causal bias and when regime-aware estimation improves over standard and temporally blocked DML. Real sensor networks provide an observational stress test across four megacities spanning distinct

residual treatment variance scales and telemetry artifacts; falsification checks are reported explicitly where temporal causal interpretation is unsupported. Together, the results show that reliable causal estimation with latent sequential confounding requires evaluating representation uncertainty jointly with the information geometry of the downstream task.

1 INTRODUCTION

Sequential machine learning systems frequently operate in environments governed by persistent, partially observed context—such as unobserved physical, macroeconomic, or operational regimes. A widespread design assumption in such systems is that predictive confidence or calibration over the latent state is sufficient for downstream reliability. However, this assumption fails: *accurate latent-state recovery does not guarantee reliable downstream inference because representation error is filtered through the information geometry of the task that consumes it*. When the downstream task is ill-conditioned—such as estimating causal effects under collinear treatments or weak overlap—even small posterior proxy error is catastrophically amplified, causing severe downstream failure despite seemingly confident latent representations. The same latent-state uncertainty can be informative for one task and nearly useless for another because the downstream information geometry differs.

Continuous-treatment causal inference in atmospheric systems provides a natural laboratory: estimating the causal effect of variation in the observed NO₂ exposure proxy on ambient PM_{2.5} (Dominici et al., 2014) (translating this observational exposure effect to an emissions policy requires structural transport assumptions). Observational time series in this domain violate two foundational assumptions: (1) **Persistent Latent Regime Confounding**: Planetary boundary-layer dynamics transition between discrete thermodynamic regimes (e.g., nocturnal inversion vs. convective clearance), jointly governing precursor dispersion (T_t)

and ambient particulate concentration (Y_t). (2) **Imperfect Overlap & Dependence:** Latent regimes are unobserved; auxiliary proxies Z_t exhibit overlap ($\varepsilon_\gamma > 0$), while autoregression breaks exchangeability (Chernozhukov et al., 2018), leaking errors across folds.

When standard Double Machine Learning (Chernozhukov et al., 2018) or tree-based plug-in estimators are applied naively to environmental time series without accounting for latent regimes, the resulting causal estimands can exhibit severe bias and, in our benchmark, sign inversions-falsely suggesting that precursor reductions exacerbate ambient pollution.

Related Work. We position our contribution across four neighboring literatures: (1) **Double Machine Learning & Orthogonal Inference:** Chernozhukov et al. (2018, 2022b); Foster and Syrgkanis (2023); Nie and Wager (2021); Kennedy (2022) achieve $\sqrt{N}$ -rate inference under unconfoundedness given X_t . (2) **Dependent & Time-Series Causal Inference:** Ciganovic et al. (2026); Semenova et al. (2023); Lewis and Syrgkanis (2020); Bai (2009); Bonhomme and Manresa (2015); Hamilton (1989) relax exchangeability under mixing, but do not address latent Markov confounding. (3) **Latent Markov & Proximal Inference:** Clouth et al. (2025); Huang et al. (2020); Miao et al. (2018); Tchetgen Tchetgen et al. (2020); Mastouri et al. (2021); Cappé et al. (2005) use proxies for unmeasured confounders, focusing on asymptotic identification rather than downstream finite-sample score conditioning. (4) **Weak Overlap & Selective Reliability:** D’Amour et al. (2021); Shalit et al. (2017); Crump et al. (2009); Guo et al. (2017); Geifman and El-Yaniv (2017); Chow (1970); El-Yaniv and Wiener (2010) address overlap and abstention, without linking latent uncertainty to coupled score conditioning.

Contributions. We separate latent-state uncertainty from task conditioning and show how their interaction controls downstream degradation:

- **Decoupling Context Uncertainty from Task Conditioning:** We prove that downstream estimation suffers two distinct degradation mechanisms: imperfect latent recovery (ε_γ) and weak task conditioning ($\lambda_{\min}(\mathbf{J})$), showing that predictive state confidence does not imply downstream task reliability (Theorems 2-3).
- **Principled Difficulty Measure & Observable Bridge:** We formulate causal difficulty $\mathcal{D}_{\text{causal}} = \frac{\varepsilon_\gamma}{\lambda_{\min}(\mathbf{J})}$ and establish rolling entropy $\mathcal{D}_{\text{operational}}(t) = \frac{H(\gamma_t)}{\lambda_{\min}(\mathbf{J}_t)}$ as an observable multiplicative surrogate under two-sided spectral comparability (Proposition 1).
- **Sequential Overlap-Aware Estimation (OR-DML):** We develop spectrally regularized OR-

DML with purged block cross-fitting, proving leading-order risk bounds and establishing centered asymptotic normality under dependent data (Theorems 4-5).

- **Representation-to-Task Reliability:** A corrected 7×7 factorial intervention independently varies proxy error and post-residualization task conditioning; $\varepsilon_\gamma / \lambda_{\min}(\mathbf{J})$ predicts causal error far better than either component alone ($\rho = 0.995$ vs. $0.626/0.521$). Across 15 shifted worlds, adding posterior NLL to state-recovery F_1 nearly doubles leave-one-world-out reliability R^2 ($0.113 \rightarrow 0.221$), while temperature scaling cuts neural ECE by up to 60% without improving downstream ROC-AUC.

2 STRUCTURAL CAUSAL MODEL & IDENTIFICATION

Let $(\Omega, \mathcal{F}, P)$ be a probability space. Consider a discrete-time longitudinal process observed at equidistant intervals $t \in \{1, \dots, N\}$.

2.1 Structural Equations

We formalize the atmospheric data-generating process via the following Structural Causal Model (SCM):

$$S_t \sim \text{Markov}(\Pi, \pi_0), \quad S_t \in \{1, \dots, K\}, \quad (1)$$

$$Z_t = \mu_Z(S_t) + \xi_t, \quad (2)$$

$$X_t = \phi(X_{t-1}, S_t) + \epsilon_t, \quad (3)$$

$$T_t = m(X_t, S_t) + V_t, \quad \mathbb{E}[V_t | X_t, S_t] = 0, \quad (4)$$

$$Y_t = \theta(S_t)T_t + g(X_t, S_t) + U_t, \quad (5)$$

with $\mathbb{E}[U_t | X_t, S_t, T_t] = 0$, where:

- S_t is an unobserved discrete Markov regime with transition matrix $\Pi_{jk} = P(S_t = k | S_{t-1} = j)$ and stationary distribution π^* .
- $Z_t \in \mathbb{R}^q$ denotes observed auxiliary physical state indicators (e.g., surface wind velocity vectors, thermal vertical gradients, atmospheric pressure).
- $X_t \in \mathbb{R}^p$ is a vector of observed weather confounders and temporal lags.
- $T_t \in \mathbb{R}$ is the continuous treatment (precursor emission proxy, e.g., NO_2).
- $Y_t \in \mathbb{R}$ is the continuous outcome (ambient pollutant, e.g., $\text{PM}_{2.5}$).
- $\theta(S_t) = \sum_{k=1}^K \theta_k^* \mathbb{I}(S_t = k)$ represents the regime-specific causal effect.




Figure 1: Temporal Causal Graph. Latent transitions $S_{t-1} \rightarrow S_t$ govern persistent regime dynamics; S_t acts as a common cause for proxy Z_t , confounders X_t , treatment T_t , and outcome Y_t . Omitting S_t induces omitted regime bias.

2.2 Identification Under Latent Regimes

Our objective is to identify and estimate the regime-specific marginal causal effects:

$$\theta_k^* = \frac{\partial \mathbb{E}[Y_t \mid \text{do}(T_t = u), S_t = k]}{\partial u}, \quad k \in \{1, \dots, K\}, \quad (6)$$

as well as the stationary Average Treatment Effect (ATE):

$$\theta_{\text{ATE}}^* = \sum_{k=1}^K \pi_k^* \theta_k^*. \quad (7)$$

Assumption 1 (Regime-Conditional Unconfoundedness & Innovation Exogeneity). Conditional on observed covariates X_t , latent regime S_t , and proxy history $Z_{1:N}$, potential outcomes are unconfounded ($Y_t(u) \perp T_t \mid (X_t, S_t, Z_{1:N}), \forall u$) and structural innovations have zero mean:

$$\mathbb{E}[U_t \mid X_t, S_t, T_t, Z_{1:N}] = 0. \quad (8)$$

Assumption 2 (Overlap, Envelopes & State Positivity). There exist constants $\eta, c_0 > 0$, $B_T, B_Y < \infty$, and $c_1, c_2 > 0$ such that:

1. **State Positivity & Envelopes:** Stationary distribution satisfies $\min_k \pi_k^* \geq \eta$, and for almost every x and all k , $P(S_t = k \mid X_t = x) \geq c_0$, $\mathbb{E}[\gamma_{tk} \mid X_t = x] \geq c_0$. Variables satisfy $|T_t| \leq B_T$, $|Y_t| \leq B_Y$ a.s., so $|\tilde{T}_{tk}| \leq M_T \equiv 2B_T$, $|\tilde{Y}_{tk}| \leq M_Y \equiv 2B_Y$.
2. **Within-Regime Treatment Overlap:** For all $k \in \{1, \dots, K\}$ and almost everywhere on $\text{supp}(X_t)$:

$$c_1 \leq \text{Var}(T_t \mid X_t, S_t = k) \leq c_2. \quad (9)$$

Assumption 3 (Latent Proxy Sufficiency). The auxiliary state proxy sequence $Z_{1:N}$ satisfies latent proxy sufficiency: $P(S_t = k \mid X_t, T_t, Z_{1:N}) = P(S_t = k \mid Z_{1:N}) \equiv \gamma_{tk}$ (or $P(S_t = k \mid Z_{1:t})$ for causal forward filtering) almost surely for all $k \in \{1, \dots, K\}$. *Discussion.* That is, $Z_{1:N}$ screens off the latent state S_t from

(X_t, T_t) , framing it as a proxy-sufficiency identification restriction. Combined with Assumption 1, this guarantees $\mathbb{E}[Y_t \mid X_t, T_t, Z_{1:N}] = \sum_{k=1}^K \gamma_{tk} [\theta_k^* T_t + g_k(X_t)]$. This identifying restriction requires $Z_{1:N}$ to be sufficiently informative that conditioning on (X_t, T_t) provides no additional regime information beyond $Z_{1:N}$.

Assumption 4 (Nuisance Estimation and Mixing Rates). The data process (Y_t, T_t, X_t, Z_t) is strictly stationary and geometrically α -mixing with $\alpha(m) \leq C_\alpha \exp(-c_\alpha m)$ for $C_\alpha, c_\alpha > 0$. Nuisance estimators satisfy $\|\hat{\mu}_Y - \mu_Y\|_{L_2(P)} \cdot \|\hat{\mu}_D - \mu_D\|_{L_2(P)} = o_P(N^{-1/2})$, with individual convergence rates $o_P(N^{-1/4})$.

3 FAILURE OF STANDARD DML: THE OMITTED REGIME BIAS

Intuition. When unobserved regimes S_t jointly shift treatment and outcome baselines, partialling out observed controls X_t cannot purge latent regime variation. Instead, by reducing residual treatment variance in the denominator without eliminating the latent covariance in the numerator, standard cross-sectional residualization actively amplifies the omitted confounding.

Theorem 1 (Decomposition of Omitted Regime Bias). *Suppose the data are generated according to SCM (1)-(5). Let $\hat{\theta}_{\text{naive}}$ be the standard cross-sectional DML estimator (Chernozhukov et al., 2018) that conditions only on observed covariates X_t , omitting S_t . Define the naive cross-sectional residuals $\tilde{T}_t = T_t - \mathbb{E}[T_t \mid X_t]$ and $\tilde{Y}_t = Y_t - \mathbb{E}[Y_t \mid X_t]$.*

Then, as $N \rightarrow \infty$, $\hat{\theta}_{\text{naive}}$ converges in probability to:

$$\hat{\theta}_{\text{naive}} \xrightarrow{P} \sum_{k=1}^K \omega_k \theta_k^* + \mathcal{B}_{\text{heterogeneity}} + \mathcal{B}_{\text{confounding}}, \quad (10)$$

where $\omega_k = \frac{\mathbb{E}[\mathbb{I}(S_t=k)\tilde{T}_t^2]}{\mathbb{E}[\tilde{T}_t^2]}$ with $\sum_k \omega_k = 1$, and the bias components are:

$$\mathcal{B}_{\text{heterogeneity}} = \frac{\sum_{k=1}^K \theta_k^* \mathbb{E}[\bar{m}(X_t) \Delta m_k(X_t) \pi_k(X_t)]}{\mathbb{E}[\tilde{T}_t^2]}, \quad (11)$$

$$\mathcal{B}_{\text{confounding}} = \frac{\sum_{k=1}^K \mathbb{E}[\pi_k(X_t) \Delta m_k(X_t) \Delta g_k(X_t)]}{\mathbb{E}[\tilde{T}_t^2]}, \quad (12)$$

with $\pi_k(X_t) \equiv P(S_t = k \mid X_t)$, $\bar{m}(X_t) \equiv \mathbb{E}[T_t \mid X_t]$, $\Delta m_k(X_t) \equiv \mathbb{E}[T_t \mid X_t, S_t = k] - \bar{m}(X_t)$, and $\Delta g_k(X_t) \equiv g(X_t, k) - \mathbb{E}[g(X_t, S_t) \mid X_t]$ (complete algebraic proof in Appendix A).

Corollary 1 (Sign Inversion & Bias Amplification). *When latent stagnation induces strong positive confounding ($\Delta m_k(X_t) \Delta g_k(X_t) \gg 0$), but residual treatment variance $\mathbb{E}[\tilde{T}_t^2]$ is attenuated toward its lower*




Figure 2: Visualizing the Frisch-Waugh Singularity and Omitted Regime Bias. (a) As residual treatment variance $\mathbb{E}[\tilde{T}_t^2] \rightarrow 0$ via partialling-out X_t , omitted regime bias $\mathcal{B}_{\text{confounding}}$ explodes hyperbolically in Standard DML, while OR-DML (Oracle Regimes, $\varepsilon_\gamma = 0$) remains strictly unbiased ($\mathcal{B} = 0$). (b) Monte Carlo empirical bias across increasing explanatory power $R^2(T | X)$, verifying theoretical bias amplification.

bound $c_1 \ll \Delta m$ via partialling out X_t (or under parameter sequences where treatment innovation variance $\sigma_V^2 \rightarrow 0$ while latent separation remains fixed), the omitted confounding bias $\mathcal{B}_{\text{confounding}}$ undergoes severe bias amplification (Pearl, 2010). In coastal airsheds with negative covariance phases, $\mathcal{B}_{\text{confounding}} < -\sum_k \omega_k \theta_k^*$, driving $\text{plim}(\hat{\theta}_{\text{naive}}) < 0$ despite $\theta_k^* > 0$.

Interpretation. Standard DML cannot eliminate omitted regime confounding: conditioning on observed controls X_t shrinks residual treatment variance $\mathbb{E}[\tilde{T}_t^2]$ toward c_1 without reducing latent regime covariance, actively amplifying bias.

4 METHODOLOGY: OVERLAP-AWARE REGIME DML

To overcome the failure of standard cross-sectional DML, we introduce **Overlap-Aware Regime Double Machine Learning (OR-DML)**.

4.1 Latent Regime Formulation & Observability Diagnostics

Let $Z_t \in \mathbb{R}^q$ denote exogenous contemporaneous meteorological covariates assumed unaffected by the exposure proxy. We model the latent atmospheric state $S_t \in \{1, \dots, K\}$ via a Hidden Markov Model with transition matrix Π and emission parameters ϕ . Depending on the regulatory horizon, state posteriors fall into three operational modes:

1. **Retrospective Smoothed Posteriors (Offline):** $\gamma_{tk}^{\text{smooth}} \equiv P(S_t = k | Z_{1:N})$ for retrospective environmental epidemiology and structural causal auditing;

2. **Training-Fitted Forward Filtered Posteriors:** $\gamma_{tk}^{\text{filter}} \equiv P(S_t = k | Z_{1:t}, \hat{\phi})$ where dynamics $\hat{\phi}$ are calibrated on a historical window strictly preceding evaluation ($\mathcal{T}_{\text{cal}} < \mathcal{T}_{\text{eval}}$), with state updates proceeding strictly causally on $Z_{1:t}$; or
3. **Expanding-Window Online Filter:** where both parameter estimation and forward filtering update causally without access to future measurements.

Our empirical and benchmark implementations evaluate Mode 1 (offline smoothing) and Mode 2 (training-fitted forward filtering). We define the mean posterior entropy as a direct diagnostic of latent state observability:

$$\bar{H}(\gamma) \equiv -\frac{1}{N} \sum_{t=1}^N \sum_{k=1}^K \gamma_{tk} \ln(\gamma_{tk} + \epsilon). \quad (13)$$

A Priori Diagnostics: Standardized separation $\Delta_Z < 0.5$ or entropy $\bar{H} \rightarrow \ln K$ signals severe proxy uncertainty and triggers a weak-identification diagnostic warning (rules in Appendix C.3), distinguishing structural non-identification ($\Delta_Z = 0$, Theorem 2) from finite-sample proxy dilution.

4.2 State-Weighted Coupled System & Spectral Regularization

To avoid heuristic decoupled regressions that suffer cross-talk bias under regime overlap, OR-DML estimates regime-specific conditional expectations via state-weighted regressions targeting population expectations $m_k^\gamma(x) \equiv \frac{\mathbb{E}[\gamma_{tk} T_t | X_t = x]}{\mathbb{E}[\gamma_{tk} | X_t = x]}$ and $\mu_k^\gamma(x) \equiv \frac{\mathbb{E}[\gamma_{tk} Y_t | X_t = x]}{\mathbb{E}[\gamma_{tk} | X_t = x]}$:

$$\hat{m}_k(X) = \arg \min_m \sum_{t \in \mathcal{T}_b} \gamma_{tk} (T_t - m(X_t))^2, \quad (14)$$

$$\hat{\mu}_k(X) = \arg \min_\mu \sum_{t \in \mathcal{T}_b} \gamma_{tk} (Y_t - \mu(X_t))^2. \quad (15)$$

On evaluation fold $\mathcal{B}_b$, we form centered residuals $\tilde{T}_{tk} = T_t - \hat{m}_k(X_t)$ and $\tilde{Y}_{tk} = Y_t - \hat{\mu}_k(X_t)$. The empirical Gram matrix $\hat{\mathbf{J}} \in \mathbb{R}^{K \times K}$ and score vector $\hat{\mathbf{S}} \in \mathbb{R}^K$ are constructed with matched weights:

$$\hat{J}_{kk} = \frac{1}{N} \sum_{t=1}^N \gamma_{tk} \tilde{T}_{tk}^2, \quad \hat{S}_k = \frac{1}{N} \sum_{t=1}^N \gamma_{tk} \tilde{T}_{tk} \tilde{Y}_{tk}, \quad (16)$$

$$\hat{J}_{jk} = \frac{1}{N} \sum_{t=1}^N \gamma_{tj} \gamma_{tk} \tilde{T}_{tj} \tilde{T}_{tk} \quad (j \neq k). \quad (17)$$

When regime separation is weak or treatments are collinear, $\lambda_{\min}(\hat{\mathbf{J}}) \rightarrow 0$, rendering the unregularized solve $\hat{\mathbf{J}}^{-1} \hat{\mathbf{S}}$ ill-conditioned. OR-DML supports both coupled regularized solve $\hat{\theta}_\lambda = (\hat{\mathbf{J}} + \lambda \mathbf{I})^{-1} \hat{\mathbf{S}}$ and decoupled regime-wise soft DML $\hat{\theta}_k = \hat{S}_k / (\hat{J}_{kk} + \lambda)$,

Algorithm 1 Overlap-Aware Regime DML (OR-DML)

Require: Series $\{Y_t, T_t, X_t, Z_t\}_{t=1}^N$, blocks B , embargo τ^* , penalty $\lambda \geq 0$.

- 1: Infer posteriors γ_t (Mode 1 smoothing or Mode 2 forward filtering).
 - 2: Cross-fit state-weighted nuisances $\hat{m}_k, \hat{\mu}_k$ with embargo buffer τ^* .
 - 3: Form centered residuals $\tilde{T}_{tk} = T_t - \hat{m}_k(X_t)$ and $\tilde{Y}_{tk} = Y_t - \hat{\mu}_k(X_t)$.
 - 4: Form Gram $\hat{\mathbf{J}} = \frac{1}{N} \sum_t \gamma_t \gamma_t^\top \odot \tilde{\mathbf{T}}_t \tilde{\mathbf{T}}_t^\top$ and score $\hat{\mathbf{S}} = \frac{1}{N} \sum_t \gamma_t \odot \tilde{\mathbf{T}}_t \odot \tilde{\mathbf{Y}}_t$.
 - 5: Compute $\hat{\boldsymbol{\theta}}_\lambda = (\hat{\mathbf{J}} + \lambda \mathbf{I})^{-1} \hat{\mathbf{S}}$ and diagnostics $(\lambda_{\min}, \kappa, \bar{H})$.
-

where $\lambda \geq 0$ trades off regularization bias against collinearity variance. To ensure scale-invariance across units, the dimensionless Gram matrix is $\hat{\mathbf{J}}_{\text{std}} \equiv \hat{\mathbf{J}} / \hat{\sigma}_T^2$ with $\hat{\sigma}_T^2 \equiv \frac{1}{K} \sum_k \frac{1}{N} \sum_t \tilde{T}_{tk}^2$, giving scale-free $\lambda_{\min}(\hat{\mathbf{J}}_{\text{std}})$ and condition number $\kappa(\hat{\mathbf{J}}) \equiv \lambda_{\max} / \lambda_{\min}$. *Scalability & Structural Abstention:* End-to-end complexity is $\mathcal{O}(BK \cdot C_{\text{nuisance}} + NK^2 + K^3)$; when data lack identifying variation entirely ($\kappa(\hat{\mathbf{J}}) > 10^4$ or $\text{CI}(\lambda_{\min})$ covers 0), systems diagnose weak identification and abstain rather than masking near-singularity with excessive shrinkage. We distinguish three causal estimand targets: (i) the true stationary population ATE $\theta_{\text{PATE}}^* = \pi^{*\top} \boldsymbol{\theta}^*$, (ii) the estimated stationary ATE $\hat{\theta}_{\text{PATE}} = \hat{\pi}(\hat{\boldsymbol{\Lambda}})^\top \hat{\boldsymbol{\theta}}_\lambda$, and (iii) the empirical sample occupation ATE $\hat{\theta}_{\text{SATE}} = \bar{\gamma}_N^\top \hat{\boldsymbol{\theta}}_\lambda$, where $\bar{\gamma}_N \equiv \frac{1}{N} \sum_{t=1}^N \gamma_t$.

4.3 Purged Block Cross-Fitting with Bartlett Embargo

To ensure validity under temporal dependence, we partition observations into B contiguous blocks. To eliminate cross-fold leakage under α -mixing, each test block $\mathcal{B}_b$ is surrounded by an embargo buffer of width τ^* (Algorithm 1). We determine τ^* from the data via the Bartlett significance threshold:

$$\tau^* = \max \left\{ h \geq 1 : \max_j |\hat{\rho}_j(h)| \geq \frac{1.96}{\sqrt{N}} \right\} + 1. \quad (18)$$

In theory, $\tau^* \geq C \log N$ with $C > 1/c_\alpha$ guarantees vanishing coupling (Theorem 5). In our 2-hourly cadence ($N \approx 14,000$, $\log N \approx 9.55$), Bartlett calibration yields 12-24 rows (24-48 hours), satisfying $\tau^* \geq 2 \log N$ (auto_embargo=True).

5 THEORETICAL ANALYSIS

We formalize three central results: (1) **Non-Identification:** latent context can induce unidentifiability; (2) **Graceful Degradation:** proxy error scales

as $\varepsilon_\gamma / \lambda_{\min}(\mathbf{J})$; and (3) **Observable Difficulty & Regularization:** task conditioning yields an observable surrogate for risk control.

Intuition. If auxiliary proxies Z_t contain zero information distinguishing latent regimes, different underlying causal structures generate observationally indistinguishable joint distributions, rendering point identification impossible regardless of sample size.

Theorem 2 (Non-Identification Under Unconstrained Overlap). *Suppose the latent state is unobserved, and auxiliary proxies Z_t satisfy $P(Z_t \mid S_t = j) = P(Z_t \mid S_t = k)$ for all j, k . When transitions are memoryless ($p = 0.5$) or under static cross-sectional proxy evaluation, the regime-specific causal effect vector $\boldsymbol{\theta}^* = (\theta_1^*, \dots, \theta_K^*)^\top$ is not point-identified from the joint distribution of observables (Y_t, T_t, X_t, Z_t) : there exist structural causal models P_1 and P_2 satisfying Assumptions 1-3 with identical joint observable laws $P_1(Y_{1:N}, T_{1:N}, X_{1:N}, Z_{1:N}) \equiv P_2(Y_{1:N}, T_{1:N}, X_{1:N}, Z_{1:N})$ for all $N \geq 1$, yet possessing strictly distinct, non-permutation causal effect vectors $\boldsymbol{\theta}^{(1)} \neq \boldsymbol{\theta}^{(2)}$.*

Interpretation. No estimator can recover regime-specific effects when the proxy provides no state information; this is a fundamental limitation of the observable information structure, not of optimization or sample size (proof in Appendix B; we contrast this with dynamic models where non-trivial sequence transitions can achieve generic identification (Allman et al., 2009)).

Intuition. Downstream estimation error has two qualitatively different sources: posterior proxy error ε_γ from imperfect state inference, and causal ill-conditioning $\lambda_{\min}(\mathbf{J})$ from collinear residual treatment variation. Theorem 3 proves that these mechanisms enter multiplicatively.

Theorem 3 (Graceful Degradation Error Bound). *Let $\mathbf{H}_t = \mathbf{e}_{S_t}$, $\gamma_t \in \Delta^K$ be a state proxy with $\varepsilon_\gamma \equiv \mathbb{E} \|\gamma_t - \mathbf{H}_t\|_1$, and $\mathbf{J}, \mathbf{S}$ be the population coupled Gram matrix and score from proxy nuisances $(m_k^\gamma, \mu_k^\gamma)$. Under Assumption 2 and $\|\boldsymbol{\theta}^*\|_2 \leq M_\theta$, if $\lambda_{\min}(\mathbf{J}) > 0$:*

- (i) $\boldsymbol{\theta}_\gamma \equiv \mathbf{J}^{-1} \mathbf{S}$ satisfies $\|\boldsymbol{\theta}_\gamma - \boldsymbol{\theta}^*\|_2 \leq C \varepsilon_\gamma / \lambda_{\min}(\mathbf{J})$ with $C = C_S + M_\theta C_J$, $C_S = \sqrt{K} (M_T M_Y + M_Y \tilde{C}_m + M_T \tilde{C}_\mu)$, $C_J = K(2M_T^2 + 2M_T \tilde{C}_m)$, $\tilde{C}_m = 2B_T / c_0^2$, $\tilde{C}_\mu = 2B_Y / c_0^2$.
- (ii) Let $r_N = \max_k (\|\hat{m}_k - m_k^\gamma\|_{L_2} + \|\hat{\mu}_k - \mu_k^\gamma\|_{L_2})$. Under purged cross-fitting and Assumption 4, on $\mathcal{E}_N = \{\|\hat{\mathbf{J}} - \mathbf{J}\|_{\text{op}} \leq \lambda_{\min}(\mathbf{J})/2\}$:

$$\|\hat{\boldsymbol{\theta}}_\gamma - \boldsymbol{\theta}^*\|_2 \leq \frac{C \varepsilon_\gamma}{\lambda_{\min}(\mathbf{J})} + \mathcal{O}_P \left(\frac{N^{-1/2} + r_N}{\lambda_{\min}(\mathbf{J})} \right). \quad (19)$$

Interpretation. $C \propto B_T(B_Y + B_T M_\theta)$ carries units

Table 1: The Latent-Confounding Difficulty Frontier: 100 Monte Carlo Replications per Grid Point (500 total evaluation runs across $\Delta_Z \in [0.2, 4.0]$, $N = 1,200$, True $\theta_{\text{ATE}}^* = 1.5278$). Nominal coverage target is 95.0% for Sample ATE (SATE) and Population ATE (PATE). Both mean and median absolute bias are reported to disclose finite-sample failure tails honestly. Coverage for hard-clustering Regime FE DML is computed via its post-clustering HAC sandwich covariance (exhibiting 0.0% coverage under weak separation $\Delta_Z \leq 1.0$ due to proxy-induced size distortion, Theorem 5). Full 11-method evaluation across all 5 grid points is tabulated in `reports/or_dml_benchmark_summary.csv`.

Δ_Z	Estimator	Mean Abs Bias	Median Abs Bias	RMSE	95% Cov (S/P)	$\lambda_{\min}(\mathbf{J})$
0.2 (Severe Overlap)	Oracle DML	0.0558	0.0440	0.0703	96.0% / 94.0%	0.4485
	Standard DML	8.4556	8.4559	8.4562	0.0% / 0.0%	4.9339
	DML + Z Controls (Ridge)	8.3072	8.3120	8.3079	0.0% / 0.0%	4.6273
	DML + Z Controls (GBM)	8.0178	8.0252	8.0190	0.0% / 0.0%	5.3254
	Regime FE DML (Hard)	8.3183	8.3238	8.3192	0.0% / 0.0%	1.7452
	Spectral OR-DML (Ours)	8.6266	8.6037	8.6276	0.0% / 0.0%	0.8044
1.0 (Moderate)	Oracle DML	0.0572	0.0431	0.0735	92.0% / 92.0%	0.4515
	Standard DML	8.4331	8.4494	8.4340	0.0% / 0.0%	4.9535
	DML + Z Controls (Ridge)	5.4379	5.4619	5.4400	0.0% / 0.0%	2.0779
	DML + Z Controls (GBM)	3.6334	3.6420	3.6465	0.0% / 0.0%	1.8845
	Regime FE DML (Hard)	2.6472	2.7005	2.6849	0.0% / 0.0%	0.6090
	Spectral OR-DML (Ours)	4.0170	4.0216	4.0367	0.0% / 0.0%	0.5570
2.0 (Clean Separation)	Oracle DML	0.0521	0.0421	0.0658	93.0% / 97.0%	0.4536
	Standard DML	8.4318	8.4508	8.4327	0.0% / 0.0%	4.9539
	DML + Z Controls (Ridge)	2.6307	2.6143	2.6345	0.0% / 0.0%	1.3505
	DML + Z Controls (GBM)	0.2066	0.1936	0.2426	72.0% / 72.0%	1.2144
	Regime FE DML (Hard)	0.0774	0.0630	0.1023	94.0% / 99.0%	0.4553
	Spectral OR-DML (Ours)	0.0812	0.0720	0.1039	92.0% / 97.0%	0.4546
4.0 (Isolated Regimes)	Filtered OR-DML (Ours)	0.1015	0.0812	0.1306	92.0% / 94.0%	0.4556
	Oracle DML	0.0540	0.0477	0.0670	96.0% / 95.0%	0.4394
	Standard DML	8.4509	8.4611	8.4518	0.0% / 0.0%	4.9123
	DML + Z Controls (Ridge)	0.8479	0.8415	0.8534	0.0% / 0.0%	1.0945
	DML + Z Controls (GBM)	0.1006	0.0817	0.1248	89.0% / 89.0%	1.1889
	Regime FE DML (Hard)	0.0540	0.0477	0.0670	96.0% / 95.0%	0.4394
	Spectral OR-DML (Ours)	0.0538	0.0462	0.0671	96.0% / 95.0%	0.4394
	Filtered OR-DML (Ours)	0.0538	0.0462	0.0671	96.0% / 95.0%	0.4394

of $T \cdot Y$; the dimensionless difficulty is $C\varepsilon_\gamma/\lambda_{\min}(\mathbf{J})$. When $\lambda_{\min}(\mathbf{J}) \approx 0$, proxy errors are catastrophically amplified.

Theorem 3 replaces heuristic claims with a rigorous bound. The cross-moment is Neyman-orthogonal to (μ_k, m_k) , but the coupled Gram part is non-orthogonal for soft assignments (Lemma 6); latent proxy uncertainty carries an irremediable price. In Appendix C, Theorem 7 proves a matching minimax lower bound $\inf_{\hat{\theta}} \sup_P \mathbb{E} \|\hat{\theta} - \theta^*\|_2 \geq c\varepsilon_\gamma/\lambda_{\min}(\mathbf{J})$, establishing that this degradation rate is tight and unimprovable.

The Difficulty Frontier and Operational Surrogate Bridge. Because true states are unobserved, we bridge unobservable $\mathcal{D}_{\text{causal}} \equiv \varepsilon_\gamma/\lambda_{\min}(\mathbf{J})$ to computable $\mathcal{D}_{\text{operational}}(t) \equiv H(\gamma_t)/\lambda_{\min}(\mathbf{J}_t)$ via the following result:

Assumption 5 (Two-Sided Spectral Comparability). There exist constants $0 < \underline{c}_J \leq \bar{c}_J < \infty$ such that rolling Gram matrices $\mathbf{J}_t$ satisfy $\underline{c}_J \leq \frac{\lambda_{\min}(\mathbf{J}_t)}{\lambda_{\min}(\mathbf{J})} \leq \bar{c}_J$ almost surely for all evaluation windows $t \geq w$.

Proposition 1 (Surrogate Entropy-Difficulty Bridge). Under Bayesian calibration in binary regimes ($K = 2$), forward-filtered posteriors γ_t satisfy $\mathbb{E}[\|\gamma_t - \mathbf{H}_t\|_1 \mid \mathcal{F}_t] = 2 \text{Gini}(\gamma_t) \leq C_K H(\gamma_t)$ with $C_K \leq 1/\ln 2 \approx 1.4427$. Defining $\mathcal{D}_{\text{causal}} \equiv \frac{\varepsilon_\gamma}{\lambda_{\min}(\mathbf{J})}$, $\mathcal{D}_{\text{global}}^H(t) \equiv \frac{H(\gamma_t)}{\lambda_{\min}(\mathbf{J})}$, and $\mathcal{D}_{\text{operational}}(t) \equiv \frac{H(\gamma_t)}{\lambda_{\min}(\mathbf{J}_t)}$: (i) Population difficulty is bounded by global entropy: $\mathcal{D}_{\text{causal}} \leq C_K \mathbb{E}[\mathcal{D}_{\text{global}}^H(t)]$. (ii) Under Assumption 5, $\mathcal{D}_{\text{causal}} \leq$

$C_K \bar{c}_J \mathbb{E}[\mathcal{D}_{\text{operational}}(t)] \leq \frac{C_K \bar{c}_J}{\underline{c}_J} \mathbb{E}[\mathcal{D}_{\text{global}}^H(t)]$. In simulation where H_t is observed, $\hat{\mathcal{D}}_{\text{causal}, N} \equiv \hat{\varepsilon}_{\gamma, N}/\lambda_{\min}(\mathbf{J})$ converges to $\mathcal{D}_{\text{causal}}$ at $\mathcal{O}_P(N^{-1/2})$.

Interpretation & Regularization. Proposition 1 grounds deployable monitoring by turning an unobservable theoretical difficulty into an observable surrogate upper bound. When $\lambda_{\min}(\mathbf{J}) \rightarrow 0$, a spectral penalty $\lambda > 0$ bounds the condition number, trading finite-sample bias for bounded variance.

Theorem 4 (Risk Bound for Spectrally Regularized OR-DML). Let $\lambda \geq 0$ with $\lambda_{\min}(\mathbf{J}) + \lambda > 0$, $\hat{\theta}_\lambda = (\hat{\mathbf{J}} + \lambda \mathbf{I})^{-1} \hat{\mathbf{S}}$, and $\theta_\lambda = (\mathbf{J} + \lambda \mathbf{I})^{-1} \mathbf{S}$. Under Theorem 3: (a) Deterministic bias satisfies $B_\lambda \equiv \|\theta_\lambda - \theta^*\|_2 \leq \frac{C\varepsilon_\gamma + \lambda M_\theta}{\lambda_{\min}(\mathbf{J}) + \lambda}$. (b) If Theorem 5(a) holds and $N \|\hat{\theta}_\lambda - \theta_\lambda\|_2^2$ is uniformly integrable, variance satisfies $V_N \equiv \mathbb{E} \|\hat{\theta}_\lambda - \theta_\lambda\|_2^2 \leq \frac{\text{Tr}(\mathbf{\Omega}_\lambda^{\text{tot}})}{N(\lambda_{\min}(\mathbf{J}) + \lambda)^2} + o(N^{-1})$, and by Minkowski's inequality:

$$\mathbb{E} \|\hat{\theta}_\lambda - \theta^*\|_2^2 \leq (B_\lambda + \sqrt{V_N})^2 \leq 2B_\lambda^2 + 2V_N. \quad (20)$$

Empirical surrogate minimizer $\hat{\lambda}_{\text{emp}}^* = \arg \min_{\lambda \in [0, \bar{\lambda}]} \hat{R}(\lambda)$ ($\bar{\lambda} < \infty$) satisfies an oracle inequality (Appendix D.1; $\lambda = 0.10$ achieves $23.0 \times$ risk reduction in benchmark simulations with 85.0% coverage).

Interpretation. Regularization provides finite-sample variance control under weak overlap, but cannot manufacture causal information absent from the proxy distribution.




Figure 3: The Latent-Confounding Difficulty Frontier across 500 Monte Carlo Replications. (a) Absolute causal bias $|\hat{\theta} - \theta^*|$ vs. regime separation Δ_Z . (b) Empirical RMSE. (c) Empirical 95% CI coverage vs. nominal 95% target. (d) Corrected factorial mechanism test: causal L_2 error vs. $\varepsilon_\gamma/\lambda_{\min}(J)$ across 49 independently controlled cells (100 replications each), where task conditioning is varied through residual treatment innovation variance after regime-aware nuisance residualization.

Theorem 5 (Asymptotic Distribution Under Dependent Purged Cross-Fitting). *Under Assumptions 1-4, 2, 6 with embargo buffer $\tau^* \geq C \log N$ ($C > 1/c_\alpha$):*

- (a) $\sqrt{N}(\hat{\theta}_\lambda - \theta_\lambda) \xrightarrow{d} \mathcal{N}(\mathbf{0}, \Sigma_\lambda^{\text{tot}})$, where $\Sigma_\lambda^{\text{tot}} = (J + \lambda I)^{-1} \Omega_\lambda^{\text{tot}} (J + \lambda I)^{-1}$, $\Omega_\lambda^{\text{tot}} = \Omega_\lambda + \kappa \mathbf{A}_\Lambda \mathcal{I}_\Lambda^{-1} \mathbf{A}_\Lambda^\top$, and $\mathbf{A}_\Lambda = \partial_\Lambda \mathbb{E}[\Psi_t(\theta_\lambda)]|_{\Lambda^*}$.
- (b) Because $\theta_\lambda = \theta^* + \text{Bias}(\lambda, \gamma)$ with $\text{Bias} = (J + \lambda I)^{-1}(\mathbf{b}_\gamma - \lambda \theta^*)$ and $\|\mathbf{b}_\gamma\|_2 \leq C\varepsilon_\gamma$, this is equivalent to $\sqrt{N}(\hat{\theta}_\lambda - \theta^* - \text{Bias}) \xrightarrow{d} \mathcal{N}(\mathbf{0}, \Sigma_\lambda^{\text{tot}})$.
- (c) Centered normality $\sqrt{N}(\hat{\theta}_\lambda - \theta^*) \xrightarrow{d} \mathcal{N}(\mathbf{0}, \Sigma_\lambda^{\text{tot}})$ holds if $\sqrt{N}(C\varepsilon_\gamma + \lambda M_\theta)/(\lambda_{\min}(J) + \lambda) \rightarrow 0$ (e.g., $\varepsilon_\gamma = o(N^{-1/2})$, $\lambda = o(N^{-1/2})$). When $\varepsilon_\gamma = O(1)$, intervals cover $\theta_\lambda \neq \theta^*$, explaining coverage breakdown under weak separation (Tables 1, 3). Estimated HMM parameters contribute $\kappa \mathbf{A}_\Lambda \mathcal{I}_\Lambda^{-1} \mathbf{A}_\Lambda^\top$ to the stacked HAC sandwich (Appendix E).

Proposition 2 (Joint Asymptotic Distribution for Sample and Population ATE). *The sample ATE $\hat{\theta}_{\text{SATE}} = \bar{\gamma}^\top \hat{\theta}_\lambda$ has conditional variance $\sigma_{\text{SATE}}^2 = \bar{\gamma}^\top \Sigma_\lambda \bar{\gamma}$, where Markov persistence ($\rho \approx 0.88$) inflates the long-run variance of sample state occupation $\bar{\gamma}_N$ by $\frac{1+\rho}{1-\rho} \approx 15.7\times$ relative to i.i.d. transitions. The population ATE $\hat{\theta}_{\text{PATE}} = \hat{\pi}(\hat{\Lambda})^\top \hat{\theta}_\lambda$ satisfies the joint delta expansion: $\sqrt{N}(\hat{\theta}_{\text{PATE}} - \theta_{\text{PATE}}^*) = \pi^{*\top} \sqrt{N}(\hat{\theta}_\lambda - \theta^*) + \theta^{*\top} \sqrt{N}(\hat{\pi} - \pi^*) + o_P(1)$, yielding asymptotic variance:*

$$\sigma_{\text{PATE}}^2 = \pi^{*\top} \Sigma_\lambda \pi^* + \theta^{*\top} \Sigma_\pi \theta^* + 2\pi^{*\top} \Sigma_{\theta\pi} \theta^*, \quad (21)$$

where $\Sigma_\pi = \nabla_\Lambda \pi^* \mathcal{I}_\Lambda^{-1} (\nabla_\Lambda \pi^*)^\top$ represents the parametric variance of the estimated stationary distribution from HMM influence function $\phi_{\Lambda,t}$, and $\Sigma_{\theta\pi} = \text{LRCov}(\phi_{\theta,t}, \phi_{\Lambda,t})$, consistently estimated via stacked joint HAC sandwich.

6 THE LATENT-CONFOUNDING DIFFICULTY FRONTIER

To test the theoretical predictions of Theorems 2-4, we conduct a comprehensive Monte Carlo evaluation across 500 evaluation runs (100 replications per grid point),

charting causal estimation error across the **Latent-Confounding Difficulty Frontier**.

Experimental Protocol. DGPs & Baselines: 100 Monte Carlo draws across $\Delta_Z \in \{0.2, 0.5, 1.0, 2.0, 4.0\}$ ($N = 1, 200$, $\theta_{\text{ATE}}^* = 1.5278$) comparing Oracle, Standard DML, Block DML, continuous proxy controls (Ridge/GBM), hard-clustering Regime FE, and OR-DML (Spectral/Filtered). **Representations & Shifts:** HMM, GRU, Transformer, and SSM under clean and compound shifts. **Reliability Target:** Predicting squared loss via difficulty $\mathcal{D}_t$ across 5 seeds with MBB. **Real-World Test:** $> 14,000$ hourly observations across four megacities.

Difficulty Frontier Findings. Table 1 and Figure 3 reveal three phenomena:

1. **Low-Observability Breakdown ($\Delta_Z \leq 1.0$):** Standard/Block DML exhibit constant omitted-regime bias (+8.43-8.49, RMSE 8.45, 0.0% coverage; Theorem 1). Under severe overlap ($\Delta_Z \leq 0.5$), all proxy estimators suffer heavy omitted-regime bias ($\hat{\varepsilon}_\gamma = 0.83$, bias 6.83-8.60). At $\Delta_Z = 1.0$, weak proxy separation leaves non-trivial residual confounding (≥ 2.65 : Decoupled Soft DML 3.29, Hard FE 2.65, Spectral OR-DML 4.02, while conditioning on treatment HMM(Z, T) achieves 1.20). Because finite-sample separation cannot eliminate proxy bias, high $\mathcal{D}_t$ acts as an operational *abstention trigger*.
2. **High-Observability Resolution ($\Delta_Z \geq 2.0$):** Clean separation resolves omitted confounding: with multi-restart EM ($n_{\text{init}} = 5$), mean absolute bias collapses to oracle levels (0.081 at $\Delta_Z = 2.0$ and 0.054 at $\Delta_Z = 4.0$), eliminating single-init EM local traps (where 2-5% of single-init runs pulled mean bias to 0.501 and RMSE to 1.905). At $\Delta_Z = 4.0$, regime estimators reach bias ≈ 0.054 with 96.0% coverage matching Oracle DML, collapsing to 0.0% under weak separation ($\Delta_Z \leq 1.0$) from proxy-induced size distortion: in an explicit null test ($\theta_1^* = 0$, Frontier B; Table 3), nominal 5% sandwich tests reject in 78.0% of runs (bias 0.342),

Table 2: Empirical Regime-Specific Exposure Effects across Four Indian Megacities ($> 14,000$ 2-Hourly Observations, Feb 2025-Jun 2026, Median Cadence 2.0h, Treatment $T_t = \text{NO}_2$, Outcome $Y_t = \text{PM}_{2.5}$). Evaluated via Gaussian HMM on meteorology († Mumbai displays 3 local optima: $R_1 \in [-13.34, -0.02]$, $R_2 \in [+3.45, +7.39]$, $p \in [0.006, 0.131]$; table reports global mode $\text{LL} = -55,557.7$). Diagnostics in Appendix G.

City	Estimator	Regime	Effect $\hat{\theta}$	Std Error	95% CI	p-value	$\lambda_{\min}(\mathbf{J})$ [κ]
Delhi ($N = 4,640$)	Spectral OR-DML	Regime 1	+0.7038	0.0513	[+0.6032, +0.8044]	< 0.0001	111.8 [1.58]
	Spectral OR-DML	Regime 2	+0.1281	0.0743	[-0.0174, +0.2737]	0.0845	111.8 [1.58]
Mumbai ($N = 4,180$) †	Spectral OR-DML	Regime 1	-7.6278	7.0540	[-21.4537, +6.1981]	0.2795	0.20 [2.52]
	Spectral OR-DML	Regime 2	+7.2007	3.9025	[-0.4482, +14.8496]	0.0650	0.20 [2.52]
Bengaluru ($N = 4,093$)	Spectral OR-DML	Regime 1	+0.0919	0.1552	[-0.2123, +0.3960]	0.5539	4.22 [1.72]
	Spectral OR-DML	Regime 2	-0.1528	0.0689	[-0.2878, -0.0178]	0.0265	4.22 [1.72]
Kolkata ($N = 1,209$)	Spectral OR-DML	Regime 1	+0.6607	1.4074	[-2.0977, +3.4191]	0.6388	0.0025 [1.43]
	Spectral OR-DML	Overall ATE	-0.0063	1.3003	[-2.5549, +2.5423]	0.9961	0.0025 [1.43]

demonstrating that without selective abstention, uncalibrated regime Wald inference is invalid.

3. **Filtering vs. Smoothing:** Forward filtering tracks retrospective smoothing closely across separated regimes ($\Delta_Z = 2.0$: bias 0.1015 vs. 0.0812; $\Delta_Z = 4.0$: identical 0.0538), incurring minimal loss once informed by dynamics.

Factorial Mechanism Validation. Independently manipulating both quantities in Theorem 3 confirms their sharp interaction: across 4,900 runs, difficulty ratio $\mathcal{D}_{\text{causal}} = \varepsilon_\gamma / \lambda_{\min}(\mathbf{J})$ ranks causal error with Spearman $\rho = 0.995$ vs. 0.626 (proxy error) and 0.521 ($1/\lambda_{\min}$) alone (cell-mean $\rho = 0.996$). With perfect states, mean error rises from 0.050 $\rightarrow$ 0.545 as residual information weakens; at $\varepsilon_\gamma = 0.01$, it amplifies from 0.686 $\rightarrow$ 131.3, confirming conditioning shifts survive regime-aware residualization.

Representation Quality vs. Downstream Reliability. State discrimination, calibration, and task reliability are distinct. Across 15 shifted worlds, adding NLL to F_1 improves leave-one-world-out reliability prediction from $R^2 = 0.113$ ($\rho = 0.374$) to $R^2 = 0.221$ ($\rho = 0.527$). Yet temperature scaling cuts GRU/Transformer ECE by 47.9%/60.3% and NLL from 0.548/1.006 to 0.331/0.382 without improving failure discrimination ($\Delta\text{AUC} = -0.004/-0.008$), proving calibration cannot rescue downstream reliability under ill-conditioned task geometry (Appendix H).

7 EMPIRICAL SENSOR NETWORK STRESS TEST

We evaluate OR-DML on our verified multi-season dataset (14,122 hourly records across Delhi, Mumbai, Bengaluru, and Kolkata, Feb 2025-Jun 2026; $< 0.7\%$ forward-fill; Table 2), treating the sensor study as an observational diagnostic stress test rather than physical causal ground truth. Complete-case and Newey-West HAC lag sensitivity ($h \in \{12, 24, 72, 168\}$ h) appear in Appendix G.

Conditioning Diagnostics. *Delhi Basin* ($\lambda_{\min} = 111.8$, $\kappa = 1.58$): Decomposes pooled effect ($+0.47 \pm 0.04$) into Regime 1 ($+0.70 \pm 0.05$, inversion) and Regime 2 ($+0.13 \pm 0.07$, clearance). *Mumbai Coastal* ($\lambda_{\min} = 0.20$, $\kappa = 2.52$): Standard DML is unstable ($+2.64 \pm 4.42$), while Spectral OR-DML yields Regime 2 $+7.20 \pm 3.90$ ($p = 0.065$) and Regime 1 -7.63 ± 7.05 ($p = 0.280$; overall $+1.08 \pm 3.91$). Telemetry screen-

ing reveals a 273-hour constant run and $985 \mu\text{g}/\text{m}^3$ saturation, causing multi-modal EM likelihood surfaces (across local modes, $R_1 \in [-13.34, -0.02]$, $R_2 \in [+3.45, +7.39]$, $p \in [0.006, 0.131]$; Table 2). *Bengaluru Plateau* ($\lambda_{\min} = 4.22$, $\kappa = 1.72$): Standard DML yields -0.001 ± 0.060 ; OR-DML isolates Regime 1 ($+0.092 \pm 0.155$) from Regime 2 (-0.153 ± 0.069 , nominal $p = 0.0265$, overall -0.026 ± 0.079). Under Frontier B null size distortion (78% Type I error, Table 3), nominal regime p -values require transport modeling before causal interpretation. *Kolkata Airshed* ($\lambda_{\min} = 0.0025$, $\kappa = 1.43$): Sparse telemetry ($N = 1,209$) and discrete quantization (31 unique NO_2 values; 376-hour constant $\text{PM}_{2.5}$ run) collapse residual variance; overall ATE is -0.006 ± 1.300 . In single-index sensitivity, unregularized DML yields -0.557 ± 1.107 , stabilized by spectral penalty ($\text{SE} \rightarrow 0.060$ at $\lambda = 0.10$, Table 5). Modest condition numbers ($\kappa = 1.43$ -2.52) confirm SE inflation stems from low residual variance and sensor defects rather than Gram collinearity. *Semi-Synthetic Megacity Benchmark:* In real atmospheric covariates with known ground truth ($\theta^* = (0.75, 2.50)$; Table 7), all proxy estimators suffer bias (3.72-5.27 for Spectral OR-DML, 1.09-2.56 for HMM(Z, T)) and 0.0% coverage, confirming real sensor confounding invalidates naive causal inference. *Dynamic Projections:* Exposure effect peaks at $h = 2\text{h}$ in Delhi (ATE $+1.31$, Regime 1 $+2.35$, decaying to $+0.28$ at 12h), but Delhi fails pre-treatment placebos at $h \in \{-6, -3, -1\}$ ($p < 0.01$). We treat this curve as an exposure diagnostic; other cities pass placebos (Appendix G).

8 DISCUSSION, LIMITATIONS & CONCLUSION

Latent-state recovery and downstream reliability are distinct: proxy errors are benign under strong geometry but catastrophic when ill-conditioned, motivating joint diagnostics over calibration alone. **Limitations:** (1) Proposition 1 assumes calibration and spectral comparability. (2) Sensor estimates are exposure tests; Delhi placebos preclude unmodeled dynamic claims. (3) Operational difficulty does not dominate entropy everywhere. (4) Financial transfer is synthetic.

AI Use Statement

We used large language models (LLMs) to assist in mathematical formulation, proof derivation drafting (Appendices A-I), codebase development (`src/or_dml.py`, `src/synthetic_dgp_benchmark.py`, `src/factorial_reliability_experiment.py`, `src/empirical_evaluation.py`), bibliographic formatting, and manuscript drafting. AI tools were not used for autonomous data generation or unverified empirical interpretation. The authors independently reviewed and verified the mathematical arguments and empirical implementations against simulation baselines with known ground truth and reproducible artifacts in `reports/`. The authors take full responsibility for all content, claims, and conclusions.

References

- Elizabeth S Allman, Catherine Matias, and John A Rhodes. Identifiability of parameters in latent structure models with independent or Markov conditions. *The Annals of Statistics*, 37(6A):3099-3132, 2009.
- Donald W K Andrews. Heteroskedasticity and autocorrelation consistent covariance matrix estimation. *Econometrica*, 59(3):817-858, 1991.
- Susan Athey, Julie Tibshirani, and Stefan Wager. Generalized random forests. *The Annals of Statistics*, 47(2):1148-1178, 2019.
- Jushan Bai. Panel data models with interactive fixed effects. *Econometrica*, 77(4):1229-1279, 2009.
- Marwa Banna, Florence Merlevède, and Pierre Youssef. Bernstein-type inequality for a class of dependent random matrices. *Random Matrices: Theory and Applications*, 5(2):1650006, 2016. arXiv:1504.05834.
- Peter L Bartlett and Marten H Wegkamp. Classification with a reject option using a hinge loss. *Journal of Machine Learning Research*, 9:1823-1840, 2008.
- Leonard E Baum and Ted Petrie. Statistical inference for probabilistic functions of finite state Markov chains. *The Annals of Mathematical Statistics*, 37(6):1554-1563, 1966.
- Christoph Bergmeir, Rob J Hyndman, and Bonsoo Koo. A note on the validity of cross-validation for evaluating autoregressive time series prediction. *Computational Statistics & Data Analysis*, 120: 70-83, 2018.
- Ioana Bica, Ahmed M Alaa, and Mihaela van der Schaar. Time series deconfounder: Estimating treatment effects over time in the presence of hidden confounders. In *International Conference on Machine Learning (ICML)*, volume 119 of *Proceedings of Machine Learning Research*, pages 884-895, 2020.
- Peter J Bickel, Chris AJ Klaassen, Ya'acov Ritov, and Jon A Wellner. *Efficient and adaptive estimation for semiparametric models*. Johns Hopkins University Press, 1993.
- Peter J Bickel, Ya'acov Ritov, and Tobias Rydén. Asymptotic normality of the maximum-likelihood estimator for general hidden Markov models. *The Annals of Statistics*, 26(4):1614-1635, 1998.
- Iavor Bojinov and Neil Shephard. Time series experiments and causal estimands: Exact randomization tests and trading. *Journal of the American Statistical Association*, 114(528):1665-1682, 2019.
- Stéphane Bonhomme, Thibaut Lamadon, and Elena Manresa. Discretizing unobserved heterogeneity. *Econometrica*, 90(2):625-643, 2022.
- Stéphane Bonhomme and Elena Manresa. Grouped patterns of heterogeneity in panel data. *Econometrica*, 83(3):1147-1184, 2015.
- Richard C Bradley. Basic properties of strong mixing conditions. A survey and some open questions. *Probability Surveys*, 2:107-144, 2005.
- Prabir Burman, Edmond Chow, and Deborah Nolan. A cross-validatory method for dependent data. *Biometrika*, 81(2):351-358, 1994.
- Olivier Cappé, Eric Moulines, and Tobias Rydén. *Inference in Hidden Markov Models*. Springer, New York, 2005.
- Victor Chernozhukov, Denis Chetverikov, Mert Demirer, Esther Duflo, Christian Hansen, Whitney Newey, and James Robins. Double/debiased machine learning for treatment and structural parameters. *The Econometrics Journal*, 21(1):C1-C68, 2018.
- Victor Chernozhukov, Juan Carlos Escanciano, Hidehiko Ichimura, Whitney K Newey, and James M Robins. Locally robust semi-parametric estimation. *Econometrica*, 90(4):1501-1535, 2022a.
- Victor Chernozhukov, Whitney K Newey, and Rahul Singh. Automatic debiased machine learning of causal and structural effects. *Econometrica*, 90(3):967-1027, 2022b.
- C K Chow. On optimum recognition error and reject tradeoff. *IEEE Transactions on Information Theory*, 16(1):41-46, 1970.
- Milos Ciganovic, Federico D'Amario, and Massimiliano Tancioni. Double machine learning for time series. *The Econometrics Journal*, 2026. doi: 10.1093/ectj/utag019. arXiv:2603.10999.
- Carlos Cinielli and Chad Hazlett. Making sense of sensitivity: Extending omitted variable bias. *Journal of the Royal Statistical Society: Series B*, 82(1):39-67, 2020.
- Felix J Clouth, Maarten J Bijlsma, Steffen Pauws, and Jeroen K Vermunt. Causal inference for latent Markov models using the parametric G-formula. *Sociological Methods & Research*, 2025. doi: 10.1177/00491241251377068.
- Corinna Cortes, Giulia DeSalvo, and Mehryar Mohri. Learning with rejection. In *Algorithmic Learning Theory (ALT)*, volume 9925 of *Lecture Notes in Computer Science*, pages 67-82. Springer, 2016.
- Central Pollution Control Board (CPCB). National ambient air quality monitoring program. Ministry of Environment, Forest and Climate Change, Government of India, 2025-2026.
- Richard K Crump, V Joseph Hotz, Guido W Imbens, and Jack K Mitnik. Dealing with limited overlap in estimation of average treatment effects. *Biometrika*, 96(1):187-199, 2009.
- Yifan Cui, Hongming Pu, Xu Shi, Wang Miao, and Eric J Tchetgen Tchetgen. Semiparametric proximal causal inference. *Journal of the American Statistical Association*, 119(546):1348-1359, 2024.
- Alexander D'Amour. On multi-cause causal inference with unobserved confounding: Counterexamples, impossibility, and alternatives. In *International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 89 of *Proceedings of Machine Learning Research*, 2019.
- Alexander D'Amour, Peng Ding, Avi Feller, Lihua Lei, and Jasjeet Sekhon. Overlap in observational studies with high-dimensional covariates. *Journal of Econometrics*, 221(2):644-654, 2021. arXiv:1711.02582.
- Arthur P Dempster, Nan M Laird, and Donald B Rubin. Maximum likelihood from incomplete data via the EM algorithm. *Journal of the Royal Statistical Society: Series B*, 39(1):1-38, 1977.
- Francesca Dominici, Michael Greenstone, and Cass R Sunstein. Particulate matter matters. *Science*, 344(6181):257-259, 2014.
- Randal Douc, Eric Moulines, and Tobias Rydén. Asymptotic properties of the maximum likelihood estimator in autoregressive models with Markov regime. *The Annals of Statistics*, 32(5):2254-2304, 2004.
- Ran El-Yaniv and Yair Wiener. On the foundations of noise-free selective classification. *Journal of Machine Learning Research*, 11:1605-1641, 2010.
- Dylan J Foster and Vasilis Syrgkanis. Orthogonal statistical learning. *The Annals of Statistics*, 51(3):879-908, 2023.
- Ragnar Frisch and Frederick V Waugh. Partial time regressions as compared with individual trends. *Econometrica*, 1(4):387-401, 1933.
- Élisabeth Gassiat, Alice Cleynen, and Stéphane Robin. Inference in finite state space non parametric hidden Markov models and applications. *Statistics and Computing*, 26(1-2):61-71, 2016.
- Yonatan Geifman and Ran El-Yaniv. Selective classification for deep neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, pages 4878-4887, 2017.
- Yonatan Geifman and Ran El-Yaniv. SelectiveNet: A deep neural network with an integrated reject option. In *International Conference on Machine Learning (ICML)*, volume 97 of *Proceedings of Machine Learning Research*, pages 2151-2159, 2019.
- Stuart K Grange and David C Carslaw. Using meteorological normalisation to detect interventions in air quality time series. *Science of the Total Environment*, 653:578-588, 2019.
- Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q Weinberger. On calibration of modern neural networks. In *International Conference on Machine Learning (ICML)*, pages 1321-1330, 2017.
- James D Hamilton. A new approach to the economic analysis of

- nonstationary time series and the business cycle. *Econometrica*, 57(2):357-384, 1989.
- James D Hamilton. Analysis of time series subject to changes in regime. *Journal of Econometrics*, 45(1-2):39-70, 1990.
- Tobias Hatt and Stefan Feuerriegel. Sequential deconfounder for estimating treatment effects in the presence of unobserved confounders. In *International Conference on Machine Learning (ICML)*, volume 139 of *Proceedings of Machine Learning Research*, 2021.
- Kilian Hendrickx, Lorenzo Perini, Dries Van der Plas, Wannes Meert, and Jesse Davis. Machine learning with a reject option: A survey. *Machine Learning*, 113(5):3073-3110, 2024. arXiv:2107.11277.
- Arthur E Hoerl and Robert W Kennard. Ridge regression: Biased estimation for nonorthogonal problems. *Technometrics*, 12(1):55-67, 1970.
- Roger A Horn and Charles R Johnson. *Matrix Analysis*. Cambridge University Press, 2nd edition, 2012.
- Biwei Huang, Kun Zhang, Jiji Zhang, Joseph Ramsey, Ruben Sanchez-Romero, Clark Glymour, and Bernhard Schölkopf. Causal discovery from heterogeneous/nonstationary data. *Journal of Machine Learning Research*, 21(89):1-53, 2020.
- Il'dar A Ibragimov. Some limit theorems for stationary processes. *Theory of Probability & Its Applications*, 7(4):349-382, 1962.
- Guido W Imbens and Donald B Rubin. *Causal Inference for Statistics, Social, and Biomedical Sciences*. Cambridge University Press, 2015.
- Nathan Kallus, Xiaojie Mao, and Angela Zhou. Interval estimation of individual-level causal effects under unobserved confounding. In *International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 89 of *Proceedings of Machine Learning Research*, 2019.
- John G Kemeny and J Laurie Snell. *Finite Markov Chains*. Springer-Verlag, 1976.
- Edward H Kennedy. Semiparametric doubly robust targeted double machine learning: a review. *arXiv preprint arXiv:2203.06469*, 2022.
- Shakeeb Khan and Elie Tamer. Irregular identification, support conditions, and inverse weight estimation. *Econometrica*, 78(6):2021-2042, 2010.
- Chang-Jin Kim and Charles R Nelson. *State-Space Models with Regime Switching: Classical and Gibbs-Sampling Approaches with Applications*. MIT Press, 1999.
- Hans R Künsch. The jackknife and the bootstrap for general stationary observations. *The Annals of Statistics*, 17(3):1217-1241, 1989.
- Manabu Kuroki and Judea Pearl. Measurement bias and effect restoration in causal inference. *Biometrika*, 101(2):423-437, 2014.
- Brian G Leroux. Maximum-likelihood estimation for hidden Markov models. *Stochastic Processes and their Applications*, 40(1):127-143, 1992.
- Greg Lewis and Vasilis Syrgkanis. Double/debiased machine learning for dynamic treatment effects. *arXiv preprint arXiv:2002.07285*, 2020.
- Christos Louizos, Uri Shalit, Joris M Mooij, David Sontag, Richard Zemel, and Max Welling. Causal effect inference with deep latent-variable models. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 30, 2017.
- Michael C Lovell. Seasonal adjustment of economic time series and multiple regression analysis. *Journal of the American Statistical Association*, 58(304):993-1010, 1963.
- Marcos López de Prado. *Advances in Financial Machine Learning*. John Wiley & Sons, 2018.
- Afsaneh Mastouri, Yuchen Zhu, Limor Gultchin, Anna Korba, Ricardo Silva, Matt J Kusner, Arthur Gretton, and Krikamol Muandet. Proximal causal learning with kernels: Two-stage estimation and moment restriction. In *International Conference on Machine Learning (ICML)*, volume 139 of *Proceedings of Machine Learning Research*, pages 7512-7523, 2021. arXiv:2105.04544.
- Wang Miao, Zhi Geng, and Eric J Tchetgen Tchetgen. Identifying causal effects with proxy variables of an unmeasured confounder. *Biometrika*, 105(4):987-993, 2018.
- Whitney K Newey and Kenneth D West. A simple, positive semi-definite, heteroskedasticity and autocorrelation consistent covariance matrix. *Econometrica*, 55(3):703-708, 1987.
- Xinkun Nie and Stefan Wager. Quasi-oracle estimation of heterogeneous treatment effects. *Biometrika*, 108(2):299-319, 2021.
- Elizabeth L Ogburn, Ilya Shpitser, and Eric J Tchetgen Tchetgen. Comment on “The blessings of multiple causes”. *Journal of the American Statistical Association*, 114(528):1611-1615, 2019.
- Open-Meteo. Historical weather API and ECMWF ERA5 meteorological reanalysis. Technical documentation and API service, 2025-2026.
- Judea Pearl. On a class of bias-amplifying variables that endanger effect estimates. In *Proceedings of the 26th Conference on Uncertainty in Artificial Intelligence (UAI)*, pages 417-424, 2010. arXiv:1203.3503.
- Maya L Petersen, Kristin E Porter, Susan Gruber, Yue Wang, and Mark J van der Laan. Diagnosing and responding to violations in the positivity assumption. *Statistical Methods in Medical Research*, 21(1):31-54, 2012.
- Lawrence R Rabiner. A tutorial on hidden Markov models and selected applications in speech recognition. *Proceedings of the IEEE*, 77(2):257-286, 1989.
- Jeff Racine. Consistent cross-validated model-selection for dependent data: hv-block cross-validation. *Journal of Econometrics*, 99(1):39-61, 2000.
- James M Robins, Miguel A Hernán, and Babette Brumback. Marginal structural models and causal inference in epidemiology. *Epidemiology*, 11(5):550-560, 2000.
- Peter M Robinson. Root-N-consistent semiparametric regression. *Econometrica*, 56(4):931-954, 1988.
- Paul R Rosenbaum and Donald B Rubin. The central role of the propensity score in observational studies for causal effects. *Biometrika*, 70(1):41-55, 1983.
- Vira Semenova, Matt Goldman, Victor Chernozhukov, and Matt Taddy. Inference on heterogeneous treatment effects in high-dimensional dynamic panels under weak dependence. *Quantitative Economics*, 14(2):471-510, 2023. arXiv:1712.09988.
- Uri Shalit, Fredrik D Johansson, and David Sontag. Estimating individual treatment effect: Generalization bounds and algorithms. In *International Conference on Machine Learning (ICML)*, volume 70 of *Proceedings of Machine Learning Research*, pages 3076-3085, 2017.
- Eric J Tchetgen Tchetgen, Andrew Ying, Yifan Cui, Xu Shi, and Wang Miao. An introduction to proximal causal learning. *arXiv preprint arXiv:2009.10835*, 2020.
- Joel A Tropp. User-friendly tail bounds for sums of random matrices. *Foundations of Computational Mathematics*, 12(4):389-434, 2012.
- Tuan V Vu, Zongbo Shi, Jing Cheng, Qiang Zhang, Kebin He, Shuxiao Wang, and Roy M Harrison. Assessing the impact of clean air action on air quality trends in Beijing using a machine learning technique. *Atmospheric Chemistry and Physics*, 19(17):11303-11314, 2019.
- Yixin Wang and David M Blei. The blessings of multiple causes. *Journal of the American Statistical Association*, 114(528):1574-1596, 2019.
- C F Jeff Wu. On the convergence properties of the EM algorithm. *The Annals of Statistics*, 11(1):95-103, 1983.
- Alexandre B Tsybakov. *Introduction to Nonparametric Estimation*. Springer Science & Business Media, 2009.

CHECKLIST

1. For all models and algorithms presented, check if you include:
 - (a) A clear description of the algorithm, mapping, or model. **[Yes]** See Section 3 and Algorithm 1.
 - (b) An analysis of the properties and complexity. **[Yes]** See Theorems 1-5 and Propositions 1-2.
 - (c) (Optional) Anonymized source code, with specification of all dependencies, including external libraries. **[Yes]** Code included in the supplementary material.
2. For any theoretical claim, check if you include:
 - (a) Statements of the full set of assumptions of all theoretical results. **[Yes]** Assumptions 1-4 are stated in Section 2 and Section 4.
 - (b) Complete proofs of all theoretical results. **[Yes]** Full proofs in Appendices A-F.
 - (c) Clear explanations of any assumptions. **[Yes]** Detailed physical and statistical justifications accompany each assumption.
3. For all figures and tables that present empirical results, check if you include:
 - (a) The code, data, and instructions needed to reproduce the main experimental results. **[Yes]** Scripts in `src/` and reproduction instructions in `README.md`.
 - (b) All the training details (e.g., data splits, hyperparameters, how they were selected). **[Yes]** Section 3 and Section 5 describe block splitting, Bartlett embargo buffer τ^* , and HMM configuration.
 - (c) A clear definition of the specific measure or statistics and error bars (e.g., confidence intervals, p-values). **[Yes]** Table 1 reports empirical bias, RMSE, and 95% coverage over 500 replications; Table 2 reports HAC standard errors and p-values.
 - (d) A description of the computing infrastructure used. **[Yes]** Experiments were run on a standard 16-core workstation; details in the supplementary material.
4. If you are using existing assets (e.g., code, data, models) or curating/releasing new assets:
 - (a) If your work uses existing assets, did you cite the creators? **[Yes]** CPCB ground stations (CPCB, 2025) and Open-Meteo reanalysis (Open-Meteo, 2025) are cited.
 - (b) Did you mention the license of the assets? **[Yes]** Open public government data terms noted.
 - (c) Did you include any new assets either in the supplemental material or as a URL? **[Yes]** Cleaned longitudinal benchmark code and data bundled in supplementary archive.
 - (d) Did you discuss whether and how consent was obtained from people whose data you're using/curating? **[Not Applicable]** Aggregate ambient physical air measurements only.
 - (e) Did you discuss whether the data you are using/curating contains personally identifiable information or offensive content? **[Not Applicable]** Contains only environmental measurements.
5. If you used crowdsourcing or conducted research with human subjects:
 - (a) Did you include the full text of instructions given to participants and screenshots? **[Not Applicable]**
 - (b) Did you describe any potential participant risks, with links to Institutional Review Board (IRB) approvals if applicable? **[Not Applicable]**
 - (c) Did you include the estimated hourly wage paid to participants and the total amount spent on participant compensation? **[Not Applicable]**

MATHEMATICAL APPENDIX: COMPLETE PROOFS AND DERIVATIONS

Supplementary Roadmap. This supplement provides complete mathematical proofs, implementation details, experimental protocols, sensitivity analyses, and multi-domain empirical evaluations supporting the claims in the main text:

- **Appendix A (Proof of Theorem 1):** Exact decomposition of omitted regime bias and Frisch-Waugh bias amplification.
- **Appendix B (Proof of Theorem 2):** Non-identification under unconstrained overlap via constructive observational equivalence with treatment confounding ($T \not\perp S$).
- **Appendix C (Proof of Theorem 3):** Graceful degradation bound $\|\theta_\gamma - \theta^*\|_2 \leq C\varepsilon_\gamma/\lambda_{\min}(\mathbf{J})$ and Surrogate Entropy Bridge (Proposition 1).
- **Appendix D (Proof of Theorem 4):** Spectrally regularized OR-DML risk bound and adaptive empirical penalty selection.
- **Appendix E & F (Proof of Theorem 5 & Proposition 2):** Dependent asymptotic normality under purged cross-fitting and joint Delta-method inference for sample and population ATE.
- **Appendix G (Empirical Sensitivity & Provenance):** Bartlett embargo diagnostics, HAC lag truncation sensitivity, missingness robustness, and megacity regime profiles.
- **Appendix H (LatentRegimeBench & Representation Zoo):** Corrected factorial mechanism validation, architectural comparisons (HMM, GRU, Transformer, SSM), calibration intervention, leave-one-world-out reliability prediction, and multi-horizon failure forecasting.
- **Appendix I (Multi-Domain Financial Stress Test):** Out-of-sample risk abstention under synthetic macroeconomic regimes with real covariate features.

A PROOF OF THEOREM 1: DECOMPOSITION OF OMITTED REGIME BIAS

Consider the Structural Causal Model (1)-(5):

$$\begin{aligned} T_t &= m(X_t, S_t) + V_t, \quad \mathbb{E}[V_t | X_t, S_t] = 0, \\ Y_t &= \sum_{k=1}^K \theta_k^* \mathbb{I}(S_t = k) T_t + g(X_t, S_t) + U_t, \quad \mathbb{E}[U_t | X_t, S_t, T_t] = 0. \end{aligned}$$

Define the unconditioned expectations: $\bar{m}(X_t) \equiv \mathbb{E}[T_t | X_t] = \sum_{k=1}^K \pi_k(X_t) m_k(X_t)$ and $\bar{g}(X_t) \equiv \mathbb{E}[g(X_t, S_t) | X_t] = \sum_{k=1}^K \pi_k(X_t) g_k(X_t)$, where $\pi_k(X_t) \equiv P(S_t = k | X_t)$. Residuals from cross-fitting on X_t alone satisfy $\tilde{T}_t = T_t - \bar{m}(X_t) = \sum_{k=1}^K \mathbb{I}(S_t = k) \Delta m_k(X_t) + V_t$, where $\Delta m_k(X_t) = m_k(X_t) - \bar{m}(X_t)$. The outcome residual is:

$$\tilde{Y}_t = Y_t - \mathbb{E}[Y_t | X_t] = \sum_{k=1}^K \theta_k^* \mathbb{I}(S_t = k) T_t - \sum_{k=1}^K \theta_k^* \pi_k(X_t) m_k(X_t) + g(X_t, S_t) - \bar{g}(X_t) + U_t. \quad (22)$$

Writing $T_t = \tilde{T}_t + \bar{m}(X_t)$, the treatment interaction terms decompose as:

$$\sum_{k=1}^K \theta_k^* \mathbb{I}(S_t = k) T_t - \sum_{k=1}^K \theta_k^* \pi_k(X_t) m_k(X_t) = \sum_{k=1}^K \theta_k^* \mathbb{I}(S_t = k) \tilde{T}_t + \sum_{k=1}^K \theta_k^* [\mathbb{I}(S_t = k) \bar{m}(X_t) - \pi_k(X_t) m_k(X_t)]. \quad (23)$$

Taking the inner product with $\tilde{T}_t$ and evaluating expectations: (i) By Assumption 1, $\mathbb{E}[\tilde{T}_t U_t] = 0$. (ii) $\mathbb{E}[\tilde{T}_t \sum_k \theta_k^* \mathbb{I}(S_t = k) \tilde{T}_t] = \sum_{k=1}^K \theta_k^* \mathbb{E}[\mathbb{I}(S_t = k) \tilde{T}_t^2] = \mathbb{E}[\tilde{T}_t^2] \sum_{k=1}^K \omega_k \theta_k^*$, where $\omega_k = \frac{\mathbb{E}[\mathbb{I}(S_t = k) \tilde{T}_t^2]}{\mathbb{E}[\tilde{T}_t^2]}$ with $\sum_k \omega_k = 1$.

(iii) Since $\mathbb{E}[\tilde{T}_t | X_t] = 0$ and $\mathbb{E}[\tilde{T}_t \mathbb{I}(S_t = k) | X_t] = \Delta m_k(X_t) \pi_k(X_t)$, the mean adjustment yields $\sum_{k=1}^K \theta_k^* \mathbb{E}[\bar{m}(X_t) \Delta m_k(X_t) \pi_k(X_t)]$. (iv) The confounding component evaluates to $\mathbb{E}[\tilde{T}_t (g(X_t, S_t) - \bar{g}(X_t))] = \sum_{k=1}^K \mathbb{E}[\pi_k(X_t) \Delta m_k(X_t) \Delta g_k(X_t)]$. Dividing by $\mathbb{E}[\tilde{T}_t^2]$ yields (11) and (12). ■

B PROOF OF THEOREM 2: NON-IDENTIFICATION UNDER UNCONSTRAINED OVERLAP

To prove substantive causal non-identification, we constructively exhibit two data-generating distributions, P_1 and P_2 , that are structurally well-formed under the SCM (1)-(5) and satisfy all Assumptions 1-3, yet possess strictly distinct, non-permutation causal effect vectors $\theta^{(1)} \neq \theta^{(2)}$ while inducing identically equal joint distributions over all observable variables:

$$P_1(Y_{1:N}, T_{1:N}, X_{1:N}, Z_{1:N}) \equiv P_2(Y_{1:N}, T_{1:N}, X_{1:N}, Z_{1:N}) \quad \forall N \geq 1. \quad (24)$$

Construction of Substantive Observational Equivalence under Latent Confounding ($T \not\perp S$): Consider a two-state Markov chain ($K = 2$) with symmetric memoryless transition probabilities $P_{01} = P_{10} = 0.5$, yielding independent stationary regimes $S_t \stackrel{\text{i.i.d.}}{\sim} \text{Bernoulli}(0.5)$ with $\pi_0^* = \pi_1^* = 0.5$. Let auxiliary meteorological proxies $Z_t \in \mathbb{R}$ satisfy complete overlap: $Z_t \mid S_t = 0 \sim \mathcal{N}(0, 1)$ and $Z_t \mid S_t = 1 \sim \mathcal{N}(0, 1)$ identically, so $Z_t \perp S_t$. Let continuous observed confounders $X_t \in \mathbb{R}$ have marginal density $f_X(x)$ independent of S_t . Let binary treatment assignment $T_t \in \{0, 1\}$ depend directly on the latent regime: $P(T_t = 1 \mid S_t = 0) = 0.75$ and $P(T_t = 1 \mid S_t = 1) = 0.25$, so that $T_t \not\perp S_t$, with propensity scores $e(X_t, S_t) \in [0.25, 0.75]$ strictly satisfying Assumption 2 (Overlap & State Positivity with $\eta = 0.25$). Marginally, $P(T_t = 1) = 0.5 \times 0.75 + 0.5 \times 0.25 = 0.5$, and by Bayes' rule, the posterior regime occupancy given treatment satisfies $P(S_t = 0 \mid T_t = 1) = 0.75$, $P(S_t = 1 \mid T_t = 1) = 0.25$, $P(S_t = 0 \mid T_t = 0) = 0.25$, and $P(S_t = 1 \mid T_t = 0) = 0.75$.

Now consider structural outcome equation $Y_t \in \{0, 1\}$ with $P(Y_t = 1 \mid T_t = t, S_t = s) = \theta_s t + g_s$ and unconfounded innovations $\mathbb{E}[U_t \mid X_t, S_t, T_t] = 0$:

1. **Model P_1 (Homogeneous Causal Effect):** Set $\boldsymbol{\theta}^{(1)} = (0.20, 0.20)^\top$ ($\theta_0^{(1)} - \theta_1^{(1)} = 0$), baselines $g_0^{(1)} = 0.30, g_1^{(1)} = 0.50$. Then $P_1(Y_t = 1 \mid T_t = 0) = 0.25(0.30) + 0.75(0.50) = 0.45$, and $P_1(Y_t = 1 \mid T_t = 1) = 0.75(0.20 + 0.30) + 0.25(0.20 + 0.50) = 0.75(0.50) + 0.25(0.70) = 0.55$.
2. **Model P_2 (Heterogeneous Causal Effect):** Set $\boldsymbol{\theta}^{(2)} = (0.10, -0.30)^\top$ ($\theta_0^{(2)} - \theta_1^{(2)} = +0.40$), baselines $g_0^{(2)} = 0.60, g_1^{(2)} = 0.40$. Then $P_2(Y_t = 1 \mid T_t = 0) = 0.25(0.60) + 0.75(0.40) = 0.45$, and $P_2(Y_t = 1 \mid T_t = 1) = 0.75(0.10 + 0.60) + 0.25(-0.30 + 0.40) = 0.75(0.70) + 0.25(0.10) = 0.525 + 0.025 = 0.55$.

Because S_t transitions independently ($p = 0.5$), observations (Y_t, T_t, X_t, Z_t) are mutually independent across time, yielding $P(Y_{1:N}, T_{1:N}, X_{1:N}, Z_{1:N}) = \prod_{t=1}^N P(Y_t, T_t, X_t, Z_t)$. Comparing single-step observable conditionals, $P_1(Y_t \mid T_t, X_t, Z_t) \equiv P_2(Y_t \mid T_t, X_t, Z_t)$ identically for all (y, t, x, z) , ensuring $P_1(Y_{1:N}, T_{1:N}, X_{1:N}, Z_{1:N}) \equiv P_2(Y_{1:N}, T_{1:N}, X_{1:N}, Z_{1:N})$ for all $N \geq 1$. Yet structural effects satisfy $\boldsymbol{\theta}^{(1)} = (0.20, 0.20)^\top \neq (0.10, -0.30)^\top = \boldsymbol{\theta}^{(2)}$, which is not a permutation. **Dynamic Identifiability:** When regimes exhibit non-trivial temporal dependence ($p \neq 0.5$), latent Markov models can achieve generic algebraic point-identification from sequence transitions of length $N \geq 3$ (Allman et al., 2009). However, in memoryless/cross-sectional regimes, or whenever proxy uncertainty leaves states unconstrained without dynamic modeling, structural non-identification holds identically for all $N \geq 1$. ■

C PROOF OF THEOREM 3: GRACEFUL DEGRADATION AND THE DIFFICULTY FRONTIER

Let $\mathbf{H}_t = \mathbf{e}_{S_t} \in \{0, 1\}^K$ denote the unobserved true regime indicator vector. Under SCM (1)-(5), the oracle regime nuisance functions are:

$$m_k^*(X_t) \equiv \mathbb{E}[T_t \mid X_t, S_t = k], \quad \mu_k^*(X_t) \equiv \mathbb{E}[Y_t \mid X_t, S_t = k] = \theta_k^* m_k^*(X_t) + g(X_t, k). \quad (25)$$

Step 1: Define proxy-weighted targets and population estimand. The empirical cross-fitted learners $\hat{m}_k$ and $\hat{\mu}_k$ minimize the posterior-weighted empirical risk targeting the proxy-weighted conditional expectations:

$$m_k^\gamma(X_t) \equiv \frac{\mathbb{E}[\gamma_{tk} T_t \mid X_t]}{\mathbb{E}[\gamma_{tk} \mid X_t]}, \quad \mu_k^\gamma(X_t) \equiv \frac{\mathbb{E}[\gamma_{tk} Y_t \mid X_t]}{\mathbb{E}[\gamma_{tk} \mid X_t]}. \quad (26)$$

Define the centered residuals at proxy targets: $\tilde{T}_{tk}^\gamma \equiv T_t - m_k^\gamma(X_t)$, $\tilde{Y}_{tk}^\gamma \equiv Y_t - \mu_k^\gamma(X_t)$, and at oracle targets: $\tilde{T}_{tk}^* \equiv T_t - m_k^*(X_t)$, $\tilde{Y}_{tk}^* \equiv Y_t - \mu_k^*(X_t)$. The population proxy Gram matrix $\mathbf{J} \in \mathbb{R}^{K \times K}$ and score vector $\mathbf{S} \in \mathbb{R}^K$ are:

$$J_{jk} \equiv \mathbb{E}[\gamma_{tj} \gamma_{tk} \tilde{T}_{tj}^\gamma \tilde{T}_{tk}^\gamma], \quad S_k \equiv \mathbb{E}[\gamma_{tk} \tilde{T}_{tk}^\gamma \tilde{Y}_{tk}^\gamma], \quad (27)$$

and their oracle counterparts $\mathbf{J}^*, \mathbf{S}^*$ are defined by:

$$J_{jk}^* \equiv \mathbb{E}[H_{tj} H_{tk} \tilde{T}_{tj}^* \tilde{T}_{tk}^*] = \delta_{jk} \mathbb{E}[H_{tk} (\tilde{T}_{tk}^*)^2], \quad S_k^* \equiv \mathbb{E}[H_{tk} \tilde{T}_{tk}^* \tilde{Y}_{tk}^*], \quad (28)$$

where $H_{tj} H_{tk} = \delta_{jk} H_{tk}$ because latent regimes are mutually exclusive.

Step 2: Orthogonality structure of the coupled score (Lemma 6).

Lemma 6 (Orthogonality Structure of the Coupled Score). *Let $\Psi_k(\boldsymbol{\theta}; \eta) = \gamma_k \tilde{T}_k (\tilde{Y}_k - \sum_j \gamma_j \tilde{T}_j \theta_j)$ with $\tilde{T}_k = T - m_k(X)$, $\tilde{Y}_k = Y - \mu_k(X)$, so $\hat{\mathbf{S}} - \hat{\mathbf{J}}\boldsymbol{\theta} = \frac{1}{N} \sum_t \boldsymbol{\Psi}_t$. At $\eta^\gamma = (m^\gamma, \mu^\gamma)$ and any $\boldsymbol{\theta}$, with $g(X) \equiv \mathbb{E}[\|\boldsymbol{\gamma} - \mathbf{H}\|_1 \mid X]$:*

- (i) $\partial_\mu \mathbb{E} \Psi_k[\Delta \mu] = 0$ for any proxy γ .

- (ii) The cross-moment part $\gamma_k \tilde{T}_k \tilde{Y}_k$ is orthogonal in both directions.
 (iii) $|\partial_m \mathbb{E} \Psi_k[\Delta m]| \leq 2M_T \|\boldsymbol{\theta}\|_1 \sqrt{2\varepsilon_\gamma} \max_i \|\Delta m_i\|_{L_2}$. For hard assignments $\gamma_t \in \{0, 1\}^K$, this derivative is identically 0.

Proof. (i)-(ii): $\mathbb{E}[\gamma_k(T - m_k^\gamma) | X] = 0$ and $\mathbb{E}[\gamma_k(Y - \mu_k^\gamma) | X] = 0$ by definition of m_k^γ, μ_k^γ . (iii): differentiating $-\gamma_k \tilde{T}_k \sum_j \gamma_j \tilde{T}_j \theta_j$ in direction Δm gives $\sum_j \theta_j \mathbb{E}[\Delta m_k \gamma_k \gamma_j \tilde{T}_j + \Delta m_j \gamma_k \gamma_j \tilde{T}_k]$. For $j = k$, $\mathbb{E}[\gamma_k^2 \tilde{T}_k | X] = \mathbb{E}[(\gamma_k^2 - \gamma_k) \tilde{T}_k | X]$ since $\mathbb{E}[\gamma_k \tilde{T}_k | X] = 0$, bounded by $M_T |\gamma_k - H_k|$. For $j \neq k$, $\gamma_k \gamma_j \leq |\gamma_k - H_k| + |\gamma_j - H_j|$. Bounding each moment by $M_T g(X)$ and using $\|g\|_{L_2} \leq \sqrt{2\varepsilon_\gamma}$ yields (iii). For hard assignments, $\gamma_k \gamma_j = \delta_{kj} \gamma_k$ so both moments vanish. $\square$

Step 3: Isolate proxy-induced bias via oracle score identity. Under SCM (1)-(5), on the event $S_t = k$ ($H_{tk} = 1$), the oracle outcome residual is $\tilde{Y}_{tk}^* = Y_t - \mu_k^*(X_t) = \theta_k^* \tilde{T}_{tk}^* + U_t$. Substituting into S_k^* :

$$S_k^* = \mathbb{E} \left[H_{tk} \tilde{T}_{tk}^* (\theta_k^* \tilde{T}_{tk}^* + U_t) \right] = \theta_k^* \mathbb{E} [H_{tk} (\tilde{T}_{tk}^*)^2] + \mathbb{E} \left[H_{tk} \tilde{T}_{tk}^* \mathbb{E}[U_t | X_t, S_t, T_t, Z_{1:N}] \right] = \theta_k^* J_{kk}^*, \quad (29)$$

where the innovation term vanishes by Assumption 1. Because $\mathbf{J}^* = \text{diag}(J_{11}^*, \dots, J_{KK}^*)$, this yields the exact vector identity:

$$\mathbf{S}^* = \mathbf{J}^* \boldsymbol{\theta}^*. \quad (30)$$

Now consider the proxy score discrepancy $\mathbf{b}_\gamma \equiv \mathbf{S} - \mathbf{J} \boldsymbol{\theta}^*$. Expanding around the oracle identity:

$$\mathbf{b}_\gamma = (\mathbf{S} - \mathbf{S}^*) - (\mathbf{J} - \mathbf{J}^*) \boldsymbol{\theta}^*. \quad (31)$$

Step 4: Bound score and Gram proxy discrepancies. We bound each component via telescoping expansions in $\varepsilon_\gamma \equiv \mathbb{E} \|\gamma_t - \mathbf{H}_t\|_1$. First, under positivity (Assumption 2), $\pi_k(x) \equiv \mathbb{E}[H_{tk} | X_t = x] \geq c_0 > 0$ and $\pi_k^\gamma(x) \equiv \mathbb{E}[\gamma_{tk} | X_t = x] \geq c_0 > 0$. For any bounded variable V_t with $|V_t| \leq M_V$ almost surely:

$$\begin{aligned} \left| \frac{\mathbb{E}[\gamma_{tk} V_t | X_t]}{\pi_k^\gamma(X_t)} - \frac{\mathbb{E}[H_{tk} V_t | X_t]}{\pi_k(X_t)} \right| &\leq \frac{\mathbb{E}[|\gamma_{tk} - H_{tk}| |V_t| | X_t]}{\pi_k^\gamma(X_t)} + \mathbb{E}[H_{tk} |V_t| | X_t] \left| \frac{\pi_k(X_t) - \pi_k^\gamma(X_t)}{\pi_k^\gamma(X_t) \pi_k(X_t)} \right| \\ &\leq \frac{2M_V}{c_0^2} \mathbb{E}[|\gamma_{tk} - H_{tk}| | X_t]. \end{aligned} \quad (32)$$

Applying this bound to $V_t = T_t$ and $V_t = Y_t$ yields $\mathbb{E}|m_k^\gamma(X_t) - m_k^*(X_t)| \leq \tilde{C}_m \varepsilon_\gamma$ and $\mathbb{E}|\mu_k^\gamma(X_t) - \mu_k^*(X_t)| \leq \tilde{C}_\mu \varepsilon_\gamma$. Second, expanding the score difference $S_k - S_k^* = \mathbb{E}[\gamma_{tk} \tilde{T}_{tk}^\gamma \tilde{Y}_{tk}^\gamma - H_{tk} \tilde{T}_{tk}^* \tilde{Y}_{tk}^*]$:

$$\gamma_{tk} \tilde{T}_{tk}^\gamma \tilde{Y}_{tk}^\gamma - H_{tk} \tilde{T}_{tk}^* \tilde{Y}_{tk}^* = (\gamma_{tk} - H_{tk}) \tilde{T}_{tk}^\gamma \tilde{Y}_{tk}^\gamma - H_{tk} (m_k^\gamma - m_k^*) \tilde{Y}_{tk}^\gamma - H_{tk} \tilde{T}_{tk}^* (\mu_k^\gamma - \mu_k^*). \quad (33)$$

Under bounded support ($\sup_t |\tilde{T}_{tk}^\gamma| \leq M_T$, $\sup_t |\tilde{Y}_{tk}^\gamma| \leq M_Y$), taking expectations yields $|S_k - S_k^*| \leq (M_T M_Y + M_Y \tilde{C}_m + M_T \tilde{C}_\mu) \varepsilon_\gamma \equiv C_S \varepsilon_\gamma$. Third, expanding the Gram difference $J_{jk} - J_{jk}^* = \mathbb{E}[\gamma_{tj} \gamma_{tk} \tilde{T}_{tj}^\gamma \tilde{T}_{tk}^\gamma - H_{tj} H_{tk} \tilde{T}_{tj}^* \tilde{T}_{tk}^*]$:

$$\begin{aligned} \gamma_{tj} \gamma_{tk} \tilde{T}_{tj}^\gamma \tilde{T}_{tk}^\gamma - H_{tj} H_{tk} \tilde{T}_{tj}^* \tilde{T}_{tk}^* &= (\gamma_{tj} \gamma_{tk} - H_{tj} H_{tk}) \tilde{T}_{tj}^\gamma \tilde{T}_{tk}^\gamma - H_{tj} H_{tk} (m_j^\gamma - m_j^*) \tilde{T}_{tk}^\gamma \\ &\quad - H_{tj} H_{tk} \tilde{T}_{tj}^* (m_k^\gamma - m_k^*). \end{aligned} \quad (34)$$

Since $|\gamma_{tj} \gamma_{tk} - H_{tj} H_{tk}| \leq |\gamma_{tj} - H_{tj}| + |\gamma_{tk} - H_{tk}|$, taking expectations gives $|J_{jk} - J_{jk}^*| \leq (2M_T^2 + 2M_T \tilde{C}_m) \varepsilon_\gamma$, so $\|\mathbf{J} - \mathbf{J}^*\|_{\text{op}} \leq C_J \varepsilon_\gamma$. Combining bounds yields:

$$\|\mathbf{b}_\gamma\|_2 \leq \|\mathbf{S} - \mathbf{S}^*\|_2 + \|\mathbf{J} - \mathbf{J}^*\|_{\text{op}} \|\boldsymbol{\theta}^*\|_2 \leq (C_S + C_J M_\theta) \varepsilon_\gamma \equiv C \cdot \varepsilon_\gamma. \quad (35)$$

Step 5: Population bound and finite-sample rate. Since $\mathbf{S} = \mathbf{J} \boldsymbol{\theta}^* + \mathbf{b}_\gamma$ and $\boldsymbol{\theta}_\gamma \equiv \mathbf{J}^{-1} \mathbf{S}$, we have $\boldsymbol{\theta}_\gamma - \boldsymbol{\theta}^* = \mathbf{J}^{-1} \mathbf{b}_\gamma$, so $\|\boldsymbol{\theta}_\gamma - \boldsymbol{\theta}^*\|_2 \leq \frac{\|\mathbf{b}_\gamma\|_2}{\lambda_{\min}(\mathbf{J})} \leq \frac{C \varepsilon_\gamma}{\lambda_{\min}(\mathbf{J})}$, establishing part (i). For part (ii): since $\mathbf{J} \boldsymbol{\theta}_\gamma = \mathbf{S}$, $\hat{\boldsymbol{\theta}}_\gamma - \boldsymbol{\theta}_\gamma = \hat{\mathbf{J}}^{-1} [(\hat{\mathbf{S}} - \mathbf{S}) - (\hat{\mathbf{J}} - \mathbf{J}) \boldsymbol{\theta}_\gamma]$. By Weyl's inequality, $\lambda_{\min}(\hat{\mathbf{J}}) \geq \lambda_{\min}(\mathbf{J}) - \|\hat{\mathbf{J}} - \mathbf{J}\|_{\text{op}} \geq \lambda_{\min}(\mathbf{J})/2$ on event $\mathcal{E}_N$, so $\|\hat{\mathbf{J}}^{-1}\|_{\text{op}} \leq 2/\lambda_{\min}(\mathbf{J})$. Lemma 6 gives $\|\hat{\mathbf{S}} - \mathbf{S}\|_2 = \mathcal{O}_P(N^{-1/2} + r_N^2)$, while non-orthogonality of the coupled Gram matrix yields $\|\hat{\mathbf{J}} - \mathbf{J}\|_{\text{op}} = \mathcal{O}_P(N^{-1/2} + r_N)$. Combining with $\|\boldsymbol{\theta}_\gamma\|_2 \leq M_\theta + C \varepsilon_\gamma / \lambda_{\min}(\mathbf{J})$ establishes the finite-sample bound in (ii). $\blacksquare$

Proposition 3 (Exact Proxy Bias in the Symmetric Flip Model). *Let $K = 2$, $\boldsymbol{\pi}^* = (\frac{1}{2}, \frac{1}{2})$, $T_t = V_t$ with $V_t \perp (S_t, X_t, U_t)$, $\mathbb{E} V_t = 0$, $\text{Var} V_t = \sigma^2$, $Y_t = \theta_{S_t} T_t + g_{S_t} + U_t$, $\mathbb{E}[U_t | T_t, S_t] = 0$, and $\boldsymbol{\gamma}_t = (1 - \frac{\varepsilon}{2}) \mathbf{H}_t + \frac{\varepsilon}{2} (\mathbf{1} - \mathbf{H}_t)$ with $\varepsilon \in [0, 1)$ ($\varepsilon_\gamma = \varepsilon$). Then $\lambda_{\min}(\mathbf{J}) = \sigma^2(1 - \varepsilon)^2/2$ and*

$$\boldsymbol{\theta}_\gamma - \boldsymbol{\theta}^* = \frac{\varepsilon}{1 - \varepsilon} \cdot \frac{\theta_1 - \theta_2}{2} (1, -1)^\top, \quad \|\boldsymbol{\theta}_\gamma - \boldsymbol{\theta}^*\|_2 = \frac{\varepsilon}{1 - \varepsilon} \frac{|\theta_1 - \theta_2|}{\sqrt{2}}. \quad (36)$$

Proof. $m_k^\gamma = 0$ since $\mathbb{E} V = 0, V \perp S$, so $\tilde{T}_k = V$ and constants in $\tilde{Y}_k$ drop out of $\mathbf{S}$. With $a = 1 - \varepsilon/2$, $b = \varepsilon/2$, $\mathbf{M} = \begin{pmatrix} a & b \\ b & a \end{pmatrix}$: $\mathbf{J} = \frac{\sigma^2}{2} \mathbf{M}^2$ and $\mathbf{S} = \frac{\sigma^2}{2} \mathbf{M} \boldsymbol{\theta}^*$. Since $\mathbf{M}$ has eigenvalues 1 (eigenvector $(1, 1)^\top$) and $1 - \varepsilon$ (eigenvector $(1, -1)^\top$), $\boldsymbol{\theta}_\gamma = \mathbf{J}^{-1} \mathbf{S} = \mathbf{M}^{-1} \boldsymbol{\theta}^*$, $\lambda_{\min}(\mathbf{J}) = \sigma^2(1 - \varepsilon)^2/2$, and $(\mathbf{M}^{-1} - \mathbf{I}) \boldsymbol{\theta}^*$ evaluates to the displayed vector. $\square$

Theorem 7 (Matching Minimax Lower Bound under Latent Confounding). *Let $\mathcal{P}(\varepsilon_\gamma, \lambda_0)$ denote the family of data-generating processes satisfying Assumptions 1-4 and Assumption 2 with binary regimes ($K = 2$), state proxy error $\mathbb{E}\|\gamma_t - \mathbf{H}_t\|_1 \leq \varepsilon_\gamma \in (0, 1/2]$, and minimal Gram eigenvalue $\lambda_{\min}(\mathbf{J}) \geq \lambda_0 > 0$. Then there exists a universal numerical constant $c > 0$ such that for any sample size $N \geq 1$:*

$$\inf_{\hat{\boldsymbol{\theta}}} \sup_{P \in \mathcal{P}(\varepsilon_\gamma, \lambda_0)} \mathbb{E}_P \left[\|\hat{\boldsymbol{\theta}} - \boldsymbol{\theta}^*(P)\|_2 \right] \geq c \frac{\varepsilon_\gamma}{\lambda_0}. \quad (37)$$

Consequently, the degradation rate $\mathcal{O}(\varepsilon_\gamma / \lambda_{\min}(\mathbf{J}))$ in Theorem 3 is minimax optimal and cannot be improved by any estimator.

Proof. We employ Le Cam’s two-point method (Tsybakov, 2009). Consider a static binary regime setup ($K = 2$) with balanced priors $\pi_1 = \pi_2 = 1/2$, $X_t \in \{0\}$, and scalar treatment innovation $V_t \sim \mathcal{N}(0, \sigma^2)$ independent of (S_t, U_t) , where $\sigma^2 = 2\lambda_0 / (1 - \varepsilon_\gamma)^2$. Let proxy $Z_t \in \{1, 2\}$ satisfy $P(Z_t = S_t | S_t) = 1 - \varepsilon_\gamma$ and $P(Z_t = 3 - S_t | S_t) = \varepsilon_\gamma$, so $\mathbb{E}\|\gamma_t - \mathbf{H}_t\|_1 = \varepsilon_\gamma$. By Proposition 3, $\lambda_{\min}(\mathbf{J}) = \sigma^2(1 - \varepsilon_\gamma)^2/2 = \lambda_0$. Construct two candidate data-generating mechanisms P_0 and P_1 : under P_0 , $\boldsymbol{\theta}^{(0)} = (\delta/2, -\delta/2)^\top$ with baselines $g_1(0) = g_2(0) = 0$; under P_1 , $\boldsymbol{\theta}^{(1)} = (-\delta/2, \delta/2)^\top$ with baseline shifts $g_1^{(1)} = \Delta g$ and $g_2^{(1)} = -\Delta g$, where $\Delta g = \frac{1-2\varepsilon_\gamma}{2\varepsilon_\gamma} \delta \mathbb{E}[|T_t|]$. Parameter separation is $\|\boldsymbol{\theta}^{(0)} - \boldsymbol{\theta}^{(1)}\|_2 = \sqrt{2}\delta$. Choosing $\delta = c_1 \frac{\varepsilon_\gamma}{\lambda_0}$ for sufficiently small $c_1 > 0$ ensures the Kullback-Leibler divergence between joint observable distributions satisfies $\text{KL}(P_0^N, P_1^N) \leq \alpha < \infty$. By Le Cam’s inequality (Tsybakov, 2009; Bickel et al., 1993), the minimax risk over $\{P_0, P_1\} \subset \mathcal{P}(\varepsilon_\gamma, \lambda_0)$ satisfies:

$$\inf_{\hat{\boldsymbol{\theta}}} \sup_{P \in \{P_0, P_1\}} \mathbb{E}_P \|\hat{\boldsymbol{\theta}} - \boldsymbol{\theta}^*(P)\|_2 \geq \frac{\|\boldsymbol{\theta}^{(0)} - \boldsymbol{\theta}^{(1)}\|_2}{4} \exp(-\text{KL}(P_0^N, P_1^N)) \geq \frac{\sqrt{2}\delta}{4} e^{-\alpha} \geq c \frac{\varepsilon_\gamma}{\lambda_0},$$

which proves the matching lower bound (37). $\square$

C.1 The Latent-Confounding Difficulty Frontier $\mathcal{D}_{\text{causal}}$

Theorem 3 demonstrates that the upper bound on structural causal estimation error decomposes into:

$$\|\hat{\boldsymbol{\theta}}_\gamma - \boldsymbol{\theta}^*\|_2 \leq C \cdot \mathcal{D}_{\text{causal}} + \mathcal{O}_P(N^{-1/2}), \quad \text{where} \quad \mathcal{D}_{\text{causal}} \equiv \frac{\varepsilon_\gamma}{\lambda_{\min}(\mathbf{J})}. \quad (38)$$

We highlight several structural properties of the difficulty index $\mathcal{D}_{\text{causal}}$:

1. **Scale-Dependent Difficulty Ratio:** Both population proxy error ε_γ and Gram conditioning $\lambda_{\min}(\mathbf{J})$ interact multiplicatively in governing downstream estimation risk (dimensionless under standardized treatment units). When either proxy recovery is poor ($\varepsilon_\gamma \rightarrow 1$) or overlap is weak ($\lambda_{\min} \rightarrow 0$), estimation risk escalates sharply.
2. **Asymmetric Sensitivity:** In our factorial simulations (Table 8), reducing $\lambda_{\min}$ by an order of magnitude produces an equivalent variance and bias inflation to a tenfold increase in state proxy error ε_γ .

C.2 Proof of Proposition 1: Surrogate Entropy-Difficulty Bridge

Let $\gamma_t \equiv P(S_t \in \cdot | \mathcal{F}_t) \in \Delta^K$ be the Bayesian posterior state distribution with respect to filtration $\mathcal{F}_t \equiv \sigma(Z_{1:t})$. Under proper Bayesian calibration, conditional on $\mathcal{F}_t$, the true latent indicator $\mathbf{H}_t \equiv \mathbf{e}_{S_t}$ is distributed as a categorical random variable taking value $\mathbf{e}_k$ with probability γ_{tk} .

Step 1: Expected L_1 Error as Gini Impurity. Evaluating the conditional expectation of the L_1 error:

$$\begin{aligned} \mathbb{E}[\|\gamma_t - \mathbf{H}_t\|_1 | \mathcal{F}_t] &= \sum_{k=1}^K \gamma_{tk} \|\gamma_t - \mathbf{e}_k\|_1 = \sum_{k=1}^K \gamma_{tk} \left((1 - \gamma_{tk}) + \sum_{j \neq k} \gamma_{tj} \right) \\ &= \sum_{k=1}^K \gamma_{tk} \cdot 2(1 - \gamma_{tk}) = 2 \left(1 - \sum_{k=1}^K \gamma_{tk}^2 \right) = 2 \text{Gini}(\gamma_t). \end{aligned} \quad (39)$$

Step 2: Entropy Bound on Gini Impurity. By the standard concave inequality $\ln(x) \leq x - 1$, for any $x > 0$:

$$H(\gamma_t) = - \sum_{k=1}^K \gamma_{tk} \ln \gamma_{tk} = \sum_{k=1}^K \gamma_{tk} \ln \left(\frac{1}{\gamma_{tk}} \right) \geq \sum_{k=1}^K \gamma_{tk} (1 - \gamma_{tk}) = \text{Gini}(\gamma_t). \quad (40)$$

Furthermore, for any binary regime partition ($K = 2$) with $\gamma_{t1} = p$ and $\gamma_{t2} = 1 - p$, $2 \text{Gini}(\gamma_t) = 4p(1 - p)$. Comparing to binary Shannon entropy $H(p) = -p \ln p - (1 - p) \ln(1 - p)$, the ratio $\frac{4p(1-p)}{H(p)}$ achieves its maximum at $p = 1/2$, where $\frac{4(1/4)}{\ln 2} = \frac{1}{\ln 2} \approx 1.4427$. Hence for all $p \in [0, 1]$:

$$2 \text{Gini}(\gamma_t) \leq C_K H(\gamma_t), \quad \text{where } C_K \leq \frac{1}{\ln 2} \approx 1.4427. \quad (41)$$

Step 3: Population and Observable Bridge. Taking the population expectation of Step 1 and applying Step 2:

$$\varepsilon_\gamma \equiv \mathbb{E}[\mathbb{E}[\|\gamma_t - \mathbf{H}_t\|_1 \mid \mathcal{F}_t]] = 2\mathbb{E}[\text{Gini}(\gamma_t)] \leq C_K \mathbb{E}[H(\gamma_t)]. \quad (42)$$

Dividing by the global minimal singular value $\lambda_{\min}(\mathbf{J})$:

$$\mathcal{D}_{\text{causal}} \equiv \frac{\varepsilon_\gamma}{\lambda_{\min}(\mathbf{J})} \leq C_K \mathbb{E}\left[\frac{H(\gamma_t)}{\lambda_{\min}(\mathbf{J})}\right] = C_K \mathbb{E}[\mathcal{D}_{\text{global}}^H(t)]. \quad (43)$$

Furthermore, under Assumption 5 (Two-Sided Spectral Comparability), the rolling Gram matrix satisfies $\lambda_{\min}(\mathbf{J}_t) \leq \bar{c}_J \lambda_{\min}(\mathbf{J})$ almost surely for all evaluation windows $t \geq w$, which implies $\frac{1}{\lambda_{\min}(\mathbf{J})} \leq \frac{\bar{c}_J}{\lambda_{\min}(\mathbf{J}_t)}$. Pointwise substitution yields:

$$\mathcal{D}_{\text{causal}} \leq C_K \bar{c}_J \mathbb{E}\left[\frac{H(\gamma_t)}{\lambda_{\min}(\mathbf{J}_t)}\right] = C_K \bar{c}_J \mathbb{E}[\mathcal{D}_{\text{operational}}(t)], \quad (44)$$

and $\mathbb{E}[\mathcal{D}_{\text{operational}}(t)] \leq \bar{c}_J^{-1} \mathbb{E}[\mathcal{D}_{\text{global}}^H(t)]$. In simulation benchmarks where true states are known and posteriors are evaluated at population parameters $\mathbf{\Lambda}^*$, the ergodic theorem and central limit theorem for stationary α -mixing chains ensure that $\hat{\varepsilon}_{\gamma, N} \equiv \frac{1}{N} \sum_{t=1}^N \|\gamma_t - \mathbf{H}_t\|_1 = \varepsilon_\gamma + \mathcal{O}_P(N^{-1/2})$ (preserved under $\sqrt{N}$ -consistent estimation of $\mathbf{\Lambda}$), providing an evaluation-only benchmark surrogate $\hat{\mathcal{D}}_{\text{causal}, N} \equiv \hat{\varepsilon}_{\gamma, N} / \lambda_{\min}(\mathbf{J}) = \mathcal{D}_{\text{causal}} + \mathcal{O}_P(N^{-1/2})$. We emphasize that these bounds establish $\mathcal{D}_{\text{operational}}(t)$ as an observable surrogate upper bound for $\mathcal{D}_{\text{causal}}$, not a two-sided equivalence. Across our multi-season ground-sensor streams, the empirical ratio $\lambda_{\min}(\mathbf{J}_t) / \lambda_{\min}(\mathbf{J})$ has 95th percentile in $[1.06, 1.72]$ and maximum ≤ 7.35 , providing empirical support for the spectral-comparability assumption with finite $\bar{c}_J$. ■

C.3 A Priori Proxy Diagnostics and Separation Screening

Before fitting OR-DML or downstream prediction heads, practitioners should assess proxy informativeness on auxiliary measurements $Z_t \in \mathbb{R}^q$. Let $\hat{\boldsymbol{\mu}}_{Z, k} \equiv \frac{1}{N_k} \sum_{t: \hat{S}_t = k} Z_t$ and $\hat{\boldsymbol{\Sigma}}_{Z, k}$ denote empirical cluster moments. We define the standardized separation index:

$$\Delta_Z \equiv \min_{j \neq k} \frac{\|\hat{\boldsymbol{\mu}}_{Z, j} - \hat{\boldsymbol{\mu}}_{Z, k}\|_2}{\sqrt{\frac{1}{2}(\text{Tr}(\hat{\boldsymbol{\Sigma}}_{Z, j}) + \text{Tr}(\hat{\boldsymbol{\Sigma}}_{Z, k}))/q}}. \quad (45)$$

When $\Delta_Z < 0.5$ (or silhouette score < 0.15), the overlap between regime emission densities $P(Z_t \mid S_t = j)$ and $P(Z_t \mid S_t = k)$ is severe, driving posterior entropy $\bar{H}(\gamma) \rightarrow \ln K$. By Theorem 2, point identification fails identically in the limit $\Delta_Z \rightarrow 0$. Deployed pipelines should enforce an automated screening gate: if $\Delta_Z < 0.5$ or mean entropy exceeds $0.80 \ln K$, the system must flag non-identifiability and refuse to emit point causal effects without external structural constraints.

D PROOF OF THEOREM 4: LEADING-ORDER RISK BOUND AND SPECTRAL REGULARIZATION

Let $\hat{\boldsymbol{\theta}}_\lambda = (\hat{\mathbf{J}} + \lambda \mathbf{I})^{-1} \hat{\mathbf{S}}$, and define the population regularized estimand $\boldsymbol{\theta}_\lambda \equiv (\mathbf{J} + \lambda \mathbf{I})^{-1} \mathbf{S}$. We decompose the estimation error around the regularized population estimand $\boldsymbol{\theta}_\lambda$ via the exact algebraic identity:

$$\mathbb{E}\|\hat{\boldsymbol{\theta}}_\lambda - \boldsymbol{\theta}^*\|_2^2 = \|\boldsymbol{\theta}_\lambda - \boldsymbol{\theta}^*\|_2^2 + \mathbb{E}\|\hat{\boldsymbol{\theta}}_\lambda - \boldsymbol{\theta}_\lambda\|_2^2 + 2\mathbb{E}[\hat{\boldsymbol{\theta}}_\lambda - \boldsymbol{\theta}_\lambda, \boldsymbol{\theta}_\lambda - \boldsymbol{\theta}^*], \quad (46)$$

which accounts explicitly for the cross term between stochastic estimation error and deterministic regularization bias. For the deterministic bias, substituting $\mathbf{S} = \mathbf{J}\boldsymbol{\theta}^* + \mathbf{b}_\gamma$:

$$\begin{aligned} \boldsymbol{\theta}_\lambda - \boldsymbol{\theta}^* &= (\mathbf{J} + \lambda \mathbf{I})^{-1} (\mathbf{S} - (\mathbf{J} + \lambda \mathbf{I})\boldsymbol{\theta}^*) \\ &= (\mathbf{J} + \lambda \mathbf{I})^{-1} (\mathbf{S} - \mathbf{J}\boldsymbol{\theta}^* - \lambda \boldsymbol{\theta}^*) \\ &= (\mathbf{J} + \lambda \mathbf{I})^{-1} (\mathbf{b}_\gamma - \lambda \boldsymbol{\theta}^*), \end{aligned} \quad (47)$$

which establishes the exact deterministic identity (47). Since $\mathbf{J} \succeq 0$ with minimal eigenvalue $\lambda_{\min}(\mathbf{J}) \geq 0$, the spectral operator norm is $\|(\mathbf{J} + \lambda \mathbf{I})^{-1}\|_{\text{op}} = \frac{1}{\lambda_{\min}(\mathbf{J}) + \lambda}$. Applying the triangle inequality:

$$\|\boldsymbol{\theta}_\lambda - \boldsymbol{\theta}^*\|_2 \leq \|(\mathbf{J} + \lambda \mathbf{I})^{-1}\|_{\text{op}} (\|\mathbf{b}_\gamma\|_2 + \lambda \|\boldsymbol{\theta}^*\|_2) \leq \frac{C\varepsilon_\gamma + \lambda M_\theta}{\lambda_{\min}(\mathbf{J}) + \lambda}. \quad (48)$$

For the stochastic variance, expanding $\hat{\theta}_\lambda - \theta_\lambda$:

$$\hat{\theta}_\lambda - \theta_\lambda = (\mathbf{J} + \lambda \mathbf{I})^{-1} (\hat{\mathbf{S}} - \mathbf{S} - (\hat{\mathbf{J}} - \mathbf{J})\theta_\lambda) + o_p(N^{-1/2}). \quad (49)$$

By Theorem 5(a), $\sqrt{N}(\hat{\theta}_\lambda - \theta_\lambda) \Rightarrow \mathcal{N}(\mathbf{0}, \Sigma_\lambda^{\text{tot}})$. Under uniform integrability of $N\|\hat{\theta}_\lambda - \theta_\lambda\|_2^2$, stochastic variance satisfies $V_N \equiv \mathbb{E}\|\hat{\theta}_\lambda - \theta_\lambda\|_2^2 = \frac{1}{N}\text{Tr}(\Sigma_\lambda^{\text{tot}}) + o(N^{-1}) \leq \frac{\text{Tr}(\Sigma_\lambda^{\text{tot}})}{N(\lambda_{\min}(\mathbf{J}) + \lambda)^2} + o(N^{-1})$, using $\text{Tr}(\mathbf{A}^{-1}\mathbf{\Omega}\mathbf{A}^{-1}) \leq \|\mathbf{A}^{-2}\|_{\text{op}}\text{Tr}(\mathbf{\Omega})$ for $\mathbf{A} = \mathbf{J} + \lambda \mathbf{I} \succ 0$. Applying Minkowski's inequality in $L_2(P; \mathbb{R}^K)$ to the decomposition $(\hat{\theta}_\lambda - \theta^*) = (\hat{\theta}_\lambda - \theta_\lambda) + (\theta_\lambda - \theta^*)$:

$$\left(\mathbb{E}\|\hat{\theta}_\lambda - \theta^*\|_2^2\right)^{1/2} \leq \|\hat{\theta}_\lambda - \theta_\lambda\|_{L_2} + \|\theta_\lambda - \theta^*\|_2 \leq \sqrt{V_N} + B_\lambda, \quad (50)$$

which establishes the finite-sample risk upper bound:

$$\mathbb{E}\|\hat{\theta}_\lambda - \theta^*\|_2^2 \leq (B_\lambda + \sqrt{V_N})^2 \leq 2B_\lambda^2 + 2V_N \leq 2\left(\frac{C\varepsilon_\gamma + \lambda M_\theta}{\lambda_{\min}(\mathbf{J}) + \lambda}\right)^2 + \frac{2\text{Tr}(\Sigma_\lambda^{\text{tot}})}{N(\lambda_{\min}(\mathbf{J}) + \lambda)^2} + o(N^{-1}), \quad (51)$$

completely avoiding any requirement to bound the stochastic bias cross-term.

Risk Bound Minimizer vs. Empirical Grid Optimum: Differentiating the leading-order risk upper bound $R_{\text{bound}}(\lambda) \equiv \left(\frac{C\varepsilon_\gamma + \lambda M_\theta}{\lambda_{\min}(\mathbf{J}) + \lambda}\right)^2 + \frac{\text{Tr}(\Sigma_\lambda^{\text{tot}})}{N(\lambda_{\min}(\mathbf{J}) + \lambda)^2}$ with respect to $\lambda \geq 0$, the bound-optimal shrinkage parameter $\lambda_{\text{bound}}^* \equiv \arg \min_{\lambda \geq 0} R_{\text{bound}}(\lambda)$ satisfies:

$$\frac{\partial}{\partial \lambda} \left[\left(\frac{C\varepsilon_\gamma + \lambda M_\theta}{\lambda_{\min}(\mathbf{J}) + \lambda}\right)^2 + \frac{\text{Tr}(\Sigma_\lambda^{\text{tot}})}{N(\lambda_{\min}(\mathbf{J}) + \lambda)^2} \right] = 0. \quad (52)$$

When the overlap Gram matrix is ill-conditioned ($\lambda_{\min}(\mathbf{J}) \approx 0$), the variance term $\frac{\text{Tr}(\Sigma_\lambda^{\text{tot}})}{N(\lambda_{\min} + \lambda)^2}$ dominates for small λ , while the deterministic shrinkage bias $\left(\frac{\lambda M_\theta}{\lambda_{\min} + \lambda}\right)^2$ dominates for large λ . The interior minimizer $\lambda_{\text{bound}}^* > 0$ balances variance reduction against shrinkage bias in the upper bound. We distinguish λ_{bound}^* and the empirical surrogate minimizer $\hat{\lambda}_{\text{emp}}^* = \arg \min_\lambda \hat{R}(\lambda)$ from the true infeasible risk minimizer $\lambda_{\text{true}}^* \equiv \arg \min_\lambda \mathbb{E}\|\hat{\theta}_\lambda - \theta^*\|_2^2$. In our empirical simulation benchmark across 100 Monte Carlo replications (`reports/or_dml_regularization_frontier.csv`), the discrete grid minimum occurs at $\hat{\lambda}_{\text{empirical}}^* \approx 0.20$ (MSE = 0.0321, representing a $25.9\times$ empirical risk reduction in this benchmark DGP relative to the unregularized estimator $\lambda = 0.0$ where MSE = 0.8313), while $\lambda = 0.10$ achieves MSE = 0.0362 ($23.0\times$ empirical reduction) with peak coverage of 85.0%. We emphasize that Theorem 4 provides the functional upper-bound characterization of the bias-variance trade-off; the exact numerical reduction factors ($23.0\times$ and $25.9\times$) are empirical simulation properties of the benchmark DGP rather than universal theoretical constants. ■

D.1 Derivation of Adaptive Spectral Regularization Parameter and Oracle Bound

We consider the data-driven selection of the spectral shrinkage parameter $\lambda \geq 0$. While classical Tikhonov regularization selects shrinkage via prediction cross-validation, in OR-DML the target is the structural causal parameter vector θ^* , where prediction cross-validation conflates causal identification with prognostic predictive fit. We distinguish two distinct regularization regimes:

1. **Inferential Regularization** ($\lambda = o(N^{-1/2})$): Required by Theorem 5 to guarantee asymptotic unbiasedness $\sqrt{N}\text{Bias}(\lambda) \rightarrow \mathbf{0}$, ensuring centered Gaussian limit laws and valid nominal confidence interval coverage.
2. **Finite-Sample Risk Minimization** ($\lambda_{\text{risk}} = \mathcal{O}(1)$): Minimizes the finite-sample expected ℓ_2 estimation risk $\mathbb{E}\|\hat{\theta}_\lambda - \theta^*\|_2^2$ under severe causal ill-conditioning ($\lambda_{\min}(\mathbf{J}) \approx 0$).

Empirical Risk Surrogate and Convexity: Define the empirical plug-in risk surrogate $\hat{R}(\lambda)$ by substituting empirical estimators $(\hat{\varepsilon}_\gamma, \lambda_{\min}(\hat{\mathbf{J}}), \hat{\mathbf{\Omega}}, \|\hat{\theta}_0\|_2)$ into the risk upper bound of Theorem 4:

$$\hat{R}(\lambda) = \frac{(C\hat{\varepsilon}_\gamma + \lambda\|\hat{\theta}_0\|_2)^2 + N^{-1}\text{Tr}(\hat{\mathbf{\Omega}})}{(\lambda_{\min}(\hat{\mathbf{J}}) + \lambda)^2}. \quad (53)$$

Let $u = \lambda_{\min}(\hat{\mathbf{J}}) + \lambda \geq \lambda_{\min}(\hat{\mathbf{J}}) > 0$. Rewriting $\lambda = u - \lambda_{\min}$, the numerator becomes:

$$\left(C\hat{\varepsilon}_\gamma - \lambda_{\min}(\hat{\mathbf{J}})\|\hat{\theta}_0\|_2 + u\|\hat{\theta}_0\|_2\right)^2 + N^{-1}\text{Tr}(\hat{\mathbf{\Omega}}) = (A + uB)^2 + V, \quad (54)$$

where $A = C\hat{\varepsilon}_\gamma - \lambda_{\min}(\hat{\mathbf{J}})\|\hat{\theta}_0\|_2$, $B = \|\hat{\theta}_0\|_2$, and $V = N^{-1}\text{Tr}(\hat{\mathbf{\Omega}})$. Dividing by u^2 :

$$\hat{R}(u - \lambda_{\min}) = B^2 + \frac{2AB}{u} + \frac{A^2 + V}{u^2}. \quad (55)$$

Differentiating with respect to u :

$$\frac{d\hat{R}}{du} = -\frac{2AB}{u^2} - \frac{2(A^2 + V)}{u^3} = -\frac{2}{u^3} [ABu + (A^2 + V)]. \quad (56)$$

Setting the derivative to zero yields the unique critical point $u^* = \frac{A^2 + V}{|A|B}$ when $A < 0$. When $A \geq 0$ (i.e., $C\hat{\varepsilon}_\gamma \geq \lambda_{\min}(\hat{\mathbf{J}})\|\hat{\boldsymbol{\theta}}_0\|_2$), $\frac{d\hat{R}}{du} < 0$ on $(0, \infty)$, so that the objective is strictly decreasing toward B^2 as $\lambda \rightarrow \infty$. Constraining the search to a compact admissible set $\lambda \in [0, \bar{\lambda}]$ with a fixed upper bound $\bar{\lambda} < \infty$ (independent of N), the bounded empirical minimizer is:

$$\hat{\lambda}^* = \begin{cases} \text{clip} \left(\frac{(C\hat{\varepsilon}_\gamma - \lambda_{\min}\|\hat{\boldsymbol{\theta}}_0\|_2)^2 + N^{-1}\text{Tr}(\hat{\boldsymbol{\Omega}})}{(\lambda_{\min}\|\hat{\boldsymbol{\theta}}_0\|_2 - C\hat{\varepsilon}_\gamma)\|\hat{\boldsymbol{\theta}}_0\|_2} - \lambda_{\min}(\hat{\mathbf{J}}), 0, \bar{\lambda} \right), & \text{if } A < 0, \\ \bar{\lambda}, & \text{if } A \geq 0. \end{cases} \quad (57)$$

Because $R'(u) = -\frac{2}{u^3}[ABu + (A^2 + V)]$ changes sign strictly from negative to positive at u^* when $A < 0$, u^* is the unique global interior minimizer on $(0, \infty)$.

Lemma 8 (Oracle Inequality for the Surrogate). *Let R_{surr} be the population surrogate, $\hat{R}$ its plug-in, and $\Delta_N = \sup_{\lambda \in [0, \bar{\lambda}]} |\hat{R}(\lambda) - R_{\text{surr}}(\lambda)| = \mathcal{O}_P(N^{-1/2})$. Then $R_{\text{surr}}(\hat{\lambda}^*) \leq \inf_{\lambda \in [0, \bar{\lambda}]} R_{\text{surr}}(\lambda) + 2\Delta_N$.*

Proof. For any $\lambda^\circ \in [0, \bar{\lambda}]$: $R_{\text{surr}}(\hat{\lambda}^*) \leq \hat{R}(\hat{\lambda}^*) + \Delta_N \leq \hat{R}(\lambda^\circ) + \Delta_N \leq R_{\text{surr}}(\lambda^\circ) + 2\Delta_N$. $\square$

This guarantee governs the surrogate, not the true risk; calibrated plug-in $\hat{\varepsilon}_\gamma = \frac{2}{N} \sum_t \text{Gini}(\gamma_t)$ preserves ranking when $C\hat{\varepsilon}_\gamma$ dominates. $\blacksquare$

D.2 State Space Scalability and Matrix Concentration for $K > 2$

While our primary experiments study binary regimes ($K = 2$), the OR-DML coupled estimator generalizes directly to arbitrary finite state dimensions $K \geq 2$. Here we formalize the computational and statistical scalability of the coupled Gram matrix $\hat{\mathbf{J}} \in \mathbb{R}^{K \times K}$.

Computational Complexity and Matrix Concentration under α -Mixing: Constructing $\hat{\mathbf{J}}$ and score $\hat{\mathbf{S}}$ requires evaluating K nuisance pairs across B folds ($\mathcal{O}(BK \cdot \text{Cost}(\text{Nuisance}))$). Forming $\hat{\mathbf{J}}$ takes $\mathcal{O}(NK^2)$ operations, and solving $(\hat{\mathbf{J}} + \lambda \mathbf{I})^{-1} \hat{\mathbf{S}}$ via Cholesky factorization takes $\mathcal{O}(K^3)$ operations, taking < 1 ms for $K \leq 20$. Statistically, the empirical Gram matrix $\hat{\mathbf{J}} = \frac{1}{N} \sum_{t=1}^N \mathbf{\Gamma}_t \tilde{\mathbf{T}}_t \tilde{\mathbf{T}}_t^\top \mathbf{\Gamma}_t$ concentrates around $\mathbf{J}$ under geometric α -mixing with spectral gap $\rho_{\text{gap}} \equiv 1 - |\lambda_2(\mathbf{\Pi})| > 0$ via the Matrix Bernstein inequality (Banna et al., 2016):

$$\|\hat{\mathbf{J}} - \mathbf{J}\|_{\text{op}} \leq C_J \sqrt{\frac{K \log(K/\delta)}{N \cdot \rho_{\text{gap}}}} + \frac{C'_J K \log(K/\delta)}{N \cdot \rho_{\text{gap}}}. \quad (58)$$

By Weyl's theorem, $|\lambda_{\min}(\hat{\mathbf{J}}) - \lambda_{\min}(\mathbf{J})| \leq \|\hat{\mathbf{J}} - \mathbf{J}\|_{\text{op}}$, ensuring that $\lambda_{\min}(\hat{\mathbf{J}})$ concentrates at the standard $\mathcal{O}_P(N^{-1/2})$ rate whenever $\lambda_{\min}(\mathbf{J}) > 0$.

Non-Identifiability vs. Shrinkage: When Systems Must Abstain: Spectral shrinkage $\hat{\boldsymbol{\theta}}_\lambda = (\hat{\mathbf{J}} + \lambda \mathbf{I})^{-1} \hat{\mathbf{S}}$ stabilizes variance under ill-conditioned overlap ($\lambda_{\min} \approx 0$). However, when data lack identifying variation entirely-diagnosed by condition number $\kappa(\hat{\mathbf{J}}) > 10^4$ or a moving-block bootstrap 95% confidence interval for $\lambda_{\min}(\hat{\mathbf{J}})$ that covers zero-the structural parameter $\boldsymbol{\theta}^*$ is non-identified (Theorem 2). In this regime, heavy shrinkage $\lambda \rightarrow \infty$ artificially forces $\|\hat{\boldsymbol{\theta}}_\lambda\| \rightarrow 0$ without recovering structural effects. Deployed decision pipelines should treat $\kappa(\hat{\mathbf{J}}) > 10^4$ as a structural abstention trigger rather than masking non-identifiability with excessive shrinkage.

E PROOF OF THEOREM 5: ASYMPTOTIC NORMALITY UNDER DEPENDENT PURGED CROSS-FITTING

Assumption 6 (HMM Regularity, Gram-Nuisance Rate, Moments). (H) The emission model is identifiable up to label permutation; $\hat{\mathbf{A}}$ is the MLE fitted on a calibration window of length N_{cal} separated by the embargo, with $N/N_{\text{cal}} \rightarrow \kappa \in [0, \infty)$ ($\kappa = 0$ for oracle filter), and $\gamma_{tk}(\mathbf{A})$ is continuously differentiable with bounded derivative. Then $\sqrt{N_{\text{cal}}}(\hat{\mathbf{A}} - \mathbf{A}^*) \Rightarrow \mathcal{N}(\mathbf{0}, \mathcal{I}_{\mathbf{A}}^{-1})$ (Leroux, 1992; Bickel et al., 1998; Douc et al., 2004). (G) $\sqrt{\varepsilon_\gamma} \max_k \|\hat{m}_k - m_k^\gamma\|_{L_2} = \mathcal{O}_P(N^{-1/2})$; true for hard assignments and under Assumption 4 when $\varepsilon_\gamma = \mathcal{O}(N^{-1/2})$. (M) $\mathbb{E}\|\Psi_t\|^{2+\delta} < \infty$ for some $\delta > 0$, and $\lambda_{\min}(\mathbf{J}) + \lambda \geq c_\lambda > 0$ is fixed.

Proof of Theorem 5. Step 1 (exact algebra, centering at θ_λ). Because $(\hat{\mathbf{J}} + \lambda \mathbf{I})\hat{\theta}_\lambda = \hat{\mathbf{S}}$ and $(\mathbf{J} + \lambda \mathbf{I})\theta_\lambda = \mathbf{S}$:

$$\hat{\theta}_\lambda - \theta_\lambda = (\hat{\mathbf{J}} + \lambda \mathbf{I})^{-1} \left[(\hat{\mathbf{S}} - \hat{\mathbf{J}}\theta_\lambda) - (\mathbf{S} - \mathbf{J}\theta_\lambda) \right] = (\hat{\mathbf{J}} + \lambda \mathbf{I})^{-1} \frac{1}{N} \sum_{t=1}^N \tilde{\Psi}_t, \quad (59)$$

with $\tilde{\Psi}_t = \Psi_t(\theta_\lambda) - \mathbb{E}\Psi_t(\theta_\lambda)$ strictly mean zero. (At θ^* the summands have non-vanishing mean $\mathbf{b}_\gamma$, so the CLT cannot be centered there without rate conditions on ε_γ .)

Step 2 (nuisance errors). Expanding in $\hat{\eta} - \eta^\gamma$: by Lemma 6(i)-(ii) the μ -direction and cross-moment have zero first-order term; the Gram part contributes $\mathcal{O}(\sqrt{\varepsilon_\gamma} \|\hat{m} - m^\gamma\|) = o_P(N^{-1/2})$ by (G); second-order remainder is $\mathcal{O}_P(r_N^2) = o_P(N^{-1/2})$ by Assumption 4.

Step 3 (cross-fitting under dependence). By Bradley’s coupling lemma (Bradley, 2005), nuisance fits on blocks separated by τ^* from test blocks match an independent copy except on an event of probability $\leq N\alpha(\tau^*) \leq MN^{1-c_\alpha C} \rightarrow 0$ when $C > 1/c_\alpha$.

Step 4 (CLT). $\tilde{\Psi}_t$ is strictly stationary, mean zero, α -mixing with geometric rate and $(2+\delta)$ moments, so $N^{-1/2} \sum_t \tilde{\Psi}_t \Rightarrow \mathcal{N}(\mathbf{0}, \Omega_\lambda)$ (Ibragimov, 1962).

Step 5 (HMM parameters). With estimated $\hat{\Lambda}$, $N^{-1/2} \sum_t \tilde{\Psi}_t(\hat{\Lambda}) = N^{-1/2} \sum_t \tilde{\Psi}_t(\Lambda^*) + \mathbf{A}_\Lambda \sqrt{N}(\hat{\Lambda} - \Lambda^*) + o_P(1)$. By (H), $\sqrt{N}(\hat{\Lambda} - \Lambda^*) \Rightarrow \mathcal{N}(\mathbf{0}, \kappa \mathcal{I}_\Lambda^{-1})$, asymptotically independent of the evaluation sum across the embargo buffer, so the asymptotic covariances add: $\Omega_\lambda^{\text{tot}} = \Omega_\lambda + \kappa \mathbf{A}_\Lambda \mathcal{I}_\Lambda^{-1} \mathbf{A}_\Lambda^\top$.

Step 6 (Slutsky). $\hat{\mathbf{J}} \rightarrow_p \mathbf{J}$ and $\lambda_{\min}(\mathbf{J}) + \lambda \geq c_\lambda$ yield $(\hat{\mathbf{J}} + \lambda \mathbf{I})^{-1} \rightarrow_p (\mathbf{J} + \lambda \mathbf{I})^{-1}$. (b) is the identity $\theta_\lambda - \theta^* = \text{Bias}$; (c) follows by Slutsky because $\sqrt{N}\|\text{Bias}\| \leq \sqrt{N}(C\varepsilon_\gamma + \lambda M_\theta)/(\lambda_{\min} + \lambda)$. ■

Operational Modes of Forward Filtering. In offline retrospective smoothing, $\gamma_{tk}^{\text{smooth}} = P(S_t = k \mid Z_{1:N})$. In real-time decision making, three distinct forward modes exist: (1) *Oracle Filter*: $P(S_t = k \mid Z_{1:t}, \Lambda^*)$ using the true HMM parameters; (2) *Training-Fitted Forward Filter*: $P(S_t = k \mid Z_{1:t}, \hat{\Lambda})$ where parameters $\hat{\Lambda}$ are fitted on a historical calibration window strictly preceding the evaluation window ($\mathcal{T}_{\text{cal}} < \mathcal{T}_{\text{eval}}$) and state inference updates strictly forward on past-to-present observations $Z_{1:t}$; and (3) *Expanding-Window Online Filter*: where both parameters $\hat{\Lambda}_t$ and state updates proceed sequentially. Our implementations evaluate Mode 1 and Mode 2, where Mode 2 enforces causal forward state filtering without future observation leakage (distinguishing forward state updating from online sequential parameter re-estimation).

E.1 Observed-Data Orthogonal Score under Estimated Latent Dynamics

In semiparametric causal inference (Chernozhukov et al., 2018; Bickel et al., 1993), when nuisance parameters consist of both non-parametric outcome/propensity regressions $\eta = (m, \mu)$ and finite-dimensional parameters Λ governing latent state transitions and emissions, one must characterize the sensitivity of the causal estimator to the first-order estimation error in $\hat{\Lambda}$.

Let the state-weighted regime score vector be:

$$\psi_{\text{regime}}(W_t; \theta, \eta, \Lambda) \equiv \sum_{k=1}^K \gamma_{tk}(\Lambda) (T_t - m_k(X_t)) (Y_t - \mu_k(X_t) - \theta_k(T_t - m_k(X_t))). \quad (60)$$

In a hidden Markov model satisfying standard regularity and mixing conditions (Bickel et al., 1993), the latent parameters Λ are estimated via sequence-level maximum likelihood with sequence score $\nabla_\Lambda \log p(Z_{1:N}; \Lambda) = \sum_{t=1}^N \phi_{\Lambda,t} + o_P(\sqrt{N})$, where $\phi_{\Lambda,t} \equiv \nabla_\Lambda \log p(Z_t \mid Z_{1:t-1}; \Lambda^*)$ denotes the sequence-level prediction-error score increment satisfying $\mathbb{E}[\phi_{\Lambda,t} \mid Z_{1:t-1}] = \mathbf{0}$ and $\frac{1}{N} \sum_{t=1}^N \mathbb{E}[\phi_{\Lambda,t} \phi_{\Lambda,t}^\top] \rightarrow \mathcal{I}_\Lambda$. Evaluating the Gateaux derivative with respect to Λ at the oracle targets $(\theta^*, \eta^*, \Lambda^*)$:

$$\mathbf{A}_\Lambda \equiv \lim_{N \rightarrow \infty} \frac{1}{N} \sum_{t=1}^N \sum_{k=1}^K \mathbb{E} \left[\frac{\partial \gamma_{tk}(\Lambda^*)}{\partial \Lambda^\top} (T_t - m_k^*(X_t)) (Y_t - \mu_k^*(X_t) - \theta_k^*(T_t - m_k^*(X_t))) \right]. \quad (61)$$

Importantly, $\mathbf{A}_\Lambda$ does not vanish pointwise across regimes. While Assumption 3 guarantees that the structural innovation has zero mean conditional on the true state, $\mathbb{E}[Y_t - \mu_k^*(X_t) - \theta_k^*(T_t - m_k^*(X_t)) \mid X_t, T_t, S_t = k] = 0$,

conditioning on observed history $Z_{1:N}$ alone yields:

$$\mathbb{E}[Y_t - \mu_k^*(X_t) - \theta_k^*(T_t - m_k^*(X_t)) \mid X_t, T_t, Z_{1:N}] = \sum_{j \neq k} \gamma_{tj}(\Lambda^*) [(\mu_j^*(X_t) - \mu_k^*(X_t)) + (\theta_j^* - \theta_k^*)(T_t - m_j^*(X_t))]. \quad (62)$$

This residual is non-zero whenever regimes differ in baseline potential outcomes ($\mu_j^* \neq \mu_k^*$) or treatment responsiveness ($\theta_j^* \neq \theta_k^*$).

Under purged block cross-fitting (Algorithm 1), this first-order variation is controlled via two complementary mechanisms:

1. **Decoupling via Purged Embargo Buffer:** The parameters $\hat{\Lambda}$ are estimated on historical calibration blocks or out-of-fold data separated by an embargo buffer $\tau^* \geq C \log N$. By the α -mixing decay of the underlying process (Assumption 4), the covariance between evaluation fold innovations and the parameter estimation error $\hat{\Lambda} - \Lambda^*$ decays as $O(N^{-C\alpha})$.
2. **Stacked Joint HAC Variance Representation:** When $\hat{\Lambda}$ is estimated at parametric rate $\sqrt{N}(\hat{\Lambda} - \Lambda^*) = \mathcal{I}_\Lambda^{-1} \frac{1}{\sqrt{N}} \sum_{t=1}^N \phi_{\Lambda,t} + o_P(1)$, first-order parameter variation is absorbed into the sequence-level influence representation $\phi_t = \phi_{\theta,t} + \mathbf{A}_\Lambda \mathcal{I}_\Lambda^{-1} \phi_{\Lambda,t}$, yielding the stacked M-estimation limiting distribution:

$$\sqrt{N} \begin{pmatrix} \hat{\theta}_\Lambda - \theta^* - \text{Bias}(\lambda) \\ \hat{\Lambda} - \Lambda^* \end{pmatrix} \xrightarrow{d} \mathcal{N} \left(\mathbf{0}, \begin{bmatrix} \Sigma_\Lambda & \Sigma_{\theta\Lambda} \\ \Sigma_{\theta\Lambda}^\top & \mathcal{I}_\Lambda^{-1} \end{bmatrix} \right), \quad (63)$$

where $\Sigma_{\theta\Lambda} \equiv \text{LRCov}(\phi_{\theta,t}, \phi_{\Lambda,t})$, and the long-run covariance Σ_Λ and cross-term $\Sigma_{\theta\Lambda}$ are consistently captured by the joint Newey-West HAC sandwich estimator. In well-separated regimes ($\Delta_Z \rightarrow \infty$), posterior uncertainty vanishes ($\gamma_{tk} \rightarrow \mathbb{I}(S_t = k)$) and $\partial \gamma_{tk} / \partial \Lambda \rightarrow \mathbf{0}$ exponentially fast, rendering $\|\mathbf{A}_\Lambda\| \leq C e^{-c\Delta_Z^2}$ and recovering asymptotic orthogonality.

E.2 Frontier B: Regime-Level Null Hypothesis Size Distortion Under Latent Proxy Error

To investigate the empirical validity of sandwich standard errors when proxy separation is imperfect ($\varepsilon_\gamma > 0$), we conduct a null hypothesis benchmark (**Frontier B**; Table 3). We simulate a DGP where the true regime-specific treatment effect in Regime 1 is strictly zero ($\theta_1^* = 0$, while $\theta_0^* = 2.0$, $N = 1,200$, 100 Monte Carlo replications; `reports/frontier_b_size_distortion_summary.csv`).

Because proxy error creates a non-vanishing mean score shift $\mathbf{b}_\gamma \neq \mathbf{0}$ (Theorem 5), the point estimator converges to $\theta^* + \mathbf{J}^{-1} \mathbf{b}_\gamma$ rather than θ^* . While the sandwich estimator correctly estimates sampling variance around the proxy limit θ_γ , it does not center at θ^* . Consequently, nominal 5% Wald tests reject the true null in 78.0% of runs for Spectral OR-DML (mean bias 0.3420) and 31.0% for Hard Regime FE (mean bias 0.2486). This demonstrates that naive regime-level Wald inference suffers severe Type I error inflation under latent proxy uncertainty, directly justifying selective causal abstention via $\mathcal{D}_t$.

Table 3: Frontier B: Null Hypothesis Size Distortion Under Latent State Proxy Error (True $\theta_1^* = 0$, Nominal $\alpha = 5\%$, $N = 1,200$, 100 Replications). Non-vanishing score shifts induce 78.0% Type I error for sandwich inference (`reports/frontier_b_size_distortion_summary.csv`).

Method	Nominal Size	Empirical Type I Error	Mean Abs Bias
Spectral OR-DML (Sandwich)	5.0%	78.0%	0.3420
Hard Regime FE (Post-Cluster)	5.0%	31.0%	0.2486

Frontier C: Sensitivity to Latent Regime Persistence. In `reports/frontier_c_persistence_stress_summary.csv`, we evaluate latent chain persistence across $\rho \in \{0.30, 0.50, 0.70, 0.88\}$. As persistence increases, annual regime transitions drop ($561.1 \rightarrow 96.0$ switches), enhancing temporal smoothing and reducing OR-DML absolute bias from 1.452 ($\rho = 0.30$) to 0.478 ($\rho = 0.88$), confirming that temporal continuity aids latent regime recovery.

F PROOF OF PROPOSITION 2: JOINT DELTA-METHOD INFERENCE FOR POPULATION ATE

We distinguish three distinct regime weighting vectors:

1. **True Stationary Distribution π^*** : The invariant measure satisfying $\pi^{*\top} \Pi^* = \pi^{*\top}$.
2. **Estimated HMM Stationary Distribution $\hat{\pi}(\Lambda)$** : The stationary vector derived from estimated HMM matrix $\hat{\Pi}$.
3. **Sample Posterior Occupation Distribution $\bar{\gamma}_N$** : The empirical average $\bar{\gamma}_N \equiv \frac{1}{N} \sum_{t=1}^N \gamma_t$.

The sample Average Treatment Effect estimator is $\hat{\theta}_{\text{SATE}} = \sum_{k=1}^K \bar{\gamma}_{N,k} \hat{\theta}_{\lambda,k} = \bar{\gamma}_N^\top \hat{\theta}_\lambda$, with conditional variance $\sigma_{\text{SATE}}^2 = \bar{\gamma}_N^\top \Sigma_\lambda \bar{\gamma}_N$. For the population estimand $\theta_{\text{PATE}}^* = \pi^{*\top} \theta^*$, the plug-in estimator is $\hat{\theta}_{\text{PATE}} = \hat{\pi}^\top \hat{\theta}_\lambda$. Expanding $\hat{\theta}_{\text{PATE}}$ around (π^*, θ^*) via the multivariate delta method: $\sqrt{N}(\hat{\theta}_{\text{PATE}} - \theta_{\text{PATE}}^*) = \pi^{*\top} \sqrt{N}(\hat{\theta}_\lambda - \theta^*) + \theta^{*\top} \sqrt{N}(\hat{\pi} - \pi^*) + o_P(1)$, yielding asymptotic variance:

$$\sigma_{\text{PATE}}^2 = \lim_{N \rightarrow \infty} \text{Var} \left(\sqrt{N}(\hat{\theta}_{\text{PATE}} - \theta_{\text{PATE}}^*) \right) = \pi^{*\top} \Sigma_\lambda \pi^* + \theta^{*\top} \Sigma_\pi \theta^* + 2\pi^{*\top} \Sigma_{\theta\pi} \theta^*, \quad (64)$$

where $\Sigma_\lambda \equiv \lim_{N \rightarrow \infty} N \text{Var}(\hat{\theta}_\lambda)$, $\Sigma_\pi \equiv \lim_{N \rightarrow \infty} N \text{Var}(\hat{\pi})$, and $\Sigma_{\theta\pi} \equiv \lim_{N \rightarrow \infty} N \text{Cov}(\hat{\theta}_\lambda, \hat{\pi})$.

Cross-Covariance Analysis and HMM Influence Orthogonality: When $\hat{\theta}_\lambda$ and $\hat{\pi}$ share estimated parameters $\hat{\Lambda}$, let $\phi_\theta(W_t) \equiv (\mathbf{J} + \lambda \mathbf{I})^{-1} \Psi_t$ and $\phi_{\Lambda,t} \equiv \nabla_\Lambda \log p(Z_t \mid Z_{1:t-1}; \Lambda^*)$. With $\sqrt{N}(\hat{\pi} - \pi^*) = \nabla_\Lambda \pi^* \mathcal{I}_\Lambda^{-1} \frac{1}{\sqrt{N}} \sum_{t=1}^N \phi_{\Lambda,t} + o_P(1)$, the cross-covariance evaluates to:

$$\Sigma_{\theta\pi} = \lim_{N \rightarrow \infty} \frac{1}{N} \sum_{t=1}^N \sum_{s=1}^N \mathbb{E} \left[\phi_\theta(W_t) \phi_{\Lambda,s}^\top \right] \mathcal{I}_\Lambda^{-1} (\nabla_\Lambda \pi^*)^\top = \text{LRCov}(\phi_\theta, \phi_\Lambda) \mathcal{I}_\Lambda^{-1} (\nabla_\Lambda \pi^*)^\top. \quad (65)$$

A sufficient condition for $\Sigma_{\theta\pi} = \mathbf{0}$ is long-run cross-orthogonality $\text{LRCov}(\phi_\theta, \phi_\Lambda) = \mathbf{0}$. In practice, σ_{PATE}^2 is consistently estimated via a joint Newey-West HAC sandwich over the stacked $(2K)$ -dimensional influence sequence $(\phi_\theta(W_t)^\top, \phi_\pi(Z_t)^\top)^\top$, where $\phi_\pi(Z_t) \equiv \nabla_\Lambda \pi^* \mathcal{I}_\Lambda^{-1} \phi_{\Lambda,t}$.

Stationary Distribution Variance vs. Sample Regime Occupancy: The parametric variance of $\hat{\pi}(\hat{\Lambda})$ is $\Sigma_\pi = \nabla_\Lambda \pi^* \mathcal{I}_\Lambda^{-1} (\nabla_\Lambda \pi^*)^\top$. For the sample occupancy process $\bar{S}_N \equiv \frac{1}{N} \sum_{t=1}^N \mathbf{H}_t$, asymptotic covariance is governed by fundamental matrix $\mathbf{Z} \equiv (\mathbf{I} - \Pi^* + \mathbf{1}\pi^{*\top})^{-1}$ (Kemeny and Snell, 1976): $\Sigma_{\text{occ}} = \mathbf{D}_\pi \mathbf{Z} + \mathbf{Z}^\top \mathbf{D}_\pi - \mathbf{D}_\pi - \pi^* \pi^{*\top}$, with $\mathbf{D}_\pi = \text{diag}(\pi^*)$. For a two-state Markov chain ($K = 2$) with persistence $\rho = \lambda_2(\Pi^*)$:

$$\text{Var}(\bar{S}_N) = \frac{\pi_0^* \pi_1^*}{N} \left(\frac{1 + \rho}{1 - \rho} \right) + o(N^{-1}). \quad (66)$$

When $\Delta\theta^* \neq 0$, Markov persistence inflates the sampling variance of $\hat{\theta}_{\text{SATE}}$ by $\frac{1+\rho}{1-\rho} \approx 15.7\times$ under persistent stagnation ($\rho \approx 0.88$) relative to independent transitions. ■

G EMPIRICAL SENSITIVITY ANALYSES AND DATA PROVENANCE

G.1 Meteorological Regime Characteristics across Megacities

Table 4 details the empirical regime profiles uncovered by the Gaussian HMM across all four megacities. Across all airsheds, the inferred regimes correspond directly to distinct synoptic meteorological states: cooler, high-humidity, low-wind nocturnal/winter inversion states versus warmer, lower-humidity daytime convective dispersion states.

Table 4: Meteorological Regime Profiling across Four Indian Megacities from Ground-Sensor Observations.

City	Regime	Share	Persistence	Wind (km/h)	Temp (°C)	RH (%)
Delhi	Regime 1 (Winter Inversion)	42.7%	0.989	5.98	19.63	61.4%
Delhi	Regime 2 (Convective Dispersion)	57.3%	0.992	7.52	30.93	56.3%
Mumbai	Regime 1 (Coastal Stagnation)	41.7%	0.789	11.49	26.69	83.2%
Mumbai	Regime 2 (Marine Breeze / Ventilated)	58.3%	0.857	16.48	28.39	77.9%
Bengaluru	Regime 1 (Cool Nocturnal Inversion)	52.2%	0.843	10.83	20.77	84.6%
Bengaluru	Regime 2 (Warm Afternoon Dispersion)	47.8%	0.828	11.22	27.67	45.3%
Kolkata	Regime 1 (Winter Inversion)	60.7%	0.871	7.32	26.73	85.3%
Kolkata	Regime 2 (Pre-Monsoon Dispersion)	39.3%	0.810	7.89	32.65	51.2%

G.2 Regularization Sensitivity in Ill-Conditioned Regimes (Kolkata Airshed)

To verify Theorem 4 under real conditions, Table 5 reports the sensitivity of regularized DML across $\lambda \in [0.0, 1.0]$ in ill-conditioned Kolkata ($\lambda_{\min} = 0.0024$). Table 5 isolates single-index DML sensitivity across λ (unregularized




Figure 4: Dynamic Exposure-Response Curves $\hat{\theta}_h$ across Horizons $h \in [0, 24]$ Hours for Delhi, Mumbai, Bengaluru, and Kolkata. Shaded bands represent 95% HAC confidence intervals. Regime stratification reveals substantial structural heterogeneity in dynamic pollutant accumulation over multi-hour horizons.

$\lambda = 0.0$ yields -0.5570 ± 1.1071 , matching Block DML), whereas Table 2 reports two-regime OR-DML ($\hat{\theta}_{\text{overall}} = -0.0063 \pm 1.3003$).

Table 5: Regularization Sensitivity Grid for Kolkata Airshed (Single-Index Evaluation, $N = 1,209$, Unregularized $\lambda_{\min} = 0.0049$, Two-Regime $\lambda_{\min} = 0.0024$). Standard errors decrease monotonically from 1.1071 at $\lambda = 0.0$ to 0.0063 at $\lambda = 1.0$, while the regularized condition number $\kappa(\hat{\mathbf{J}} + \lambda \mathbf{I})$ contracts monotonically from $1.28 \rightarrow 1.00$.

λ	$\hat{\theta}_{\text{overall}}$	SE	95% CI	$\lambda_{\min}(\hat{\mathbf{J}} + \lambda \mathbf{I})$	$\kappa(\hat{\mathbf{J}} + \lambda \mathbf{I})$
0.0000	-0.5570	1.1071	$[-2.7269, +1.6130]$	0.0049	1.28
0.0005	-0.5172	1.0170	$[-2.5105, +1.4760]$	0.0054	1.25
0.0024*	-0.4068	0.7776	$[-1.9309, +1.1173]$	0.0073	1.19
0.0100	-0.2191	0.4022	$[-1.0074, +0.5691]$	0.0149	1.09
0.0500	-0.0639	0.1140	$[-0.2872, +0.1595]$	0.0549	1.02
0.1000	-0.0339	0.0601	$[-0.1517, +0.0840]$	0.1049	1.01
1.0000	-0.0036	0.0063	$[-0.0160, +0.0088]$	1.0049	1.00

As predicted by Theorem 4, when $\lambda \rightarrow 0$, standard errors explode to 1.1071 due to near-singularity in residual treatment variation. Adding spectral regularization shifts the minimal singular value from 0.0049 to 0.1049 at $\lambda = 0.10$, reducing standard errors by an order of magnitude ($1.1071 \rightarrow 0.0601$) while shrinking point estimates conservatively toward zero. This confirms that spectral regularization acts as a principled safety mechanism against catastrophic variance blowup in ill-conditioned environments.

G.3 Data Imputation Sensitivity Analysis (Strictly Observed vs. Cleaned)

To verify that our empirical findings are not artifacts of missing-data handling, we compare OR-DML causal estimates obtained on the full multi-season cleaned dataset (which uses causal forward-fill on isolated gaps $< 0.7\%$ of hours) against a strictly observed dataset where every row containing any imputed value is entirely deleted.

Table 6: Sensitivity Analysis Comparing Full Dataset ($< 0.7\%$ Causal Forward-Fill) vs. Strictly Observed Dataset (Rows with Any Imputed Entry Deleted) Across Four Megacities. Point estimates and standard errors exhibit overlapping confidence bands across both sample definitions.

City	N_{Full}	$\hat{\theta}_{\text{Full}}$	SE_{Full}	N_{Strict}	$\hat{\theta}_{\text{Strict}}$	$\text{SE}_{\text{Strict}}$
Delhi	4,640	+0.3742	0.0642	4,602	+0.4014	0.0662
Mumbai	4,180	+1.0806	3.9079	4,146	-0.6048	5.1383
Bengaluru	4,093	-0.0258	0.0790	4,084	-0.0364	0.0741
Kolkata	1,209	-0.0063	1.3003	1,206	+0.1502	1.3835

As documented in Table 6, point estimates retain overlapping confidence bands across all cities under complete-case deletion: Delhi ($+0.4014 \pm 0.0662$ vs. $+0.3742 \pm 0.0642$), Mumbai (-0.6048 ± 5.1383 vs. $+1.0806 \pm 3.9079$, both statistically indistinguishable from zero), Bengaluru (-0.0364 ± 0.0741 vs. -0.0258 ± 0.0790), and Kolkata ($+0.1502 \pm 1.3835$ vs. -0.0063 ± 1.3003). These results confirm that our qualitative conclusions and directional signs are robust to missing data and do not depend on the causal forward-fill protocol.

G.4 HAC Lag Truncation Sensitivity and Physical Telemetry Audit

To evaluate the sensitivity of statistical inference to temporal dependence bandwidths, we evaluate Newey-West HAC standard errors across Bartlett truncation lags $h \in \{12, 24, 72, 168\}$ hours (spanning diurnal, multi-day synoptic, and weekly cycles). As lag truncation increases to capture persistent weather autocorrelation, standard errors widen moderately: in Delhi Regime 1, standard errors rise from 0.0513 ($h = 12$) to 0.0734 ($h = 168$); in Bengaluru Regime 2, from 0.0689 to 0.0931 (p -value widening from 0.0265 to 0.101); in Mumbai Regime 2, from 3.9025 to 4.7150 (p -value widening from 0.065 to 0.152); and in Kolkata, from 1.3003 to 1.4510.

Physical Telemetry Screening. Audit of `data/processed_clean/combined_hourly_clean.csv` identifies physical telemetry artifacts: (1) *Kolkata Discretization*: Kolkata NO₂ records only 31 distinct values across 1,209 hours, with a 376-hour constant run in PM_{2.5}. (2) *Mumbai Constant Runs*: Mumbai exhibits a 273-hour constant run in NO₂ and saturation at 985.0 $\mu\text{g}/\text{m}^3$. (3) *Conditioning Interpretation*: These sensor artifacts deplete residual treatment variance, driving unstandardized $\lambda_{\min}$ to near-zero (0.0025 in Kolkata, 0.20 in Mumbai). Crucially, condition numbers remain modest ($\kappa \in [1.43, 2.52]$), showing that variance inflation stems from low residual variance and stuck sensor runs rather than Gram collinearity.

Table 7: Real-Covariate Semi-Synthetic Benchmark: Evaluated on Megacity Meteorological Distributions with Injected Ground Truth ($\theta^* = (0.75, 2.50)^\top$, $N = 1,200$, 30 Replications). All proxy estimators suffer substantial bias and 0% coverage under real atmospheric sequential confounding (`reports/megacity_semisynthetic_summary.csv`).

Estimator	Delhi Airshed			Mumbai Coastal		
	Bias	RMSE	95% Cov	Bias	RMSE	95% Cov
Oracle DML	0.0134	0.0170	93.3%	0.0126	0.0166	93.3%
Standard DML	6.0154	6.0156	0.0%	5.9961	5.9962	0.0%
Regime FE DML (Hard)	0.7320	2.0807	60.0%	3.5628	3.9008	0.0%
HMM(Z, T) OR-DML	1.2437	1.5954	0.0%	2.5556	2.7378	0.0%
Spectral OR-DML (Ours)	4.0444	4.1188	0.0%	5.2679	5.2804	0.0%

Estimator	Bengaluru Plateau			Kolkata Airshed		
	Bias	RMSE	95% Cov	Bias	RMSE	95% Cov
Oracle DML	0.0148	0.0195	93.3%	0.0332	0.0432	86.7%
Standard DML	5.6822	5.6824	0.0%	5.7069	5.7074	0.0%
Regime FE DML (Hard)	2.4796	2.4851	0.0%	0.3870	0.4161	43.3%
HMM(Z, T) OR-DML	1.4618	1.4638	0.0%	1.0888	1.0961	0.0%
Spectral OR-DML (Ours)	3.9198	3.9204	0.0%	3.7201	3.7218	0.0%

H LATENTREGIMEBENCH: MULTI-DIMENSIONAL BENCHMARK RESULTS

To systematically evaluate machine learning under persistent, partially observed, and shifting latent regimes, we conduct comprehensive evaluations on **LatentRegimeBench** across five diagnostic configurations ($N = 1,200$, 30 Monte Carlo replications each) and four representation architectures (Table 8 and Table 9, with summary visualizations in Figure 5).

H.1 Independent Factorial Validation of the Graceful-Degradation Mechanism

The earlier treatment-mean-separation diagnostic is intentionally not used because correctly specified regime-aware nuisance residualization removes regime mean shifts. Instead, we independently manipulate (i) controlled posterior proxy error ε_γ and (ii) state-specific treatment innovation variance, which changes residual information and hence $\lambda_{\min}(\mathbf{J})$ after oracle regime-aware cross-fitting. The 7×7 design uses 100 independent Markov sequences per cell ($N = 1,200$; 4,900 runs). Across runs, causal L_2 error has Spearman correlation 0.9955 with $\varepsilon_\gamma/\lambda_{\min}(\mathbf{J})$, versus 0.6260 with ε_γ alone, 0.5209 with $1/\lambda_{\min}$ alone, and 0.5355 with an additive combination $\varepsilon_\gamma + 1/\lambda_{\min}(\mathbf{J})$; cell-mean correlations are 0.9963, 0.6249, 0.5294, and 0.5483, re-




Figure 5: Universal LatentRegimeBench Evaluation. (a) Bias calibration against theoretical degradation ratio $\mathcal{D} = \bar{\varepsilon}_\gamma / \lambda_{\min}(\mathbf{J})$. (b) Decision abstention safety comparing unconstrained vs. gated MSE. (c) Representation zoo OOD generalization across HMM, GRU, Transformer, and Linear SSM.

spectively (`reports/factorial_reliability_correlations.csv`). Selective retention of the lowest-difficulty 40% of runs reduces mean causal error from 48.81 to 2.61 and the 95th percentile from 228.10 to 8.39 (`reports/factorial_risk_coverage.csv`). These results validate an interaction, not dominance of either constituent.

Failure score	Run-level ρ	Cell-mean ρ
$\varepsilon_\gamma / \lambda_{\min}(\mathbf{J})$	0.9955	0.9963
ε_γ only	0.6260	0.6249
$1/\lambda_{\min}(\mathbf{J})$ only	0.5209	0.5294
Additive: $\varepsilon_\gamma + 1/\lambda_{\min}(\mathbf{J})$	0.5355	0.5483

H.2 Calibration Intervention and Leave-One-World-Out Reliability

Temperature scaling is fit on a held-out calibration partition and evaluated on a disjoint compound-shift test partition. It sharply improves conventional probabilistic metrics but does not rescue downstream failure discrimination: GRU ECE falls $0.1050 \rightarrow 0.0569$ and NLL $0.5475 \rightarrow 0.3312$, while difficulty ROC-AUC changes $0.5644 \rightarrow 0.5609$; Transformer ECE falls $0.1379 \rightarrow 0.0541$ and NLL $1.0058 \rightarrow 0.3816$, while ROC-AUC changes $0.5356 \rightarrow 0.5282$. Thus calibration is not a sufficient intervention for task reliability. Conversely, probabilistic quality is still informative across domains: in 15-fold leave-one-world-out evaluation, adding NLL to F_1 improves reliability-prediction R^2 from 0.1133 to 0.2207 and Spearman correlation from 0.3742 to 0.5265 (`reports/representation_zoo_lowo_evaluation.csv`).

Table 8: LatentRegimeBench: Stress Test across 5 Strategic Configurations (30 Replications). In well-conditioned regimes (C1, C3), OR-DML achieves $> 5.5\times$ bias reduction over Standard DML; under severe proxy error (C2, C4) or persistence inversion (C5), degradation tracks $\mathcal{D}$.

Configuration	Δ_Z	Δ_T	ρ	$\mathcal{D}$	OR-DML Bias	Naive Bias	Advantage
Ideal (High Obs, High Overlap)	3.5	5.0	0.85	0.001	0.2165	1.2061	$5.57\times$
Adversarial A (Low Obs, High Overlap)	0.3	5.0	0.85	0.709	1.1950	1.2061	$1.01\times$
Adversarial B (High Obs, Weak Overlap)	3.5	0.4	0.85	0.001	0.2498	1.3886	$5.56\times$
Challenging (Low Obs, Weak Overlap)	0.5	0.5	0.90	2.748	1.3237	1.3862	$1.05\times$
High-Persistence Inversion	2.0	2.0	0.96	0.159	1.5934	1.3175	$0.83\times$

Representation Accuracy vs. Downstream Task Reliability. To evaluate whether latent representations translate to reliable prediction under shift, we fit a Ridge predictor $f_m : (X_t, T_t, \gamma_t^{(m)}) \mapsto Y_t$ on training data ($N = 1,200$) and freeze it. Residualized treatments $\tilde{T}_{tk} = T_t - \hat{m}_k(X_t)$ yield rolling coupled Gram matrix $\mathbf{J}_t = \frac{1}{w} \sum_{s=t-w+1}^t (\gamma_s \odot \tilde{T}_s)(\gamma_s \odot \tilde{T}_s)^\top$ ($w = 48$) and minimum eigenvalue $\lambda_{\min}(\mathbf{J}_t)$. On held-out Shift C ($N = 1,200$, 5 seeds), a calibration split ($N_{\text{cal}} = 400$) fixes failure loss threshold $q_{0.90}$ out-of-sample. We evaluate future lead loss $L_{t+1} = (Y_{t+1} - \hat{Y}_{t+1})^2$ against: (i) State Entropy $H_t = H(\gamma_t)$, (ii) Conditioning $1/\lambda_{\min}(\mathbf{J}_t)$, and (iii) Task Difficulty $\mathcal{D}_t = H_t/\lambda_{\min}(\mathbf{J}_t)$. Moving-Block Bootstrap ($B = 1,000$, $L = 24\text{h}$) quantifies uncertainty over dependent series; sensitivity across $L \in \{12, 24, 48, 72\}$ confirms CI widths are stable (0.142-0.147; `reports/representation_zoo_mbb_sensitivity.csv`).

Offline Representation Diagnostics vs. Prospective Failure Signals. We formalize a three-level Reliabil-

Table 9: Latent Representation Zoo under Disentangled Shifts and Downstream Reliability ($N = 2,000$, 5 seeds; mean $\pm$ std). Panel A: Latent recovery under dynamics shift (A, $\rho = 0.85 \rightarrow 0.96$), observation noise (B, $\Delta_Z = 3.5 \rightarrow 1.0$), and compound shift (C). Panel B: Downstream failure anticipation under compound shift (C) with 95% MBB CIs [$L = 24$]. Neural encoders achieve competitive recovery ($F_1 \approx 0.84$ - 0.86) but near-chance failure detection (ROC-AUC ≈ 0.515 - 0.549), whereas Gaussian HMM reaches ROC-AUC = **0.6153 ± 0.0246** (MBB CI [$0.540, 0.683$]). Post-hoc temperature scaling improves ECE/NLL without improving failure ROC-AUC.

Panel A: Latent Representation Zoo under Disentangled Distribution Shifts							
Architecture	In-Dist F_1	Shift A F_1	Shift B F_1	Shift C F_1	ΔF_1 Gap	Shift C ECE	
Gaussian HMM (Causal Filter)	0.9986 $\pm$ 0.0004	1.0000 $\pm$ 0.0000	0.8770 $\pm$ 0.0114	0.9170 $\pm$ 0.0118	+ 0.0817 $\pm$ 0.0115	0.0397 $\pm$ 0.0242	
Neural GRU Encoder	0.9997 $\pm$ 0.0007	0.9997 $\pm$ 0.0004	0.8426 $\pm$ 0.0128	0.8481 $\pm$ 0.0136	+0.1516 $\pm$ 0.0136	0.1047 $\pm$ 0.0129	
Causal Transformer Encoder	0.9997 $\pm$ 0.0004	1.0000 $\pm$ 0.0000	0.8328 $\pm$ 0.0125	0.8372 $\pm$ 0.0105	+0.1625 $\pm$ 0.0108	0.1561 $\pm$ 0.0136	
Linear State Space Model (SSM)	0.9995 $\pm$ 0.0007	1.0000 $\pm$ 0.0000	0.8563 $\pm$ 0.0132	0.8580 $\pm$ 0.0118	+0.1415 $\pm$ 0.0119	0.0558 $\pm$ 0.0159	
Panel B: Downstream Reliability: Evaluating Task-Conditioned Difficulty $\mathcal{D}_t$ across Sequence Encoders (Shift C)							
Architecture	Shift C F_1	ROC-AUC $\mathcal{D}$	MBB 95% CI	Baseline H AUC	PR-AUC $\mathcal{D}$	MDA $\rho(\mathcal{D}, L)$	
Gaussian HMM (Causal Filter)	0.9170 $\pm$ 0.012	0.6153 $\pm$ 0.025	[0.540, 0.683]	0.6209 $\pm$ 0.021	0.1300 $\pm$ 0.017	+ 0.0656 $\pm$ 0.050	
Neural GRU Encoder	0.8481 $\pm$ 0.014	0.5228 $\pm$ 0.036	[0.446, 0.600]	0.5217 $\pm$ 0.031	0.0987 $\pm$ 0.019	-0.0032 $\pm$ 0.040	
Causal Transformer Encoder	0.8372 $\pm$ 0.011	0.5150 $\pm$ 0.055	[0.438, 0.589]	0.5173 $\pm$ 0.063	0.1044 $\pm$ 0.027	+0.0009 $\pm$ 0.056	
Linear State Space Model (SSM)	0.8580 $\pm$ 0.012	0.5492 $\pm$ 0.021	[0.478, 0.625]	0.5521 $\pm$ 0.018	0.0938 $\pm$ 0.027	+0.0369 $\pm$ 0.036	

ity Triangle: (1) *State Identification* (F_1), (2) *Probabilistic Trustworthiness* (ECE, NLL), and (3) *Downstream Reliability* (ROC-AUC, $\mathcal{D}_t$). While retrospective diagnostics (F_1 , ECE) require ground truth, online monitoring relies on prospective signals ($H_t, \lambda_{\min}(\mathbf{J}_t), \mathcal{D}_t$). In Table 9 (Panel B), neural sequence models (GRU, Transformer, SSM) achieve competitive recovery ($F_1 \approx 0.84$ - 0.86) but near-chance failure detection (ROC-AUC ≈ 0.515 - 0.549). In contrast, Gaussian HMM causal filtering achieves well-calibrated states (ECE = 0.0397) and informative failure anticipation (ROC-AUC = **0.6153 ± 0.0246** , MBB 95% CI [$0.540, 0.683$], PR-AUC = 0.1300). In fixed-effects regressions across 15 benchmark worlds ($N = 60$; `reports/representation_zoo_hierarchical_regression.csv`), classification F_1 explains substantial reliability ($\beta_1 = +0.9817 \pm 0.2011, p = 0.0002, R^2_{\text{within}} = 0.4313$), while posterior log-loss provides incremental explanatory power ($\Delta R^2_{\text{within}} = +0.0924, \beta_{\text{NLL}} = -0.0653 \pm 0.0295, p = 0.0442$, reaching $R^2_{\text{within}} = 0.5237$). Calibration ECE exhibits negative association ($\Delta R^2_{\text{within}} = 0.0647$, reaching $R^2_{\text{within}} = 0.4960$). Numerical eigenvalue floor sensitivity across $\epsilon \in [10^{-6}, 10^{-3}]$ confirms invariance (AUC 0.6153).

H.3 Prospective Downstream Failure Forecasting: Multi-Horizon Lead-Time Evaluation

To test prospective reliability at deployment, we strictly utilize prospective information available at time t ($\mathcal{D}_t = H(\gamma_t)/\lambda_{\min}(\mathbf{J}_t)$, $H_t = H(\gamma_t)$, and $1/\lambda_{\min}(\mathbf{J}_t)$ -without access to ground-truth state labels, ECE, or NLL) to forecast extreme downstream forecasting failures (worst decile, top 10% loss $L_{t+h} = (Y_{t+h} - \hat{Y}_{t+h})^2$) across advance lead horizons $h \in [1, 24]$ hours on real-world sensor networks and synthetic streams (Table 10). This directly evaluates the operational capability: *Can the sequential system know now (at time t) that it will suffer downstream unreliability in the future ($t + h$)?*

Table 10: Multi-Horizon Failure Prediction: Task Difficulty $\mathcal{D}_t$ vs. State Entropy H_t across Lead Times $h \in [1, 24]$. In variance-depleted airsheds (Mumbai, $\lambda_{\min} = 0.23$), $\mathcal{D}_t$ outperforms state entropy across all horizons. In Bengaluru ($\lambda_{\min} = 4.29$), raw state entropy dominates, demonstrating complementary, airshed-dependent failure mechanisms.

Airshed / Domain	h	Spearman ρ	Failure ROC-AUC $\mathcal{D}$	Baseline ROC-AUC H	Δ AUC	PR-AUC $\mathcal{D}$
Mumbai Coastal ($\lambda_{\min} = 0.23$)	$h = 1$	+0.0519	0.5309	0.4855	+ 0.0454	0.1078
	$h = 3$	+0.0458	0.5418	0.5020	+ 0.0398	0.1113
	$h = 6$	+0.0175	0.5640	0.5267	+ 0.0373	0.1176
	$h = 12$	+0.0358	0.5366	0.5029	+ 0.0337	0.1088
	$h = 24$	+0.0418	0.5550	0.5303	+ 0.0247	0.1150
Bengaluru Plateau ($\lambda_{\min} = 4.29$)	$h = 1$	-0.1445	0.4249	0.6737	-0.2488	0.0788
	$h = 6$	-0.1504	0.4089	0.6504	-0.2414	0.0768
	$h = 24$	-0.1547	0.4237	0.6744	-0.2507	0.0788
Delhi Basin ($\lambda_{\min} = 111.8$)	$h = 1$	-0.1265	0.3827	0.3293	+0.0533	0.0845
	$h = 3$	-0.1161	0.4303	0.4148	+0.0155	0.1012
	$h = 6$	-0.0841	0.4114	0.4075	+0.0039	0.0858
Synthetic Benchmark ($\Delta_Z = 2.0$)	$h = 1$	+0.0311	0.5382	0.5249	+0.0133	0.1103
	$h = 2$	+0.0127	0.5530	0.5477	+0.0053	0.1175
	$h = 8$	+0.0186	0.4974	0.4910	+0.0064	0.1001

In coastal airsheds with severe residual variance depletion (Mumbai, $\lambda_{\min} = 0.2288$), state entropy H_t acts as a coin flip (ROC-AUC ≈ 0.4855 at $h = 1$), while coupled difficulty $\mathcal{D}_t$ restores failure anticipation across all lead times (0.5640 at $h = 6$). Conversely, in well-conditioned airsheds (Bengaluru, $\lambda_{\min} = 4.29$), state




Figure 6: Macro Financial Regime Transfer. (a) Factor returns by regime. (b) Dynamic difficulty index $D_t = H(\gamma_t)/\lambda_{\min}(\mathbf{J}_t)$. (c) Coverage vs. tail loss frontier under abstention.

transitions dominate, rendering raw entropy H_t superior (ROC-AUC ≈ 0.67), confirming their complementary domain-dependent roles.

I SYNTHETIC FINANCIAL MACRO STRESS TEST WITH REAL COVARIATE CONTEXT

To verify that the latent regime confounding and difficulty framework reflects a general mathematical reality of sequential decision systems under hidden persistent context rather than an idiosyncratic atmospheric feature, we evaluate the Regime Intelligence engine on a synthetic macro financial stress test with real macroeconomic covariate context (daily FRED macroeconomic interest rate, spread, and liquidity indicators; $N = 1,500$ daily trading periods; Figure 6). Regimes reflect low-volatility normal market dynamics versus severe liquidity stress / crisis regimes.

We partition into calibration ($N = 900$, 60%) and test ($N = 600$, 40%) splits. Threshold $\tau^* = 4.0$ is fixed out-of-sample on calibration. Table 11 partitions test periods into risk quartiles $R_t = H(\gamma_t)/\lambda_{\min}(\mathbf{J}_t)$. In Q4 (Severe Risk), tail 95% forecast MSE surges to 4.2931 (vs. 1.7461 in moderate periods) under acute ill-conditioning ($\lambda_{\min} \rightarrow 0.211$). Table 12 reports three-way decision regret for adaptive policy ($D_t \leq \tau$) vs. Always-Active (regret 0.1994) and Always-Fallback (MSE = 11.3749). Calibration threshold $\tau^* = 4.0$ abstains on 1.0% of periods, slashing regret to 0.0025 (**98.74%** reduction) at 99.0% coverage; $\tau = 2.5$ yields **95.94%** reduction while avoiding fallback loss (MSE = 11.3749).

Table 11: Synthetic Financial Stress Test: Dynamic Risk Quartiles ($N = 1,500$). Tail 95% MSE escalates to 4.2931 under causal ill-conditioning ($\lambda_{\min} = 0.211$).

Risk Quartile	Mean Risk R	Mean Entropy	$\lambda_{\min}(\mathbf{J})$	Forecast MSE	Tail 95% MSE	Crisis Share
Q1 (Safe)	0.000	0.000	0.365	0.9278	3.3514	96.5%
Q2 (Low Risk)	0.000	0.000	0.306	0.6191	2.8449	27.7%
Q3 (Moderate)	0.000	0.000	0.324	0.4899	1.7461	3.2%
Q4 (Severe Risk)	19.164	0.103	0.211	0.9419	4.2931	18.9%

Table 12: Three-Way Decision Regret on Financial Macro Regimes ($N = 600$). Risk-Aware Policy ($D_t \leq \tau$) vs. Always-Active (Regret 0.1994) and Always-Fallback (MSE = 11.3749).

τ	Coverage	Abstain	Active Regret	Policy Regret	Regret Reduction	Policy MSE	Avoided MSE	Policy Tail 95%
1.0	95.8%	4.2%	0.1994	0.7542	-278.32%	1.5955	1.1500	3.4752
1.5	96.3%	3.7%	0.1994	0.6635	-232.82%	1.5048	1.2993	3.3296
2.0	97.8%	2.2%	0.1994	0.1911	+4.15%	1.0324	0.8576	3.2681
2.5	98.2%	1.8%	0.1994	0.0081	+95.94%	0.8494	0.4944	3.2581
3.0	98.5%	1.5%	0.1994	0.0081	+95.94%	0.8521	0.4281	3.2581
4.0*	99.0%	1.0%	0.1994	0.0025	+98.74%	0.8496	0.2900	3.2269